
Own Your Cloud

From Cloud Tenant to Infrastructure Owner
a sovereign-Kubernetes story told in chapters

Mahipal

Second draft · September 2026

Meridian scenes are composite fiction. Build-log sections, and every number in them,
were measured on the companion platform repository.

Contents

Prologue — The Invoice

PART I — THE BILL

Chapter 1 — The Great Cloud ROI Myth

Chapter 2 — The Landlord Problem

Chapter 3 — The Sovereign Imperative

Chapter 4 — Kubernetes That Runs Kubernetes

PART II — THE FACTORY

Chapter 5 — The First Cluster

Chapter 6 — Ten Minutes Becomes Two

Chapter 7 — Thirty-Six Percent Before the Kernel

Chapter 8 — Two Lines and a CPU Quota

Chapter 9 — Warmth Is Mandatory

PART III — THE FABRIC

Chapter 10 — Hardcoded IPs and Other Confessions

Chapter 11 — One Label

Chapter 12 — The Model Is Just Another Service

PART IV — THE BRAIN

Chapter 13 — An Agent as a First-Class Object

Chapter 14 — Rules the Model Cannot Break

Chapter 15 — Trust, but Verify

Chapter 16 — The Fleet

PART V — THE DIVIDEND

Chapter 17 — The Same Cluster Twice

Chapter 18 — From the Lab to the Data Centre

Chapter 19 — The Day We Upgraded Everything

Chapter 20 — Reclaiming the Margin

Epilogue — Own Your Cloud

Appendices — The Owner's Reference

Prologue — The Invoice

Is moving to the cloud actually saving you money, or are you just swapping one landlord for another?
— *Reclaiming Margins with Sovereign Cloud, strategic briefing, February 2026*

MERIDIAN · A MONDAY IN SEPTEMBER 2025 · THE CTO'S OFFICE

The invoice came on the first working day of the month, as it always did, and Anita Rao read it the way she had read it every month for four years: not the total, but the slope.

The total told her what Meridian had spent. The slope told her what Meridian was becoming. For the sixth quarter running, the cloud bill had grown faster than revenue. Nobody had made a bad decision. The migration had worked; the teams shipped faster than they ever had on the old virtualization estate; the auditors were satisfied. The bill simply grew — with every new service, every successful launch, every environment someone created on a Friday and never deleted.

The FinOps team had attached their monthly tagging report. The last page listed resources that no team would admit to owning. It was a long page.

She called Vikram Iyer, who ran the platform team, and asked him a question she had been saving.

"If we stopped paying tomorrow," she said, "what would we still own?"

He thought about it properly, which was why she had asked him. "The code. The data, if we could get it out in time. The people who know how it all fits together."

"Machines?"

"Not one. The data centre went in the migration. The virtualization estate is still there for the old systems, but we license that too."

"So we are tenants," Anita said. "Everywhere."

"Good tenants. We pay on time."

She did not smile. She wrote a single line at the top of the invoice, above the total, and underlined it: *What would it take to own this instead of rent it?*

* * *

Four numbers, and a fifth

BRIEFING · FEBRUARY 2026

The strategy deck behind this book makes the case against renting in four numbers. It calls them "*Cloud Waste Reality*":

The deck's category	The deck's figure
Overprovisioning	35% — <i>"Resources provisioned but unused"</i>
Idle resources	28% — <i>"Running 24/7 with no workload"</i>
Orphaned resources	18% — <i>"Forgotten but still billing"</i>
Premium service markup	40% — <i>"Above open-source alternatives"</i>

They are vivid, and they are exactly the kind of number this book will not repeat without checking. The deck names no source for any of them. And they cannot all be shares of the same bill: the first three alone add up to 81%, which would leave one dollar in five doing useful work, before the markup.

PUBLIC RECORD · MARCH 2026

The nearest published figure is smaller, and more modest in how it was obtained. Flexera's *2026 State of the Cloud Report*, a survey of more than 750 cloud decision-makers and users, found that respondents estimate 29% of their cloud spend is wasted — up from 27% the year before, the first increase in five years. Managing cloud spend remained a top challenge, named by 85%.

Twenty-nine percent is a self-estimate, not a measurement. It is still roughly three dollars in every ten. And it is the kind of number a CTO can test: not against an industry survey, but against her own invoice.

How this book keeps score

Meridian is not a real company. It is a composite of the enterprises this book is written for: an organization with a long virtualization history, a cloud bill growing faster than its revenue, and a board asking whether it has a choice. Its two recurring characters — Anita Rao, the CTO, and Vikram Iyer, who leads the platform team — are fiction. Scenes set at Meridian are marked with its name.

Everything else is not fiction, and the book marks that too.

Build log sections come from a real repository: a working sovereign-cloud platform, built from open-source parts, beginning on 11 February 2026. It has a management cluster that manufactures Kubernetes clusters whose machines are virtual machines; a service mesh joining two clusters; and AI agents that run on models on hardware in the room. Every number in a build log was measured on that platform, most of them again on 17 September 2026, while this book was being written. When a measurement contradicted the repository's own documentation, the book reports the measurement and says so.

Public record and **Briefing** sections quote named sources: published reports, or the strategy deck. When a source gives no evidence for a number, the book says that too.

The platform runs on a single laptop — an AMD Ryzen AI 9 HX 370 with 29 GiB of memory and an NVIDIA RTX 4050 laptop GPU with 6 GB of video memory. That is not a data centre, and the book never pretends it is. What transfers from a laptop to a data centre is not scale. It is the pattern, the measurements, and — more often than the vendor brochures would suggest — the failures.

Every chapter ends the same way:

- **The ledger** — what was built, measured, or found, in numbers only where something was measured.
- **Ask your team** — three questions a CxO can put to their own organization the next morning.
- **Open the repo** — the files and commands that let an engineer check the chapter for themselves.

The book is in five parts. **The Bill** is why. **The Factory** is how clusters get made. **The Fabric** is how they talk to each other. **The Brain** is what AI can safely do with them. **The Dividend** comes back to the money.

How to read the build logs

The platform described in the build logs is a real repository, and the **Open the repo** block at the end of each chapter names the files and commands that reproduce what the chapter describes. An engineer with the repository open can check every claim in it.

Most readers will not have it open — and a book whose argument depends on a file the reader cannot see is asking to be taken on faith, which is the one thing this book is trying not to do. So everything load-bearing is quoted inline: the manifest that declares a cluster, the two configuration lines that made it faster, the check that could not fail, the agent defined as a Kubernetes object, the version variable that controlled nothing. Where a quote is too long for the page it sits in, it moves to **Appendix G — Key files, quoted**, with its path as the caption.

Three rules hold for every quoted block. It is verbatim, never paraphrased — if code is worth citing as evidence, it is worth showing as written. Where a file is trimmed to fit, the cut is marked, and nothing is silently tidied: comments that admit uncertainty stay in, because those are frequently the most honest lines in the repository. And a quote is evidence of what the file said on the date given, not a promise about what it says now.

An engineer who wants to run any of it will need the repository. A reader who only wants to know whether the book is telling the truth should not have to.

It starts where Anita did: with an invoice, and a question written above the total.

PART I

The Bill

Why. Business first; the machinery arrives only at the end.

Chapter 1 — The Great Cloud ROI Myth

You're crazy if you don't start in the cloud; you're crazy if you stay on it.

— Sarah Wang and Martin Casado, Andreessen Horowitz, *"The Cost of Cloud, a Trillion-Dollar Paradox"*, 2021

MERIDIAN · A THURSDAY IN OCTOBER 2025 · THE FINANCE COMMITTEE

The chief financial officer had brought the original business case for the cloud migration. It was seven years old, and one line was highlighted in yellow: *lower total cost of infrastructure within three years.*

"I'm not here to blame anyone," he said. "The migration was right. I signed it. But this line was the reason, and it hasn't been true for a while. So what are we paying for now?"

Anita had expected the question. "Speed," she said. "The migration did everything else it promised. Teams get infrastructure in minutes, not months. We launched in two new markets without building anything. That was real, and we'd make the same choice again."

"And the cost?"

"Nobody promised it would stay cheap. We assumed it. The price of an hour of computing didn't go up much. We bought many more hours — because it was easy, and because the business grew. The bill grows with our success."

The CFO looked at the yellow line. "Then I want to know which parts of the bill are buying speed, and which parts are just the rent."

"So do I," Anita said. "Give me a quarter."

* * *

The promise was real

BRIEFING · FEBRUARY 2026

The strategy deck behind this book opens its first chapter fairly. It sets out *"The Promise"* before it gets to the problem: speed and agility, *"rapid deployment, instant scalability"*; cost efficiency through *"pay-as-you-go models, reduced capital expenditure"*; and *"access to cutting-edge services, global infrastructure, continuous updates."*

None of that was marketing. A team that needs a database for a two-week experiment should not wait a quarter for hardware. A company that does not know whether a product will have a hundred users or a million should not buy for a million. A business entering a country where it has no data centre should not build one first. For uncertain, spiky or brand-new demand, renting is not merely convenient. It is the correct economics.

So the myth this chapter is about is not "the cloud is expensive." It is narrower, and more costly to believe: *that renting stays cheaper as you grow.*

The paradox

PUBLIC RECORD · MAY 2021

The clearest statement of the problem came, perhaps surprisingly, from a venture capital firm whose portfolio companies are among the cloud's biggest customers. In "The Cost of Cloud, a Trillion-Dollar Paradox", Sarah Wang and Martin Casado of Andreessen Horowitz put it in one sentence — the epigraph above — and then tried to measure it.

Among the public software companies they benchmarked, committed cloud spend averaged **about 50% of cost of revenue**. Practitioners who had moved workloads back to their own infrastructure told them that doing so cost **one-third to one-half** as much as running the same workloads in the cloud. The authors then estimated what cloud costs were doing to company valuations, through margins: **around \$100 billion** of market value lost across fifty top public software companies, and **more than \$500 billion** across a wider universe.

Treat the valuation figures as what they are: a model, built from margins and market multiples, and contested at the time. The operational claim is the one that matters for a CTO, and it has a simple shape. *At scale, for steady workloads, the rent can cost two to three times the cost of owning.*

Two companies that left

PUBLIC RECORD · 2018 AND 2024

Two companies made the case with their own accounts rather than a model.

Dropbox began moving user files off Amazon's storage service onto its own infrastructure in 2015. Its IPO filing in 2018 disclosed that the project, which it called *Infrastructure Optimization*, had reduced operating costs by **\$74.6 million over two years**.

37signals, the company behind Basecamp and HEY, announced in 2022 that it would leave the cloud. In October 2024 its co-founder David Heinemeier Hansson reported the result. The cloud bill had fallen from an original run rate of **\$3.2 million a year to \$1.3 million**, a saving of almost two million dollars a year, for new servers costing about \$700,000. With storage to follow, the company projected savings of *"well over ten million dollars over five years."* He also addressed the objection every CFO raises: *"There were no hidden dragons of additional workload associated with the exit that required us to balloon the team."*

Read both cases carefully, because they are easy to over-sell. Both companies run large, steady, predictable workloads — file storage and web applications with well-understood demand. Both have unusually strong engineering teams. Neither is a bank with a fifteen-year virtualization estate and a regulator. They prove that the paradox is real for *some* workloads. They do not prove it for yours.

Two taxes

BRIEFING · FEBRUARY 2026

The strategy deck gives the problem two names, and both are useful in a boardroom.

The **Innovation Tax** is *"money spent maintaining old environments instead of building new capabilities."* The deck states the paradox bluntly: *"You're paying a premium for the privilege of being stuck in an ecosystem that limits your ability to innovate."*

The **Success Tax** is the Meridian CFO's complaint, put as a principle: *"The better you perform, the more you pay — disproportionately so."*

The names are worth keeping. Several numbers the deck attaches to them are not, at least not without evidence. It claims a *"2–3x cost multiplier"*, *"15–25% annual increases"*, and that *"80% of IT budgets"* go to *"cloud maintenance vs. innovation."* None of these has a source in the deck. Its headline market figure has a problem too. The deck puts global cloud spend at *"\$1.3T ... in 2025"*. Gartner's published forecast for worldwide public cloud end-user spending in 2025 was **\$723.4 billion**, up 21.5% on the year. The deck may be counting a wider market. Without a source, it cannot be checked.

That growth figure carries the chapter's most practical distinction. A bill can rise because prices rise, or because consumption rises. They call for opposite responses. Rising prices are a supplier problem, answered by negotiation or by leaving. Rising consumption is a management problem, and changing suppliers doesn't fix it. The Success Tax is real only for the share of growth that cost more than the business value it bought.

Waste doesn't care who owns the hardware

The published waste figure from the prologue — 29% of cloud spend, by respondents' own estimate — is often used as an argument for leaving. It is really an argument for managing.

Owned hardware can be wasted too. Chapter 9 of this book measured a *warm standby*: a spare cluster kept running so that a user could claim one in under half a second instead of waiting for one to be built. It is a good design. It also holds eight processor cores and 14 GiB of memory idle, around the clock, waiting. On the platform this book describes, that standby is switched off by default, for exactly that reason.

Idle capacity costs money wherever it runs. The difference is who sends the bill, and whether anyone reads it.

The promise that is worth testing

BUILD LOG · JUNE–SEPTEMBER 2026

If owning is going to compete with renting, it has to match the cloud's signature promise: capacity in seconds, not months. That is testable, and the rest of this book tests it.

On the platform in this book, a complete Kubernetes cluster — a control plane and a worker, each running as its own virtual machine — is built from a written declaration to both machines ready in a median of **33.7 seconds**, measured on 17 September 2026 (Chapter 7). Tearing one down takes **45.7 seconds** (Chapter 17; an earlier single-run figure of 18.9 s did not survive being measured five times). Claiming a pre-built standby took **467 milliseconds**, median of three runs measured 18 September 2026 — about 72 times faster than the cold build (Chapter 9).

Speed, it turns out, is not something only a landlord can offer. But the same chapters also show what it cost to get there. The first measured build took 50 seconds. Four attempts to make it faster failed. The two changes that worked took months of instrumentation to find (Chapter 8). And the dashboard reported the worker as ready for months while it was unable to run anything (Chapter 5).

That is the honest shape of the ROI question. Owning can match renting on speed. What it costs in people and attention is the part the rest of this book counts.

* * *

MERIDIAN · THE FOLLOWING JANUARY

Anita's quarter produced a spreadsheet, not a recommendation. It had one row for each major workload, and four columns.

Shape of demand: steady, seasonal, spiky, or unknown. *Cost trend against the revenue it supports:* flat, rising, or rising faster. *What leaving would take:* the proprietary services each workload depends on. And *what we could test.*

The CFO ran his finger down the first column. "Most of it is steady."

"Most of it is steady," Anita agreed. "Which means most of it is exactly where renting stops being the obvious answer. It doesn't mean owning is the answer. It means we're allowed to ask."

"So what do you recommend?"

"Nothing yet. I'm not going to move a production system on the strength of a venture capital essay and two companies' blog posts." She turned to the last column. "I recommend a test. Small, on hardware we own, measured the way we'd measure a supplier. If it can't build infrastructure as fast as we rent it, we stop."

The ledger

- **The promise:** speed, elasticity and managed services are real, and are the right economics for uncertain or spiky demand.
- **The paradox (public record):** committed cloud spend averaging about 50% of cost of revenue among benchmarked public software companies; repatriation reported at one-third to one-half of the cloud cost for equivalent workloads (Andreessen Horowitz, 2021).

- **The evidence (public record):** Dropbox, \$74.6 million saved over two years; 37signals, a cloud bill cut from \$3.2 million to \$1.3 million a year, with more than \$10 million projected over five years.
- **Not supported:** the deck's 2–3x multiplier, 15–25% annual increases, and "80% of IT budgets" have no source. Its \$1.3 trillion 2025 market figure does not match Gartner's \$723.4 billion forecast.
- **Measured on owned hardware (build log):** a two-VM Kubernetes cluster ready in 33.7 s median; torn down in 45.7 s.

Ask your team

1. For each major workload, is its demand steady, seasonal, spiky or unknown — and who decided?
2. How much of last year's cloud growth came from higher prices, and how much from using more — and did the extra use pay for itself?
3. What would it cost, in services and in months, to move our steadiest workload off its current supplier?

Open the repo

- `docs/sub-60s-cluster-strategy.md` — how cluster build time was measured and brought down, including the dead ends.
- `scripts/phase-timings.sh` (`make phase-timings`) — the per-phase measurement behind the 33.7-second figure.
- `docs/warm-pool-strategy.md` and `ui/backend/handlers/pool.go` — the standby, and what it costs to keep warm.
- The strategy deck, slides 3–9 — its first two chapters. (Not a repository file; see Appendix H — Sources.)

Chapter 2 — The Landlord Problem

You're playing in the vendor's sandbox, and they set the rules on how high you can build.
— *Reclaiming Margins with Sovereign Cloud, strategic briefing, February 2026*

MERIDIAN · A TUESDAY IN NOVEMBER 2025 · THE CTO'S OFFICE

Two letters arrived in the same week, and Anita put them side by side on her desk.

The first was from Meridian's internal audit team. A new control required a particular setting on the Kubernetes clusters that ran the payments platform: a rule that the clusters' control planes check every request against, before it is accepted. The managed Kubernetes service Meridian rented did not offer that setting at the level they used. The supplier's account manager had been helpful and apologetic. The setting might be possible on a higher tier. It might arrive in a future release. It was not on the roadmap anyone could share.

The second letter was the renewal quote for Meridian's virtualization estate — fifteen years of it, running the systems that had never moved to the cloud. The supplier had been acquired. The licensing model had changed. The number at the bottom was one nobody had budgeted for.

Vikram read both. "One's a cloud and one's our own data centre," he said.

"And they're the same letter," Anita said. "In both, somebody else decides what we can change, and what it costs to stay. We don't own either one. We just thought we owned the second."

* * *

The analogy, taken seriously

BRIEFING · FEBRUARY 2026

The strategy deck makes the argument with an analogy most board members will have lived through. *"Managed Cloud = Renting an Apartment"*: convenient, occupied immediately, but *"you cannot renovate the walls, upgrade the infrastructure, or optimize the layout for your specific needs. The landlord controls everything."* Against it: *"Sovereign Cloud = Owning Your Property"*.

It then names the three things that make a landlord a problem:

1. **No architectural control** — *"Cannot customize control planes, networking, or security configurations."*
2. **Rent increases.**
3. **No equity building** — *"Years of payments build no lasting infrastructure value."*

The first is the one Meridian's auditors ran into. The deck puts it most sharply about Kubernetes: *"With managed services like AKS, you cannot customize the control plane — the brain of your cluster. Need specific RBAC configurations for compliance? You might hit a brick wall because vendors abstract that layer to keep things simple."*

That is fair, with one precision worth keeping. Managed Kubernetes services do let customers change some control-plane settings. The problem is that the *supplier* chooses which ones, on the supplier's timetable, at the supplier's price tiers. The wall is not that nothing can be changed. It is that the list of things you are allowed to change belongs to somebody else.

What the analogy leaves out

Analogies persuade by what they leave out, and this one leaves out the boiler.

A tenant who wakes to a broken boiler calls the landlord. An owner fixes it. Owning the brain of your cluster means owning its mistakes too, and this book records plenty of them:

- In September 2026, Kubernetes and everything around it were upgraded on the platform this book describes. It took a week of rebuilding. The rebuild turned up a dead container image that never started, behind a setup script that reported success anyway, a version setting that had never controlled what it claimed to, and a mesh release running outside its tested range. That is Chapter 19.
- For months, the dashboard reported a worker machine as `Ready` while a single configuration line prevented it from running any workload at all. That is Chapter 5.
- A software component pinned to a *moving* version tag quietly updated itself one day and then crashed about every thirty seconds, with no log output, for thousands of restarts before anyone diagnosed it. That is recorded in the repository's own list of pitfalls.

None of these would have been Meridian's problem in a rented managed service. All of them were the platform team's problem here. The landlord problem is real, but so is the owner's problem, and a CxO should be sold both.

Rent increases come from every landlord

PUBLIC RECORD · 2023–2024

Meridian's second letter is not fiction in its shape. In November 2023, Broadcom completed its \$61 billion acquisition of VMware, the virtualization software under a large share of the world's corporate data centres. Broadcom then ended the sale of perpetual licences for VMware's products in favour of subscriptions.

The dispute that made this public was AT&T's. In a lawsuit filed in August 2024, AT&T said Broadcom had proposed a price increase of 1,050% for continued support under its existing arrangement. The two companies later settled.

The lesson for a board is not about any one supplier. It is that "on-premises" and "owned" are not the same thing. An organization can own every server in its data centre and still rent the layer that makes those servers useful. The landlord problem belongs to any single supplier controlling a layer you cannot easily leave, whether that layer runs in someone else's building or your own.

The exit fee is being legislated away. The lock-in isn't.

PUBLIC RECORD · 2024–2027

The deck's third chapter prices the cost of leaving: an *"exit cost multiplier"* of 3–5x, *"18–24 months to migrate"*, and an *"average migration cost"* of \$2–5 million. It gives no source for any of the three, so they are best read as illustrations.

One part of the exit cost is falling, and falling by law. The European Union's Data Act — Regulation (EU) 2023/2854 — entered into force on 11 January 2024, and its obligations became applicable on **12 September 2025**, from which point a provider may recover only its costs for a switch. From **12 January 2027** switching charges are to be withdrawn entirely. The largest cloud providers moved ahead of the law rather than behind it: Google Cloud stopped charging data-transfer fees to departing customers in January 2024, and Amazon Web Services followed in March 2024, with Microsoft following suit — each of them a commercial decision taken before any obligation had bitten.

So the most visible exit fee — paying to take your own data out — is going away, at least for customers covered by the Data Act. What regulation does not remove is the part of the exit cost the deck gets exactly right. It calls this *"The Dependency Spiral"*: basic infrastructure has a *"low exit cost"*; add platform services and there is *"moderate lock-in"*; add software services and there is *"high entanglement"*; adopt the full ecosystem and the exit becomes *"prohibitive"*. Every proprietary database, queue, identity service or AI platform a workload depends on is another thing that has to be rebuilt somewhere else before it can move.

That is where Meridian's real lease lives. Not in the egress price, but in the list of services nobody else sells.

What owning the brain looks like

BUILD LOG · JUNE 2026

The strongest argument for owning a layer is being able to change it the day you need to.

Chapter 8 tells how the platform in this book brought the time to build a cluster down from 50.0 seconds to 34.5 seconds. Every attempt to make the golden image smarter failed. The two changes that worked were one line each, and both were in layers a tenant usually cannot touch:

- **A control-plane startup setting** — telling the cluster's control plane which address to advertise to its workers. That one line removed a ten-second network timeout and took the build from 50.0 to 42.7 seconds.
- **A virtualization-layer resource setting** — raising the processor allowance that KubeVirt gives the small helper containers beside each virtual machine. With a tiny default allowance, the operating system throttled them while each VM started. That took the build from 42.7 to 34.5 seconds.

Neither is dramatic. Both are exactly the kind of change the deck's apartment analogy is about: knocking down a wall. On rented infrastructure, the first would have been a feature request. The second would not have been possible at all.

The owner's ledger

The deck is honest about what ownership demands, and this chapter should be too. Its own list of *"Infrastructure Requirements"* for a self-managed platform is: hardware; data-centre facilities; a platform engineering team of *"Kubernetes experts, DevOps, SRE specialists"*; automation and tooling; and operational processes for incidents, changes and security. It estimates *"12–18 months to full production readiness"* and *"\$2–5M for mid-size enterprise deployment"* — again without a source.

This book adds three items from experience:

- **Upgrades are yours.** On a managed service, a Kubernetes upgrade is a button and a maintenance window. Here it was a week (Chapter 19).
- **Measurement is yours.** Nobody tells an owner how long things take, or whether a dashboard is telling the truth. Chapters 5, 7 and 11 are about building that for yourself.
- **Pinning is yours.** An owner who does not fix software versions in place has quietly handed the landlord's role to whoever publishes the latest release.

Owning does not remove the landlord's work. It moves the work in-house, where it can be done on your schedule and to your standard — or not done at all.

* * *

MERIDIAN · THE FOLLOWING WEEK

Vikram's answer to both letters was a table. One row for every layer of Meridian's technology, from physical buildings up to the AI assistant the service desk was piloting. Four columns: *who supplies it, what we cannot change that we need to, what leaving would take, and who at Meridian could run it if we owned it.*

Anita read the fourth column first. Most of it was blank.

"That's the real finding," Vikram said. "It's not the lock-in. We know about the lock-in. It's that for most of these layers, if the landlord went away tomorrow, we wouldn't have anyone who could fix the boiler."

"Then that's our lease," Anita said. "The second column is what we pay the landlord in control. The fourth is what we'd pay ourselves in people." She circled two rows: the Kubernetes clusters, and the virtualization estate. "Start there. Those are the two where somebody else owns the brain, and where learning to own it ourselves would be worth something even if we never left."

The ledger

- **The landlord problem (briefing):** no architectural control, rising rent, no equity. On managed Kubernetes, which control-plane settings a customer may change is decided by the supplier.
- **Every landlord (public record):** Broadcom completed its \$61 billion VMware acquisition in November 2023 and ended perpetual licences; AT&T alleged a proposed 1,050% increase in a 2024 lawsuit, later settled.
- **The exit fee (public record):** EU Data Act in force 11 January 2024, applicable from 12 September 2025 (cost-only switching charges), with charges withdrawn entirely from 12 January 2027; Google Cloud and AWS dropped exit data-transfer fees in 2024, ahead of the obligation.
- **The lock-in that stays (briefing):** the dependency spiral of proprietary services.
- **Owning the brain (build log):** two one-line changes, in layers a tenant rarely reaches, cut cluster build time from 50.0 s to 34.5 s.
- **The owner's cost (build log):** a week-long upgrade; a dashboard wrong for months; a moving version tag that crashed a service thousands of times.
- **Not supported:** the deck's 3–5x exit multiplier, 18–24 months, and \$2–5 million figures have no source.

Ask your team

1. For each layer of our technology, which settings do we need to change but aren't allowed to — and who decides?
2. Which of our suppliers could raise our price tenfold, and what would we do the next morning?
3. If a supplier disappeared tomorrow, which layers could we run ourselves, and which could we not even diagnose?

Open the repo

- `docs/sub-60s-cluster-strategy.md` — the two one-line changes, and the four that failed.
- `03-target-cluster/target-cluster-warm.yaml` — the control-plane setting, in the manifest (`advertise-address`).
- `scripts/configure-kubevirt-perf.sh` — the virtualization-layer setting (`supportContainerResources`).
- `CLAUDE.md`, "Common Pitfalls" — the owner's side of the ledger, including the upgrade record and the moving-tag crash.
- The strategy deck, slides 5, 8, 9 and 16. (Not a repository file; see Appendix H — Sources.)

Chapter 3 — The Sovereign Imperative

Local-LLM-first sovereignty (the feature enterprise Europe / India / public sector buyers actually ask for)

— *Sovereign_Cloud_Agentic_Strategy.md*, "The Big Picture"

MERIDIAN · A FRIDAY IN DECEMBER 2025 · THE PROCUREMENT REVIEW

The tender was for a public-sector client, and Meridian's bid team had been confident until they reached section eleven.

Section eleven asked where the client's data would be processed. That was a question Meridian had answered many times. It then asked where *derived* data would be processed: logs, support tickets, configuration, and — a new line this year — *"any content submitted to artificial-intelligence services, including prompts and outputs."* All of it had to stay within the country, under operators subject to its law.

The head of the bid team brought the question to Anita. "Our service desk has an AI assistant now," he said. "It reads tickets and logs to suggest fixes. Where does it send them?"

Anita knew the answer before Vikram confirmed it. The assistant called a model hosted by a foreign provider, in a region outside the country. The tickets and logs went with every call.

"It's the most useful thing we've deployed this year," Vikram said. "It's also the most complete map of our operations anyone has ever built, and we're mailing it abroad a paragraph at a time."

* * *

Three kinds of sovereignty

BRIEFING · FEBRUARY 2026

The strategy deck divides sovereignty into three pillars, and the division is the most useful thing in its third chapter:

- **Data sovereignty** — *"Data resides within specific national or regional borders, complying with local laws (GDPR, data localization). Organizations retain legal jurisdiction over their data."*
- **Technical sovereignty** — *"Control over infrastructure, software, and hardware technology stacks. Freedom to customize, optimize, and innovate without vendor-imposed limitations."*
- **Operational sovereignty** — *"Transparency and control over cloud operations. Customer-managed encryption keys, geo-fencing, isolated data centers, and complete operational autonomy."*

The distinction matters because the three fail independently. Data can sit inside the country's borders while being operated by people under another country's jurisdiction. A stack can be entirely open source and still be run by a supplier who decides when it changes. Most "sovereign" offerings deliver one pillar well and describe all three.

The deck also sizes the market: from "\$96.77B" in 2024 to "\$648.87B" by 2033, growing at 23.8% a year. Its 2033 figure matches the current published estimate from Grand View Research, which now puts the market at \$117.53 billion in 2025 and the growth rate at 24.1% a year from 2026 to 2033. Treat both as what market forecasts are: a signal of direction, not a measurement.

Why it is a constraint, not a feature

PUBLIC RECORD · 2018–2027

Sovereignty tends to be sold as a feature. For many of the organizations this book is written for, it is a constraint, imposed from outside, and constraints do not negotiate.

A few public facts explain why:

- **The General Data Protection Regulation** has applied across the European Union since May 2018, and governs how personal data may leave it.
- In July 2020, the EU's Court of Justice, in the case known as **Schrems II**, struck down the framework many companies had relied on to transfer personal data to the United States. A replacement framework was adopted in July 2023. For three years in between, a legal basis that thousands of companies depended on was simply gone. That is the risk a board should notice: not that a rule is strict, but that it can change underneath you.
- The United States' **CLOUD Act**, passed in 2018, allows US authorities to require US-based providers to disclose data in their possession, custody or control, *whether it is stored inside or outside the United States*. Where the data sits is not the same as whose law reaches it.
- India's **Digital Personal Data Protection Act** was passed in 2023.
- The **EU Data Act**, whose cloud-switching rules began applying in September 2025, makes it a legal requirement that customers can move between providers (Chapter 2).

None of this is legal advice, and the details belong with counsel. The pattern belongs with the board. For regulated organizations, where data is processed, and by whom, is decided partly by legislatures and courts. The architecture has to be able to answer that decision, not merely hope for a favourable one.

The AI question makes it urgent

Meridian's section eleven is where many organizations are meeting sovereignty for the first time. The reason is AI.

An operations assistant is only useful if it can see the estate: logs, configuration, alerts, identities, network addresses, and the history of what went wrong. It is, in Vikram's words, the most complete map of your operations anyone has ever built. When that assistant calls a model hosted elsewhere, every prompt carries part of the map with it.

That is why the strategy document behind this book's platform lists **local-LLM-first** as its first non-negotiable, and treats it as a business decision before a technical one:

Treat any cloud LLM as an optional accelerator for specific non-sensitive personas (e.g. documentation drafting), never the default.

Two more of its non-negotiables turn AI operations into something a regulator can inspect. Every action an agent takes must be recorded, because *"who told the agent to do what, and what did it do" is a regulatory record, not a debugging aid.* And the model that produced each output must be recorded too — *"If you ever need to reproduce a decision three years later, you'll need this."*

The document is also candid about who is asking. The combination it is building is *"the feature enterprise Europe / India / public sector buyers actually ask for."*

What sovereignty looks like when you measure it

BUILD LOG · 17 SEPTEMBER 2026

Sovereignty is easy to assert and tedious to check. This is what the platform in this book actually does, checked on one day in September.

Where it holds:

- **The models are in the room.** All three of the platform's serving agents send their model calls to Ollama, running on the laptop's own GPU, through an address inside the cluster (Chapter 12). The host holds fourteen models, and none of the agents calls a hosted model service.
- **A cluster boots without the internet.** Everything a new cluster needs — Kubernetes, its container images, its system settings — was downloaded once, when the golden image was built. Nothing is downloaded when a cluster is created (Chapter 6).
- **The mesh's root of trust is local.** The certificate authority that issues every workload identity in the service mesh is generated on the platform itself. There is no external certificate authority, *"and it also means the demo works air-gapped"* (Chapter 11).
- **The software is open source.** Cluster API, KubeVirt, k3s, Istio, Symposium and Ollama are all open projects. No component requires a licence key from a supplier.

Where it leaks:

- **The observability stack is downloaded at install time.** Kiali and Prometheus are fetched from a public GitHub branch when Act 5 runs, not from a local copy (Chapter 12).
- **The model gateway has no lock.** Anything that can reach it can list, inspect and call the models. The model server itself listens on every network interface of the host (Chapter 12).
- **The agents' network rule is looser than it reads.** Agent egress is allowed on a few ports to *any* destination, not to named ones (Chapter 14).
- **Some keys are in the repository.** The cluster certificate authorities used by the fast-build image are committed to the repository, private keys included, under the label *"demo use only."*
- **No provenance is recorded.** Nothing in the agent manifests records which model, or which version of its weights, produced an answer. The non-negotiable exists on paper.

No one of these is a scandal on a demonstration platform, and each has a known fix. Together they make the point this chapter exists to make: **sovereignty is not something you buy. It is an audit you keep passing.** A platform can be built from open source, run on owned hardware and host its own models, and still ship a map of its estate to anyone on the network who asks.

The price of keeping it at home

There is a cost to local-LLM-first, and a board should hear it before a vendor hides it.

The models that fit on hardware you own are smaller than the frontier models available as a service. This book's Part IV shows what that means in practice:

- A 3-billion-parameter model, when its tool silently failed, invented plausible cluster output and presented it as fact (Chapter 15).
- A 7-billion-parameter model, told in its instructions to refuse a class of question, did not reliably refuse (Chapter 14).
- The same 7-billion-parameter model, in a scored triage exercise, found the correct root cause of a planted fault in 3 runs out of 3 (Chapter 15).
- The laptop's GPU holds one model at a time, so which model is "warm" is a shared resource that agents compete for (Chapter 9).

Sovereign AI today means smaller models, owned hardware sized for them, and more verification around every answer. None of those is a reason not to do it. All of them are reasons to budget for it.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Anita's reply to section eleven was two sentences long, and she wrote them before she knew how Meridian would deliver them.

All processing of client data, including derived data and content submitted to AI services, will take place in-country, on infrastructure operated by Meridian staff under national law. AI services will use models hosted on that infrastructure, and every AI action will be recorded, with the model that produced it.

"We can't do the second sentence today," Vikram said.

"I know. That's why I wrote it." She pushed the page across. "Turn the service-desk assistant off for this client. Then find out what it takes to run it on a model we host ourselves. It'll be a smaller model. It'll make more mistakes. So we check every answer it gives, until we know how often it's wrong."

"And if the smaller model isn't good enough?"

"Then we will have learned that for the price of a pilot," Anita said, "instead of for the price of a tender."

The ledger

- **Three pillars (briefing):** data, technical and operational sovereignty fail independently; most offerings deliver one.
- **The market (briefing / public record):** the deck's \$648.87 billion 2033 estimate matches Grand View Research's current forecast (\$117.53 billion in 2025; 24.1% a year from 2026). A forecast, not a measurement.
- **The constraint (public record):** GDPR (2018); Schrems II struck down an EU-US transfer framework (2020), with a replacement in 2023; the US CLOUD Act (2018) reaches data by provider, not by location; India's DPDP Act (2023); the EU Data Act (cloud switching from September 2025).
- **Held (build log, 17 September 2026):** all three agents on local models; clusters that boot without the internet; a locally generated mesh root of trust; an open-source stack.
- **Leaked (build log):** observability downloaded at install; an unlocked model gateway and a model server on every interface; port-only agent egress; demo certificate keys committed; no model provenance recorded.
- **The price:** smaller models, one model at a time on this GPU, and verification around every answer.

Ask your team

1. For each AI service we use, where do the prompts and outputs go — and under whose law?
2. Which of our "sovereign" claims have we checked this quarter, and which are we still asserting?
3. If a court struck down the legal basis for one of our data flows tomorrow, which systems would have to stop?

Open the repo

- `Sovereign_Cloud_Agentic_Strategy.md` — §4 "Sovereignty Considerations (Non-Negotiables)" and §7 "The Big Picture".
- `06-symposium/cluster2-agent.yaml` and the other agent manifests — `baseUrl`, where model calls go.
- `07-istio-advanced/01-install/gen-mesh-ca.sh` — the locally generated root of trust.
- `06-symposium/agent-egress-ollama.yaml` — the agents' egress rule, and its ports.
- `git ls-files 03-target-cluster/warm-ca` — the committed demo certificate authorities.
- `07-istio-advanced/act5-observability/run.sh` — where observability is downloaded from.

Chapter 4 — Kubernetes That Runs Kubernetes

Cluster API (CAPI) is elegantly defined as "using Kubernetes to manage Kubernetes."

— Reclaiming Margins with Sovereign Cloud, strategic briefing, February 2026

MERIDIAN · A SATURDAY IN JANUARY 2026 · THE BOARD STRATEGY DAY

The board gave the technology session forty minutes, and Anita had been told to use no more than one slide.

The slide showed three roads. On the left, *stay with what we know*: Meridian's virtualization supplier, which now also sold Kubernetes built into its platform. In the middle, *rent the pieces and run them ourselves* on one cloud. On the right, *build our own*.

The chair asked the question she had been waiting for. "I understand the left-hand road. I understand renting. What is the right-hand one, in a sentence? What would we actually be building?"

"A factory," Anita said. "Not for products — for computing environments. We write down what we want: how many machines, what size, which software. The factory reads the order and builds it, and it keeps checking that what exists still matches what we wrote. If a machine dies, it builds another. If we change the order, it changes the environment."

"And who builds the factory?"

"The same factory builds itself, the same way. That's the point, and it's the part that sounds like a trick." She looked around the table. "The technology is called Cluster API. It's Kubernetes, managing Kubernetes."

* * *

A factory for clusters

BRIEFING · FEBRUARY 2026

Kubernetes is already a factory of a kind. The strategy deck explains it in one line: "*Normally, you write YAML to say 'I want 3 web servers running' — Kubernetes makes it happen automatically.*" You declare what you want, and a control loop keeps working until reality matches the declaration.

Cluster API applies the same idea one level down. In the deck's words: "*With CAPI, you write code to say 'I want 3 Kubernetes clusters with these hardware settings' — and it creates them.*" A central **management cluster** "*acts as the command center, bootstrapping workload clusters across providers.*"

The words that matter for a board are *declaration* and *reconciliation*:

- **Declaration** means a cluster is described in a file. The file can be reviewed, approved, versioned and audited like any other document. The question "how is the payments cluster configured?" has an answer you can read.
- **Reconciliation** means the management cluster keeps comparing the description with what exists, and acting on any difference. A cluster is not built once and then left to drift. It is continuously held to its description.

The deck calls CAPI "*a universal translator*", because the same declaration style works whether the machines come from a public cloud, a private one, or bare metal. A different **provider** plugin translates the declaration for each kind of infrastructure. That is where the escape from lock-in lives: not in the promise that nothing changes when you move, but in the fact that the *shape* of the declaration stays the same.

The spectrum

BRIEFING · FEBRUARY 2026

The deck lays the options out as a spectrum from convenience to control, using three real products as markers:

- **VMware Tanzu, "The Familiar Choice"** — aimed at *"organizations with 15+ years of VMware estate"*; Kubernetes built into the virtualization platform they already run. *"No staff retraining required."* The warning attached: *"High vendor lock-in persists. If Broadcom raises prices or you want bare metal elsewhere, you're still tangled in proprietary web."*
- **CAPZ, "The DIY Power User"** — the open-source Cluster API provider for Azure. *"You manage Azure VMs, networking, and cluster resources yourself. No Microsoft support ticket for internal cluster logic."*
- **Kubermatic Kubernetes Platform (KKP), "The Multi-Cloud Orchestrator"** — *"a unified control plane across AWS, Azure, GCP, and on-premises infrastructure."*

The deck's summary line: *"From safe but locked (Tanzu) → powerful but complex (CAPZ) → truly sovereign (KKP)."* Past the end of the spectrum, it adds a fourth option — *"The Self-Managed Sovereign Platform": "building your own cloud platform using open standards."*

Two cautions before a board uses this slide. First, the labels are the deck's opinions, not measurements; *"truly sovereign"* in particular depends on who runs the platform, under what licence, and with what support. Second, the four options are not exclusive. The platform in this book is the fourth kind, built from the same open-source Cluster API the middle options use. Meridian's three roads share more pavement than the slide suggests.

What the slide leaves off

NOT INDEPENDENTLY VERIFIED · SEPTEMBER 2026

A three-marker spectrum is a slide, not a market. A board evaluating this decision will meet several other options, and it is worth naming them so that nobody mistakes the deck's three for the field.

A caution that applies to this whole section. What follows is a map, assembled from each project's own public description of itself. **None of it was installed, benchmarked, or otherwise verified for this book.** Only one option here was actually built and measured — the one the rest of the book is about — and a reader should treat every other line as a pointer to evaluate, not a finding.

Capabilities change release by release; check current documentation before deciding anything.

- **Harvester** is the closest thing to a productized version of what this book builds by hand: an open-source hyperconverged infrastructure product from SUSE that packages KubeVirt with storage and networking, and integrates with Rancher for cluster management. An organization that wants this architecture without assembling it is looking for this category. The trade is the usual one — less to build, less to change.
- **OpenStack** is the long-established open-source answer to "run our own infrastructure-as-a-service," predating the Kubernetes era and widely deployed in telecommunications and public-sector estates. It solves a broader problem than this book's platform and is correspondingly larger to operate. Organizations that already run it should be asking how Kubernetes sits *on* it, not whether to replace it.
- **Talos Linux** takes a different cut at the same sovereignty argument: a minimal, immutable, API-managed operating system built only to run Kubernetes, with no shell and no package manager. It addresses the layer this book mostly inherits — the golden image and what is inside it (Chapter 6).
- **Distributions** — RKE2, OpenShift, Anthos, EKS Anywhere — are the "supported Kubernetes you run yourself" tier, and they are the most likely real alternative for an enterprise with a support requirement. They differ enormously in licence and price, and two of them are paths back toward the vendor whose bill started this book. That is not an objection; it is the trade being made deliberately, which is the whole subject of Chapter 2.
- **OpenTofu and Terraform** are the declarative alternative to Cluster API itself, and the comparison is genuinely interesting for a board. Both describe infrastructure as code. The difference is *where the loop lives*: Terraform and OpenTofu run a plan when a person or a pipeline runs them, while Cluster API runs a controller inside a cluster that continuously reconciles reality against the declaration. Chapter 18 makes the case for the continuous kind, and Chapter 19 shows what happens without it. Many organizations use both — one to build the foundation, the other to manage what runs on it.

The book's own choice is not presented as the winner of this list. It is the option that best matched one set of constraints — open source throughout, virtual machines as first-class objects, no dependency on a provider's control plane — and the rest of the book is what happened when someone actually built it.

Why virtual machines still matter

For an organization like Meridian, the most important word on the Tanzu marker is *estate*. Fifteen years of virtualization means thousands of virtual machines running systems that will never be rewritten as containers: vendor appliances, licensed databases, old applications nobody dares to touch.

A Kubernetes strategy that has nothing to say about virtual machines asks those systems to wait for a future that will not arrive. That is why the platform in this book is built on **KubeVirt**, an open-source project that runs virtual machines *as Kubernetes objects*. A virtual machine sits beside containers, is declared in the same kind of file, and is scheduled, networked and monitored by the same system.

The platform then does something a little dizzying with that. The workload clusters it manufactures are Kubernetes clusters **whose machines are themselves virtual machines**, run by KubeVirt on the management cluster. Kubernetes manages virtual machines, which run Kubernetes.

One honest limit belongs here, before Part II begins. The repository contains no tooling for *moving* existing VMware virtual machines — none of the usual import or conversion tools appears anywhere in it. This book shows the destination, and measures it. The migration of an existing estate is a separate project, and this book does not pretend to have done it.

That tooling does exist, and a board should know the names even though this book cannot report on them. **Forklift** is the open-source migration project for bringing virtual machines from VMware and other sources into a KubeVirt environment; **virt-v2v** is the older, lower-level converter that turns a virtual machine's disk into a form a different hypervisor can boot. Neither was run here, so this book has nothing to say about how well they work on a real estate, how long a migration takes, or what fraction of machines need hand-holding — and those three questions, not the destination architecture, are what determine whether a migration is affordable. Treat the names as the start of an evaluation. An organization serious about this decision should migrate a representative sample early, because the answer will shape the business case far more than the thirty-four seconds in Chapter 7.

The factory, running

BUILD LOG · 17 SEPTEMBER 2026

This is the factory as it stood on the day this chapter was checked.

The machinery is open source, with version numbers. On the management cluster: Cluster API v1.14.2; the KubeVirt infrastructure provider (CAPK) v0.11.2; the k3s bootstrap and control-plane providers v0.3.0; KubeVirt v1.9.0; and CDI v1.66.0, which imports disk images for the virtual machines.

The order is a 258-line file. The default declaration, `03-target-cluster/target-cluster-warm.yaml`, contains seven objects:

Object	What it declares
<code>Cluster</code>	The cluster itself, and what it is made of
<code>KubevirtCluster</code>	The infrastructure it runs on — here, KubeVirt
<code>KThreesControlPlane</code>	The control plane: one k3s server, version v1.37.0+k3s1
<code>KubevirtMachineTemplate</code> ×2	The virtual-machine shapes, one for the control plane and one for workers
<code>MachineDeployment</code>	The worker pool: how many, from which template
<code>Secret</code>	A pre-shared join token, so the worker can boot in parallel with the control plane

The factory built what the order said. Cluster API recorded two Machines. Each is backed by a KubeVirt virtual machine, `target-cluster-cp-s9j8h` and `target-cluster-workers-rz9xt-m5x6z`, both `Running` and ready. From the moment the file is applied, both machines are ready in a median of **33.7 seconds**, and the whole cluster is torn down in **45.7 seconds** (Chapters 7 and 17).

The factory fits in a laptop. The management cluster, the cluster it built, a second cluster for the service mesh, and the AI agents all ran on one host (the prologue lists it). The manufacturing pattern is the same one a data centre would use. The capacity is not.

The gauge that disagrees

One more reading from the same day belongs in a chapter for a board, because it is the first appearance of this book's most persistent theme.

While the target cluster was serving workloads normally — including the measurements in Chapter 10, run that morning — Cluster API's own summary of it read:

```
AVAILABLE    False
ControlPlaneAvailable=False (WaitingForKthreesServer)
ControlPlaneComponentsHealthy=False Following machines are reporting control plane errors
```

The factory reported its product as *not available* while the product was working. The k3s control-plane provider, at version 0.3.0, is a young project, and the cause was not investigated for this chapter. Chapter 5 will find the mirror image of the problem: a status that said *Ready* about a machine that could not do its job.

The lesson for a board is the same in both directions. A factory's gauges are part of the factory. They need to be checked against the product as carefully as the product is checked against the order.

A smaller example of the same drift sits at the top of the repository. Its front-page architecture diagram still shows the platform as it was in March — a namespace, addresses and a cross-cluster proxy that no longer exist (Chapters 10 and 12). Documentation, like a gauge, describes the factory only until the factory changes.

From money to machinery

Part I has argued from the invoice outwards: that renting stops being obviously cheaper for steady work (Chapter 1); that every landlord controls something you need (Chapter 2); and that for some organizations, where data and AI live is not a preference but a constraint (Chapter 3). This chapter is the bridge: a way to own the brain of your clusters, declaring them in files and manufacturing them on hardware you control.

The rest of the book is the machinery, measured. **Part II, The Factory**, is how the clusters are made, and how long it takes. **Part III, The Fabric**, is how they talk to each other. **Part IV, The Brain**, is what AI agents can safely do with them. **Part V, The Dividend**, returns to the money.

For a CxO, four questions carry through all of it. How long does the factory take to build? How long to destroy? What exactly does the order say? And how do you know the gauges are telling the truth?

* * *

MERIDIAN · THE SAME SATURDAY, FORTY MINUTES LATER

The board did not choose a road. It chose a test.

"Small," the chair said. "Owned hardware, open-source software, a single team. Show us that the factory can build an environment from a written order faster than we can rent one, and show us how you measured it. Then we'll talk about roads."

Anita wrote the commitment into the minutes herself, in words the board could hold her to.

The platform team will build a small cluster factory on hardware Meridian owns, using open-source Cluster API. It will demonstrate an environment built from a written declaration, and report how long it takes, measured, by February.

Vikram read it over her shoulder. "*Measured* is doing a lot of work in that sentence."

"It's the only word in it I care about," Anita said.

The ledger

- **The idea (briefing):** Cluster API extends Kubernetes' declaration-and-reconciliation model to whole clusters, run from a management cluster, with providers for different infrastructure.
- **The spectrum (briefing):** Tanzu, CAPZ and KKP as markers from convenience to control, and a fourth option — self-managed on open standards. The labels are the deck's opinions.
- **Virtual machines (build log):** KubeVirt runs VMs as Kubernetes objects; this platform's workload clusters run on VMs. There is no VMware migration tooling in the repository.
- **The factory, 17 September 2026 (build log):** Cluster API v1.14.2, CAPK v0.11.2, k3s providers v0.3.0, KubeVirt v1.9.0, CDI v1.66.0. A 258-line declaration of seven objects produces two VM-backed machines, ready in 33.7 s median and torn down in 45.7 s.

- **Found:** Cluster API reported the working cluster as `Available=False` ; the repository's architecture diagram still shows March.

Ask your team

1. Can we point to the file that describes each of our production environments — and does the environment still match it?
2. Which of our virtual machines will never become containers, and what is our plan for running them next to the ones that will?
3. When a status screen says an environment is healthy or unavailable, how do we know it is right?

Open the repo

- `03-target-cluster/target-cluster-warm.yaml` — the order: seven objects, 258 lines.
- `kubectl --context kind-cluster2 get clusters,machines,kubevirtmachines -A` and `kubectl --context kind-cluster2 get vm -A` — the factory and what it built.
- `kubectl --context kind-cluster2 get clusters.cluster.x-k8s.io target-cluster -o wide` — the gauge.
- `02-capi-init/init-management-cluster.sh` — how the factory itself is set up.
- `README.md` — the architecture diagram, and how far it has drifted.
- The strategy deck, slides 14–16; the recorded factory demonstration (8 min 56 s), end to end.

PART II

The Factory

The cluster-provisioning story. Every number in this Part was measured.

Chapter 5 — The First Cluster

Spins up a full Kubernetes cluster where each node is a Virtual Machine running inside our management cluster.

— `demo.sh`, "What this demo does"

MERIDIAN · A WEDNESDAY IN SEPTEMBER · THE FOURTH-FLOOR BOARDROOM

The newest member of the board had asked a fair question at dinner — *what exactly did we buy?* — and Anita Rao had promised him the demo that had won the pilot its funding in February.

"From the beginning," she told Vikram. "The version where the machine explains itself."

Vikram started the script. The screen filled with a banner, a sentence about virtual machines running inside a management cluster, and a checklist.

```
✓ KubeVirt installed
✓ CDI installed
○ Container registry (172.18.0.2:5000) — not ready
```

Then nothing. The script had stopped.

"It's the demo," Vikram said, into the silence. "Not the platform. The platform's fine."

The board member smiled politely. Anita did not.

"Then show him the platform," she said, "and afterwards, tell me how the thing we show people stopped working without anyone noticing."

* * *

BUILD LOG · FEBRUARY–SEPTEMBER 2026

The demo that stopped at step one

`demo.sh` was written on 11 February 2026, the platform's first day in the repository, and last changed on 28 February. It is a good piece of work: nine steps, each announced with a banner and a paragraph explaining what is about to happen and why — check the prerequisites, make sure the load balancer is running, make sure the Cluster API providers are ready, pre-pull the VM image, show the golden image, deploy the cluster, wait for the VMs, wait for the API server, verify.

Run on 17 September, it printed two green checks and stopped at the third.

The cause is Chapter 6's. The script confirms the registry is up by sending a plain HTTP request to `172.18.0.2:5000` — the alias every image reference uses, which only the container runtime knows how to translate. A plain request to that address now reaches a Kubernetes node with no registry

on it. The check fails, and because the script is written to stop on the first error, the demo ends before it begins.

Nothing in the demo is wrong about February. Its narration describes VMs with two CPUs and 4 GiB of memory; today's default VMs have four cores and 8 and 6 GiB. It pre-pulls the image tagged `:latest`, which Chapter 6 found was last built in March. The platform kept moving, and the story about the platform stayed where it was.

What follows walks the same ground on the running platform instead.

Three clusters, one laptop

The platform runs on a single workstation, and it is built from three Kubernetes clusters.

`cluster1` and `cluster2` are Kind clusters: Kubernetes whose "nodes" are Docker containers. `cluster1` hosts the service-mesh demonstrations in Part III. `cluster2` is the **management cluster** — the factory. It runs Cluster API, which manages clusters as Kubernetes objects; KubeVirt, which runs virtual machines as Kubernetes workloads; CDI, which imports disk images; MetalLB, which hands out load-balancer addresses; and the agent platform from Part IV.

The third cluster, `target-cluster`, is the product. It is a k3s cluster whose control plane and worker are virtual machines running *inside* `cluster2`.

The repository builds the factory in numbered stages — `00-prereqs` installs the `clusterctl` tool, `01-metallb` the load balancer, `02-capi-init` the Cluster API providers, `03-target-cluster` the target cluster, `04-verify` a health check — with `make all` running the first four. As Chapter 19 found, the stages assume a few things already exist: the Kind clusters themselves, KubeVirt, and CDI are installed by hand.

One apply

The default target cluster is one file: `03-target-cluster/target-cluster-warm.yaml`, 258 lines declaring seven objects.

Object	What it declares
<code>Cluster</code>	the cluster exists, and which objects implement it
<code>KubevirtCluster</code>	its machines are KubeVirt virtual machines
<code>KThreesControlPlane</code>	a one-node k3s control plane
<code>KubevirtMachineTemplate</code> ×2	the shape of each VM: cores, memory, disk image
<code>Secret</code>	the worker's startup configuration
<code>MachineDeployment</code>	one worker, kept at one

Before it is applied, a script seeds four more secrets — certificate authorities and a join token — so the cluster's identity is known in advance (Chapter 8 explains why). Then `kubect1 apply`, and the file's work is done. From here on, the cluster exists only as a declaration, and controllers make it true.

Four controllers take part: Cluster API's core controller, the KubeVirt infrastructure provider, and k3s's bootstrap and control-plane providers — with KubeVirt itself underneath. Counted on the running platform on 17 September, those seven objects and four secrets had produced about eighteen more: a bootstrap configuration, a MachineSet, two Machines, two KubevirtMachines, two VirtualMachines, two running VM instances, two launcher pods, a LoadBalancer Service holding the API address `172.18.255.215`, and five further Secrets, including the kubeconfig that grants access to the new cluster.

Followed through its owner references, the control plane's chain reads from the top down:

```
Cluster/target-cluster
└─ KThreesControlPlane/target-cluster-control-plane
   └─ Machine/target-cluster-control-plane-qkck1
      └─ KubevirtMachine/target-cluster-cp-s9j8h
         │ (no owner reference)
         VirtualMachine/target-cluster-cp-s9j8h
            └─ VirtualMachineInstance/target-cluster-cp-s9j8h
               └─ Pod/virt-launcher-target-cluster-cp-s9j8h-...
```

One link is missing. The VirtualMachine has no owner reference back to the KubevirtMachine that created it; the provider manages it without one. So deleting a cluster is not a single cascade that Kubernetes performs on its own. It is a sequence of controllers noticing, finalizing, and deleting, one object at a time — which is part of why, in Chapter 19, tearing a cluster down took nearly nineteen seconds.

Six layers deep

The chain above describes ownership. The picture below describes where the software actually runs, and it is the thing to hold in mind for the rest of Part II:

1. **A workstation.**
2. **A Docker container**, `cluster2-control-plane` — a Kind node running Kubernetes 1.37.
3. **A pod in that cluster**, `virt-launcher-target-cluster-cp-...`, whose job is to run one virtual machine.
4. **A virtual machine** inside the pod: QEMU running Ubuntu 24.04.5, with four cores and 8 GiB.
5. **A Kubernetes node** inside the VM: `target-cluster-cp-s9j8h`, running k3s v1.37.0 and its own container runtime.
6. **Pods** on that node — today, CoreDNS, the metrics server, a storage provisioner.

The layers show up in small, strange ways. The target cluster reports its nodes' internal addresses as `10.244.0.115` and `10.244.0.114`, and those are not addresses in the target cluster at all: they are pod addresses in the management cluster, because each VM lives in a pod. Everything outside reaches the target cluster's API through one load-balancer address that `cluster2` hands out.

Each layer is a boundary, and the virtual-machine layer is the one that matters for tenants: it gives each target cluster its own kernel, not just its own containers.

The victory lap

`show-cluster.sh`, written on 17 February, is the demo's second half. It walks every object in the chain, looks inside the target cluster, then proves the cluster works by deploying a web server. It was run as-is on 17 September.

Most of it still told the truth. It listed both VMs — four cores and 8 GiB for the control plane, four cores and 6 GiB for the worker — and the load-balancer Service at `172.18.255.215`. Some of it had aged. Its disk section announced that *"each VM disk is a PVC cloned from the golden image via CDI DataVolumes"* and then found none, because disks have been container images since February (Chapter 6).

It also reused a kubeconfig file it found in `/tmp`, written the previous afternoon for a cluster that had since been deleted and rebuilt six times. The file still worked. The warm image's fixed certificates (Chapter 6) give every rebuilt cluster the same identity, so a credential issued for yesterday's cluster opens today's. For a single-cluster demo that is convenient. It is also a credential that outlives the thing it was issued for.

Then the web server. The script created a namespace, deployed two replicas of nginx, and printed its heading for the result:

```
==> Pods (should be spread across nodes):
  nginx-demo-...-7kmpx   1/1   Running   0    8s    10.42.1.7    target-cluster-cp-s9j8h
  nginx-demo-...-pf24z   1/1   Running   0    8s    10.42.1.6    target-cluster-cp-s9j8h
```

Both replicas ran on the control plane. Neither ran on the worker. The script then reported success — *"✓ nginx responded: <title>Welcome to nginx!</title>"* — from a request that went to `localhost` inside one of the pods, which proves that nginx starts, and nothing about the second node or the Service in front of it. The namespace was deleted afterwards, as the script instructs.

Ready, but not allowed to work

The worker was not busy and not broken. It was forbidden.

```
target-cluster-workers-rz9xt-m5x6z:
  taints=[{"effect": "NoSchedule",
           "key": "node.cloudprovider.kubernetes.io/uninitialized",
           "value": "true"}]
```

That taint tells the Kubernetes scheduler to place no ordinary workload on the node. On 17 September the worker ran nothing at all: every pod in the target cluster — the four system pods and both web servers — was on the control plane.

The cause is one line in the worker's startup configuration, the Secret in the manifest:

```
kubelet-arg:
- cloud-provider=external
```

That setting tells the worker's kubelet that an external *cloud controller* will initialize the node, and to keep the node tainted until it does. But the control plane is configured with its cloud controller disabled. Nothing on the platform ever initializes the worker, so the taint never comes off.

And every signal says the worker is fine. Its Machine is `Running`. Its node is `Ready`. Its provider ID matches the Machine's, so Cluster API links them. The verification script lists it as `Ready`. The stopwatch in Chapter 7 waits for it. By every automated measure on the platform, the target cluster has two working nodes. It has one — and a second virtual machine holding four cores and 6 GiB of memory that has been given nothing to do.

This static worker configuration was introduced in June, to let the worker boot at the same time as the control plane (Chapter 8), and the same line is in the default path's manifest today. This chapter verified only the cluster that was running on 17 September; whether every cluster since June has had an idle worker was not checked, and the fix — removing the setting, or running something that initializes nodes — has not been tried.

The only check on the platform that ever gave the worker real work was the victory lap. Its heading said the pods *should be spread across nodes*. It printed that they were not, and declared success anyway.

* * *

MERIDIAN · THE FOLLOWING MONDAY

"It's a two-node cluster," Vikram said, "with one node allowed to work. It's been that way since June, as far as I can tell. Every dashboard we have says both are Ready."

Anita let that sit. "And the demo?"

"The demo stopped because of a registry address. The victory lap still runs. It even prints the evidence — it says the pods should be on both nodes and shows them on one. Nobody reads that line."

"So we have been timing, to a tenth of a second, the arrival of a machine that isn't allowed to do anything." She was quiet for a moment. "Change one word on the dashboard. Stop saying *Ready*. Say *working*, and don't show it until something we gave that node has actually run on it."

"And the demo?"

"Fix it, and run it from the first step every Monday." She stood. "The thing we show people is the thing people believe. It had better still be true."



Figure 5.2 — Six layers: a pod inside a VM inside a pod inside a container.

The ledger

- **Declared:** seven objects in a 258-line manifest, plus four pre-seeded secrets.
- **Produced by controllers:** about eighteen more — Machines, VMs, launcher pods, a load-balancer Service, and Secrets including the new cluster's kubeconfig.
- **Nesting:** six layers, from the workstation to a pod inside a VM inside a pod inside a container.
- **Found on 17 September:** the origin demo stops at its first check; the victory lap narrates a disk design replaced in February; a day-old credential opens a newly rebuilt cluster; and the worker node, `Ready` by every measure, carries a taint that keeps every ordinary workload off it.
- **Idle capacity:** one VM, four cores and 6 GiB.
- **Still owed:** a demo that runs from step one; a definition of "done" that schedules work on every node; a worker configuration that does not wait for a cloud controller the platform does not run.

Ask your team

1. When a dashboard says a system is ready, what did anyone check besides the system's own report?
2. Does our flagship demo still run today, from the first step?
3. How much capacity are we paying for that has never done any work?

Open the repo

- `demo.sh` — the nine narrated steps; `show-cluster.sh` — the victory lap.
- `00-prereqs` through `04-verify`, and `make all`.
- `03-target-cluster/target-cluster-warm.yaml` — the seven objects, and the worker's `kubelet-arg` in its bootstrap Secret.
- On the management cluster: `kubectl get cluster,kthreescontrolplane,machinedeployment,machine,kubevirtmachine,vm,vmi`.
- On the target cluster: `kubectl get nodes -o jsonpath='{range .items[*]}}{{.metadata.name}}{{.spec.taints}}{"\n"}}{end}'`.
- **Without the repo:** Appendix G.1 quotes the manifest's seven objects, including the worker bootstrap Secret and its `kubelet-arg`.

Chapter 6 — Ten Minutes Becomes Two

The golden image pre-bakes k3s binary, airgap images, CA certificates, and kernel modules into a DataVolume, reducing cluster spin-up from ~10 minutes to ~2 minutes.

— `README.md`

MERIDIAN · A WEDNESDAY IN SEPTEMBER · THE CTO'S OFFICE

The supply-chain questionnaire from the regulator had forty questions, and Anita Rao had flagged one.

For each new compute environment you create: where does its software come from, and when was it obtained?

"Walk me through a new cluster," she said.

"Nothing is downloaded when a cluster is created," Vikram said. "Everything it needs — Kubernetes, its container images, its system settings — was downloaded once, when we baked the golden image. Every cluster boots from a copy of that image. No internet at boot."

"That's a good answer for a sovereign platform." She wrote it down. "Second half. When was the image baked?"

"The one we deploy from, yesterday. The one the standby cluster uses —" He checked. "The third of March."

Anita stopped writing. "So some of the clusters we create today are six months old the moment they're born."

"In their software, yes."

"Then an image isn't a file," she said. "It's a decision someone made on the day it was baked, copied into every cluster after. I want to know what's inside each one, and when it was decided."

* * *

BUILD LOG · FEBRUARY–SEPTEMBER 2026

Why a cluster took ten minutes

A virtual machine has to get its software from somewhere. In the platform's earliest form, each VM booted a stock Ubuntu cloud image and then set itself up: fetched k3s and the container images it needs, and prepared itself to run a cluster — over the network, on every boot, for every node.

On 28 February 2026 the repository gained two scripts that changed that: `bake-golden-image.sh` and `build-containerdisk.sh`, committed together as "*optimized images for cdi and container images*." The README states the result in one line — cluster spin-up "*from ~10 minutes to ~2 minutes*" — and

the same pair of numbers appears in the developer guide and the strategy document. The user guide is more cautious: "*~2-5 minutes (with golden image) or ~10 minutes (without).*"

Those figures deserve the same honesty as the rest of Part II. They are stated in four documents, and no run log or script output in the repository records the measurements behind them. They date from before the platform had an instrument (Chapter 7 tells that story). What is not in doubt is the direction, and the reason: the slowest thing a new machine can do is download its own software.

Baking

Baking a golden image is a four-step job run by `bake-common.sh`, and on 16 September it was run from scratch during the rebuild in Chapter 19.

1. **Clone the base image.** A stock Ubuntu Noble cloud image, already imported into the cluster as a `DataVolume` by the Containerized Data Importer, is cloned into a new volume. On 16 September: about 50 seconds.
2. **Boot a bake VM** from the clone, with a cloud-init script that does the real work.
3. **Wait for the VM to finish and power itself off.** On 16 September: 160 seconds.
4. **Delete the VM and keep the disk.**

Inside the VM, the script's own log lines tell the recipe. It installs the k3s binary. It downloads k3s's air-gapped container images. It writes the systemd units k3s will run under. It pre-loads kernel modules and sets kernel parameters. It masks services that slow a boot down and are useless here — `snapped`, `multipathd`, the automatic `apt` timers, `motd-news`, unattended upgrades. Depending on the image variant, it starts k3s once to warm its datastore, then stops it. It creates the login user and an air-gapped install script. And finally it resets cloud-init's state and the machine ID, so that every VM booted from this disk believes it is a new machine, and powers off.

That last step is worth noticing for what it does *not* reset. It wipes the operating system's memory of the bake VM, not k3s's. The warm image's datastore still contains a Node object for the machine it was baked on — the ghost that, in Chapter 7, fooled a stopwatch into stopping early.

The sovereignty consequence is the one Vikram gave the regulator. The internet is touched during the bake — k3s from its install endpoint, the air-gapped images from the project's release page — and never again. The cluster manifests declare `airGapped: true`, and every cluster after that boots from what the image already holds. The dependency on the outside world moves from every boot of every node to one event, on one day, that can be reviewed.

A disk that ships like a container

A baked disk still has to reach each new VM. The first approach cloned the golden `DataVolume` for every virtual machine — a full copy of a 20 GiB volume per node. `build-containerdisk.sh` replaced that with a KubeVirt `containerDisk`: the VM's disk packaged as an ordinary container image.

The script does it in four steps. A helper pod mounts the baked volume and converts the disk to a compressed `qcow2` file. The file is copied out. A two-line container image is built around it:

```
FROM scratch
ADD --chown=107:107 disk.qcow2 /disk/disk.qcow2
```

And the image is pushed to the platform's local registry. On 16 September, the rebuilt warm image came out at **1.35 GB — against a 20 GiB volume**.

A containerDisk behaves like any container image. It is pulled to a node once and reused; it needs no per-VM volume clone; and it is ephemeral — every start of the VM begins from the pristine image, which is why, in Chapter 7, a restarted VM counts as a genuine first boot.

Measured on 17 September, pulling a 1.29 GB containerDisk that was not yet on the node took **4.2 seconds** from the local registry, and pulling it again once cached took **0.1**. That is the value of the Makefile's `pre-pull` targets on this host: a few seconds per image, paid once. Worth having. Not, as Chapter 8 showed, where the time in a cluster build actually goes.

The containerDisk script's header makes its own promise: *"~30-60 second cluster spin-up (vs 1.5-2.5 min with DV cloning)." It was written before the instrument existed. When the repository wrote down its baseline in June, cold builds took 90 to 150 seconds, with the worker waiting for the control plane before it even started. The images were not the reason. Chapter 8 tells what was.*

The address that isn't

Every image reference on the platform names its registry the same way:

```
172.18.0.2:5000/ubuntu-noble-k3s:warm
```

`172.18.0.2` is not the registry's address. It never has been.

The two Kind clusters and the registry all run as containers on one Docker network, and Docker hands out addresses in the order containers join it. `.2` goes to whichever joined first. On 17 September the network looked like this:

cluster1-control-plane	172.18.0.2
cluster2-control-plane	172.18.0.3
registry	172.18.0.4

Every image reference on the platform points at the control-plane node of `cluster1` — a Kubernetes node that runs no registry. Before the rebuild in Chapter 19, the repository's notes recorded `.2` as `cluster2`'s node instead, and in June as `cluster1`'s. The address in the image names has moved from one neighbor to another and back. In every record the repository keeps, it has never been the registry's own.

Images still pull, because the container runtime on each `cluster2` node is given a small translation file:

```
[host."http://172.18.0.4:5000"]
capabilities = ["pull", "resolve", "push"]
skip_verify = true
```

containerd reads it and quietly sends every request for `172.18.0.2:5000` to `172.18.0.4:5000` instead. Nothing else knows about the file. Asked directly from the same node on 17 September, a plain HTTP request to the alias failed to connect at all; the same request to the registry's real address returned `200`.

That arrangement fails in two ways, and both have happened.

The registry's address drifts. When containers reconnect to the network, the registry can come back at a different address. The translation file still points at the old one, and new VMs sit in `ImagePullBackOff`, failing to reach `172.18.0.2:5000`.

The file disappears. It lives inside the node's container, so recreating a node — as Chapter 19 did to both clusters — deletes it.

The repair is `scripts/fix-registry-hosts.sh`, added in June. It looks up the registry's current address and rewrites the translation file on every `cluster2` node. The Makefile runs it before the deploy and pre-pull targets. The UI's own deploy path does not.

The alias claimed one more casualty. The UI has a Registry tab that lists images, and it dialed the alias the way any program would — over plain HTTP, without containerd's translation. It reported `failed to connect to registry` and showed the registry as disconnected, both from the host and from inside the cluster. The fix split one setting into two: `REGISTRY_URL`, a real address to dial, and `REGISTRY_ALIAS`, the name to *display* — because, in the handler's words, the alias *"is what image references use."*

The platform works. But hundreds of references name a machine that is not the registry, and every consumer that needs to reach it must be told, separately and locally, where the registry really is.

An image is a decision, frozen

By 17 September the registry held three tags of the platform image, and the gap between them is the chapter's real subject.

`:warm` was built on 16 September, the day of the upgrade. `:preinit` was built on 18 April. `:latest` was built on 3 March. After the rebuild, only `:warm` had been pulled onto the new `cluster2` node.

Everything inside an image was decided on the day it was baked, and every cluster built from it inherits those decisions:

- **A Kubernetes version.** Chapter 19 had to rebake `:warm` because every tag still carried k3s v1.31, and the variable the bake script documented for choosing the version turned out never to have controlled anything.
- **A ghost.** The warm image's Node object for its own bake VM (Chapter 7).

- **Fixed certificates and a join token**, baked into the warm image on purpose, safe only while one cluster runs at a time.
- **An Ubuntu base** as it stood on the day it was imported.

And the checks that decide whether to rebuild an image look only for its presence. `ensure-warm-image` asks whether an image tagged `warm` exists on the node, not what it contains. The standby cluster in Chapter 9 is built from `:latest` — whose name promises "the newest," and whose contents are six months old.

For a CxO, a golden image is what a regulator would call a baseline: the approved state every new system starts from. It earns its keep only if its age and contents are as visible as the systems it produces.

* * *

MERIDIAN · THE FOLLOWING WEDNESDAY

Vikram's answer to question fourteen ran to a table: one row per image, with the date it was baked, the Kubernetes version inside, the base operating system, and what had been downloaded to build it.

Anita read the row for `:latest` twice. "March."

"It's scheduled for a rebake, or for deletion."

"And the address in every image name that belongs to another machine?"

"It works because each cluster node has a file that translates it. We've added the file to the list of things a rebuild must recreate."

"Then call it what it is." Anita made a note in the margin. "Anything that only works because something else translates it for us is a dependency. It goes on the dependency list, not in a script comment." She signed the page. "And every image gets a birth certificate: what's in it, where it came from, and the date after which we stop building from it."

"An expiry date?"

"Every cluster we build is as old as the image it came from," she said. "I'd like to decide how old that's allowed to be."

The ledger

- **Claimed:** cluster spin-up from about ten minutes to about two with the golden image — stated in four documents, with no measurement recorded.
- **Measured (16–17 September):** base-image clone about 50 s; bake VM 160 s; a 20 GiB disk packaged as a 1.35 GB containerDisk; a cold pull of a 1.29 GB containerDisk in 4.2 s from the local registry, 0.1 s once cached.
- **Moved:** the platform's internet dependency, from every node's boot to one bake.

- **Found:** every image reference names an address that belongs to a Kubernetes node, not the registry; pulls work through a per-node translation file that the UI's deploy path never refreshes.
- **Frozen inside the images:** a Kubernetes version, a ghost Node, fixed certificates, and — in the tag the standby uses — a build date of 3 March.
- **Still owed:** a recorded bake date and contents for each image; checks that verify what an image contains, not just that it exists; a registry reference that is the registry's own address.

Ask your team

1. When was the image behind our newest system built, and what is inside it?
2. Which of our systems only work because something else quietly translates a name or an address for them?
3. When a new machine boots, what does it download, and from whom?

Open the repo

- `bake-golden-image.sh` and `bake-common.sh` — the bake; the in-VM recipe is the sequence of `[bake] Step log` lines.
- `build-containerdisk.sh` — from baked volume to container image.
- `scripts/import-base-images.sh` — the base Ubuntu images the bake starts from.
- `scripts/fix-registry-hosts.sh` and `make registry-fix` — the alias translation, rewritten on demand.
- `ui/backend/handlers/registry.go` — `REGISTRY_URL` versus `REGISTRY_ALIAS`.
- `docker network inspect kind` — who actually holds `172.18.0.2` today.

Chapter 7 — Thirty-Six Percent Before the Kernel

`scripts/time-to-ready.sh` prints ONE aggregate number, which is why two optimisation dead ends ... each cost a full bake cycle to disprove.

— the header of `scripts/phase-timings.sh`

MERIDIAN · A MONDAY IN SEPTEMBER · THE PLATFORM TEAM'S ROOM

The number on the wall display was `48.5 s`, in type large enough to read from the door. Anita Rao stood under it for a while before anyone noticed she was there.

"That's the cluster build?" she asked.

"Apply to ready," Vikram said. "Every time we change something, we run it and watch that number."

"And when it doesn't move?"

"Then the change didn't work."

"Or it worked, and something else got slower." She turned to him. "If finance sent me our cloud bill as one number — no line items, no accounts, just a total — what would I do with it?"

"Send it back."

"So why is this one allowed to be a single number?"

Vikram looked at the display. "Because nobody has built the itemized version."

"Then build that first," Anita said. "Before anyone tries to make it smaller."

* * *

BUILD LOG · SEPTEMBER 2026

One number

For most of 2026 the platform's cluster-build time was measured by `scripts/time-to-ready.sh`: seventy lines of shell that note the time, apply the cluster manifest, poll until two nodes report `Ready`, and print one line.

```
RESULT: 48.520s (2 nodes Ready)
```

It is a perfectly reasonable script, and it shaped a summer of decisions. On its word, two golden-image experiments were built, baked, measured, and abandoned — the dead ends Chapter 8 tells in full. Each took a full bake cycle to disprove, because a single total cannot say *which part* of a cluster

build a change affected, only whether the whole moved. It also could not separate a real improvement from a quiet afternoon: an early run of the warm image recorded 35 seconds that turned out to be *"quiet-host luck rather than a real improvement."*

In September, before anyone was allowed to optimize anything, the brief described in Chapter 8 made one step mandatory: *"Step 1 (mandatory, before any optimisation): instrument the phases."* It listed the phases it wanted — VM creation to running, kernel start to cloud-init finished, k3s start to the API's first healthy answer, worker join, each node Ready — and it distrusted its own starting estimate: *"The 17/17/18 split above is a coarse estimate — confirm or correct it before acting on it."*

A cluster you cannot log into

The brief also suggested how to see inside the virtual machines: log in over SSH and run `systemd-analyze`, since *"SSH key and user `ubuntu` are baked in."*

That was not possible. On this cluster, the KubeVirt API service that `virtctl console` and `virtctl ssh` depend on reported `Available=False`. The obvious way into the guests was closed.

So the instrument was built from what the platform already writes down, without logging in to anything. Its header lists every source:

- **Creation timestamps** on the Cluster API and KubeVirt objects — how long the controllers took to turn a manifest into a virtual machine.
- **Phase transition timestamps** on each VM instance — scheduled, running.
- **Container start times** inside each VM's launcher pod — the disk container, the console container, the one that runs QEMU.
- **The guest's serial console**, which KubeVirt already copies into a pod log — kernel, initramfs, systemd, cloud-init, and k3s milestones, as the guest printed them.
- **Each node's Ready transition** — the metric itself.
- **A probe on the host** that tries the control plane's virtual IP every quarter-second, recording the first TCP connection and the first healthy `/readyz` answer.

The result, `scripts/phase-timings.sh`, is read-only and needs nothing installed inside the guest. It also has a `--collect-only` mode that changes nothing at all and reconstructs the table for whatever VMs are already running.

Anchoring a clock that starts at zero

One problem had an elegant answer, and it is worth a moment.

A guest kernel does not timestamp its log lines with the time of day. It uses seconds since boot — `[15.738316]` — so a serial console full of milestones says how long each took after the machine started, but not *when* the machine started. Without that, the guest's timeline cannot be placed on the same clock as the Kubernetes objects around it.

The answer was already in the log. Cloud-init prints lines that carry *both* clocks at once:

```
running 'init' at Wed, 09 Sep 2026 06:44:19 +0000. Up 6.81 seconds.
```

Subtract the uptime from the wall-clock time, and you have the moment the kernel started, to the precision of the log. Every other console timestamp can then be placed on the same timeline as the objects, the pods, and the host probe. No clock synchronization, no agent, no login.

Small lies in the plumbing

Most of the instrument's 320 lines are not measurement. They are defenses against its own data, and each one is recorded in a comment because each one produced a wrong answer first.

- **The login prompt interleaves with the kernel log.** A console line can read `target-cluster-cp login: [ 15.738316] cloud-init[976]: . . .`. A search anchored to the start of the line *"silently drops exactly the late lines we care about."* The fix was to search anywhere in the line, and to require a decimal timestamp so that a process ID like `cloud-init[976]` is not mistaken for one.
- **Two clocks disagree by a second.** VM phase timestamps and container start times are both rounded to the second, and are written by different components. Mixing them *"yields negative durations."* The launcher pod's own creation time became the anchor for that phase instead — and even so, a one-second rounding can still produce a small negative number in a table.
- **The disk container is gone by the time you look.** It runs once and exits, so its start time lives in a different status field from the console container's, which keeps running.
- **A background probe can freeze its own caller.** The host probe is started in the background and its process ID is read back through command substitution, which waits until every process holding the output pipe has exited. Without redirecting the probe's output, the script *"deadlocks the caller before it can even apply the manifest."*

None of this is glamorous, and all of it is the job. An instrument is mostly the code that stops it from lying.

The first reading

On 9 September the instrument ran three clean builds of the warm path, re-measuring the baseline in the same session as the brief required. The baseline was not 48.5 seconds but 50.0 — runs of 52.7, 50.0, and 46.6. Here is where those seconds went, median of three:

Phase	Seconds	Share
<code>kubect1 apply</code> → VM instance created (Cluster API controllers)	6.6	14%
Launcher pod created → QEMU running	10.0	21%
Firmware, kernel, initramfs	3.7	8%
systemd and cloud-init → k3s start	5.5	12%
Control-plane k3s start → first <code>200</code> from <code>/readyz</code>	7.0	15%
Worker joins, after the API is already answering	11.8	25%
Median time-to-ready	50.0	52.7 / 50.0 / 46.6

Two findings came out of the first reading, and together they redirected the whole effort.

The coarse estimate had given the control plane's k3s startup seventeen seconds. The instrument measured seven. Both image experiments had been built on the theory that k3s startup was the slow part — and that theory had bet on a phase worth fifteen percent of the total.

And the top two rows: **thirty-six percent of the wall clock — 16.6 seconds — elapsed before the guest kernel printed its first character.** No change to the golden image could touch any of it, because none of it happens inside the image. Of those seconds, the table labelled eight as *"containerDisk unpack"* — a label that Chapter 8 shows was a guess, and wrong.

What the team did with that table is the next chapter. This chapter has one more thing to say about the instrument itself, because it was discovered while this chapter was being written.

* * *

BUILD LOG · 17 SEPTEMBER 2026

The stopwatch that stopped early

After the platform upgrade in Chapter 19, the same instrument reported a cluster build of **23.9 seconds** — 23.9, 24.1, and 23.7 across three runs, a spread of only 0.4 seconds. Ten seconds faster than before the upgrade, and steadier. The number went into the platform's operating notes, into two chapters of this book, and into an answer about how the platform got under forty seconds.

It was wrong, and the instrument's own output said so to anyone who looked closely.

In all three tables, and in a fourth run later that day, one row was empty:

```
[WORKER]
```

```
...
```

```
E node Ready
```

```
- -
```

And the total at the bottom of each table — labelled "*TOTAL time-to-ready (both nodes)*" — was identical, to the tenth of a second, to the **control plane's** Ready time. Reading the code explained why. The total is the later of the two nodes' Ready times, and when one of them is missing, it is silently treated as zero. With the worker's timestamp absent, the "both nodes" total was the control plane alone.

That raised a second question: why was the worker's timestamp missing? Queried by hand the next morning, the same worker node reported its Ready time without trouble — seven seconds after the control plane's. So the instrument had looked for it *before the worker was Ready*. And it had looked then because its wait loop — like the older one-number script — stops when *two nodes* report Ready.

On 17 September, the platform's canonical timer, `time-to-ready.sh`, was run three times and reported 26.1, 24.9, and 25.6 seconds, each with "*(2 nodes Ready)*". Immediately afterwards, the cluster's node list was read by name. It did not contain a worker at all. It contained the control plane and this:

```
node ubuntu-bake-vm-warm    Ready at 2026-09-16T07:10:37Z
```

`ubuntu-bake-vm-warm` is the virtual machine the warm golden image was *baked* on. Chapter 6's warm image carries a copy of k3s's datastore from bake time, and that datastore still contains the bake machine's Node object — with a `Ready` status stamped at 07:10 on 16 September, while the image was being baked. Every new cluster booted from that image begins life believing it already has a second, healthy node.

The repository knew about the ghost. Since June, the control plane has run a cleanup job at startup that deletes it: a background unit that tries every two seconds, for up to two minutes, and does not block anything while it waits. That keeps the cluster tidy. But it runs *concurrently* with every stopwatch that counts Ready nodes. In every run measured on 17 September, when the control plane turned Ready, the ghost was still there. The count reached two, and the clock stopped — with the real worker still booting.

Measured again on 17 September, naming each node rather than counting them, three clean runs from `kubecttl apply`:

	Run 1	Run 2	Run 3	Median
Control plane Ready	23.4 s	25.1 s	24.7 s	24.7 s
Worker Ready — both nodes	30.4 s	35.1 s	33.7 s	33.7 s

In every run, at the moment the control plane first reported Ready, the node list read: control plane `Ready`, `ubuntu-bake-vm-warm Ready`. No worker.

The upgraded platform builds a two-node cluster in **33.7 seconds**, median of three. Before the upgrade the figure was 34.5. The upgrade did not make the cluster ten seconds faster. Within the noise of three runs, it changed nothing.

The tight 0.4-second spread that made 23.9 look so trustworthy was not a sign of a stable pipeline. Measured by name, the spread is 4.7 seconds. The instrument had simply been timing the steadier half of the cluster.

Two more things turned up in the same hour. The worker's install path never prints the *"k3s started"* line the instrument searches for — the control plane prints it at 12.8 seconds, the worker never does — so that row could never have been filled. And the KubeVirt API that was unavailable on 9 September, and that forced the instrument's no-login design, reports `Available=True` on the rebuilt cluster. The constraint that shaped the instrument did not survive the rebuild. The instrument's virtues did: nothing on the guests, nothing changed by measuring.

What this does not overturn deserves the same precision. The 9 September tables report a worker-join phase, which means the worker's Ready time was captured in those runs — so, as far as the surviving tables show, the chain from 50.0 to 34.5 seconds in Chapter 8 was a both-nodes measurement. But the raw files behind those tables were not kept, so that cannot be re-checked. The 16 September files were kept, in a temporary directory — which is the only reason this chapter could find its own mistake.

What an instrument owes you

The lessons are unglamorous, and they apply to any number a business steers by.

- **Name what you count.** "Two nodes Ready" was satisfied by a node that did not exist. "The control plane and the worker are Ready" could not have been.
- **A missing value is not zero.** A dashboard that renders a gap as a zero will one day report an improvement that never happened.
- **Keep the raw data behind the headline.** The correction was possible only because one run's files survived in a temporary directory.
- **Re-verify the instrument whenever the thing it measures changes.** A new image, a new version, a new cluster — each can quietly change what a count means.
- **Be most suspicious of good news.** A number that improves by thirty percent with no change anyone can point to is not a result. It is a question about the stopwatch.

* * *

MERIDIAN · THE FOLLOWING MONDAY

The wall display had changed. It no longer showed one number. It showed two, each with a node name beside it, and a date underneath.

"Thirty-three point seven," Anita read. "Last week it said twenty-four."

"Last week it was timing half the cluster," Vikram said. "There's a machine baked into our golden image that doesn't exist, and for the first few seconds of every build it counts as a node. Our stopwatch counted it."

"And nobody noticed because the number looked better."

"Better, and steadier. That's what made it convincing."

Anita looked at the display for a moment. "When a number improves and nobody can say why," she said, "it gets audited before it gets celebrated. Put that on the wall too."

"Under the numbers?"

"Above them," she said. "An instrument is a claim too."




Figure 7.1 — Where fifty seconds went, and where the golden image could not help.

The ledger

- **Built:** a read-only, 320-line instrument that reconstructs a cluster build from records the platform already keeps — object timestamps, pod status, the guest's serial console, node transitions, and a host-side probe.
- **First reading (9 September):** 36% of the wall clock passed before the guest kernel started; the "slow k3s startup" theory behind two failed experiments had bet on 15% of it.
- **Enabled:** the two fixes in Chapter 8.
- **Found wrong (17 September):** a node count satisfied by a ghost Node baked into the golden image; a "both nodes" total that treated a missing worker timestamp as zero; a worker milestone the worker never logs.
- **Corrected:** the upgraded platform builds both nodes in 33.7 seconds (median of three; control plane alone, 24.7) — not 23.9, and no measurable change from before the upgrade.
- **Still owed:** both timing scripts still count nodes rather than naming them; the 9 September raw files no longer exist to re-check.

Ask your team

1. When a metric improves, what do we check first — the change, or the instrument?
2. What do our dashboards do with missing data: show a gap, or show a zero?
3. Do we keep the raw data behind our headline numbers long enough for someone to prove them wrong?

Open the repo

- `scripts/phase-timings.sh` — the instrument; its header lists every data source, and `--collect-only` changes nothing.
- `scripts/time-to-ready.sh` — the one-number timer it replaced, and whose "two nodes Ready" condition it shares.
- `docs/fable-cluster-speed-prompt.md` — the brief's mandatory Step 1.
- `docs/sub-60s-cluster-strategy.md` — "Phase-level measurement", with the 9 September budget.
- `03-target-cluster/target-cluster-warm.yaml` — the `warm-ghost-node-cleanup` step, and why it must run first.
- The ghost, live: list the target cluster's nodes the moment the control plane turns `Ready`, and look for `ubuntu-bake-vm-warm`.
- **Without the repo:** Appendix G.2 quotes the `warm-ghost-node-cleanup` step and the ordering comment above it.

Chapter 8 — Two Lines and a CPU Quota

Do not claim anything is faster until you have pasted the timing output that shows it.

— from the brief, `docs/fable-cluster-speed-prompt.md`

MERIDIAN · A TUESDAY IN SEPTEMBER · THE FOURTH-FLOOR BOARDROOM

The slide said `< 40 s`.

It had said so since June: in the platform team's quarterly review, on the internal wiki, and in the one-page summary Anita Rao had sent the board when she asked them to fund the sovereign-cloud pilot instead of another three-year committed-spend agreement. Forty seconds from *I need a cluster* to a cluster. It was the number that made the pilot sound like a product.

Vikram Iyer put a single sheet of paper on the table. It was not a slide. It was the raw output of a script, and its last line read `48.520s`.

"That's the cold build," he said. "Clean host, both nodes ready."

"When?"

"June. Nothing has moved it since."

Anita read the sheet twice. "Then where did forty come from?"

"Forty was the target. Somewhere between the plan and the wiki, somebody wrote the target down in the column for results."

She did not ask who. She asked the more expensive question. "The standby pool hands a team a cluster in two hundred milliseconds. You showed me. Why do I care what the cold build does?"

"Because the pool is one deep." He had rehearsed this part. "The first team that asks gets the standby. The second team that asks in the same minute gets the cold build. And the standby is eight cores and fourteen gigabytes of memory sitting idle, on purpose, all day. Every second I take off the cold build is idle capacity I don't have to buy to hide it."

"So the cold build is an inventory problem."

"It's an inventory problem."

Anita turned the sheet over, as if there might be a better number on the back. "We spent June making the golden image smarter. Twice. What did that buy us?"

"Two things we now know don't work," Vikram said. "And a rule. Nobody changes a line until we can say where the forty-eight seconds go."

* * *

Making the image smarter

Chapter 6 described the golden image: a virtual-machine disk with k3s already installed, baked once and copied for every cluster, so that a new node does not spend its first minutes downloading software. By early summer the default build — control plane first, worker only once the control plane reported ready — took between 90 and 150 seconds. The obvious next move was to put even more into the image.

The theory was attractive and, on paper, sound. The first time k3s starts, it does a great deal of work: it creates its datastore, generates certificate authorities and signs certificates with them, registers its built-in resource types, and installs its system components. All of that is identical on every cluster. Do it once, at bake time, and every cluster afterwards inherits it for free.

Two experiments followed. Both failed, for different reasons, and both failures are worth the space.

Dead end one: `:preinit`

The first attempt preserved k3s's entire state inside the image — the datastore with its resource definitions and access rules, the certificate directory, the air-gapped container images — and added a first-boot step that ran `k3s server --cluster-reset` to give each new cluster a fresh datastore identity. The estimated saving was 30 to 60 seconds.

What actually happened was a cluster stuck at `WaitingForKthreesServer` for more than twenty minutes, while a log inside the control-plane VM repeated one line:

```
CA cert validation failed: Get "https://127.0.0.1:6444/cacerts":
  tls: failed to verify certificate: x509: certificate signed by unknown authority
```

The cause was a disagreement between two systems, each doing exactly what it was built to do. The bake had generated certificate authorities and signed every certificate on disk with them. Then the Cluster API provider for KubeVirt booted the VM and wrote *fresh* per-cluster certificate authorities into it — because managing cluster identity from the management cluster is its job. It overwrote the authority files and left in place every certificate the old authorities had signed. k3s started, presented a certificate signed by the old authority, validated it against the new one, and retried forever.

The repository's post-mortem puts it plainly: *"This failure mode is invisible in design docs that only reason about etcd identity. It only showed up under real end-to-end boot."*

The fix was to delete every certificate that was not an authority before k3s started, and let k3s regenerate them. That worked. It also threw away most of what the image had preserved. The reset still ran, the certificates were still generated at boot, and the one measurement taken came in at about 136 seconds — back inside the 90-to-150-second band the whole exercise had been meant to escape.

The change that did work

The same week, a different kind of change was tried. Cluster API normally holds the worker back until the control plane reports ready, so the two VMs boot one after the other. Three settings removed that wait: a join token created in advance, a static bootstrap configuration for the worker, and an annotation telling Cluster API to skip its readiness gate. The worker's agent simply retries the control plane's virtual IP until it answers.

Three runs: 46.6, 49.7, and 50.0 seconds. It was the largest single improvement of the summer, and it was not an image change. Nobody drew that conclusion yet.

Dead end two: `:warm`

The second image experiment attacked the certificate conflict directly. If the provider insists on writing certificate authorities into the VM, make sure it writes the *same* ones the image was baked with. The bake used a fixed set of authorities and a fixed join token; before each deploy, a script pre-loaded those same authorities into the management cluster's secrets; the provider found them there, adopted them, and wrote back identical files. No conflict, no purge, no reset.

It worked exactly as designed. k3s kept the pre-placed authority byte for byte. And the first few runs came in at around 35 seconds.

That was the dangerous moment. Thirty-five seconds is under forty. It is the number on the slide.

The team ran it again, head to head, on the same host, against the plain parallel build it was supposed to beat. The warm image landed between 46 and 56 seconds. The plain build had landed at 46.6, 49.7, and 50.0. A clean cold run of the warm path, measured a few days later, took 48.520 seconds. The two were statistically indistinguishable. The 35-second runs had happened on a quiet host, and this host is rarely quiet.

In hindsight the reason is simple. Everything the warm image pre-computed — certificates, resource definitions, system components — sat *off the critical path*. A node reports ready well before k3s has finished installing its optional components, which apply in the background anyway. The image had made fast a set of steps nobody was waiting for.

The warm image became the default regardless. It was no slower, it was more reliable than the preinit path, and it was what the parallel build now ran on. What went onto the wiki next to it was a line describing it as the path that *"targets time-to-ready < 40s."*

Two capable attempts, each built on a plausible theory, each costing a full bake cycle to disprove. Neither theory had been measured before it was built.

This is one of the most common ways engineering budgets disappear, and it rarely looks like failure while it is happening. The team optimized the part of the system it could reason about most easily — k3s's own startup. Nobody had yet asked where the fifty seconds actually went.

* * *

The brief

Before anyone changed a line in September, someone wrote a brief. It was addressed to an AI coding agent — the same kind of worker Part IV of this book is about — and it lives in the repository as `docs/fable-cluster-speed-prompt.md`. It is worth reading in full, because nearly everything that went right afterwards is already in it.

It named **one number** and defined it precisely: the wall time from `kubectl apply` of the cluster manifest until both nodes report `Ready`. *"It is the only metric that counts."*

It listed what had already failed, under a heading that leaves no room for interpretation: **"Do NOT retry these — they are settled negative results."** Baking more state into the image. The warm image. Parallelizing the worker's boot, which had already been done.

It called out its own documentation. The wiki's under-forty claim was *"a stale aspiration, never achieved. Where CLAUDE.md and the table above disagree, the table wins."*

It made instrumentation mandatory before any optimization: *"Every optimisation attempt so far has been evaluated against a single aggregate number, which is why two dead ends took a full bake cycle each to disprove."* It even distrusted its own starting estimate — a rough split of seventeen seconds for the VMs to boot, seventeen for the control plane to start, eighteen for the worker to join — and asked for it to be checked: *"If it is wrong, say so; that finding alone is worth more than a speculative patch."*

It set a measurement protocol and called it non-negotiable. Tear the old cluster down and confirm it is gone. Record host load before every run. Three runs minimum; report the median *and the full spread*. Re-measure the baseline in the same session as the candidate, never against a number from another day. Change one variable at a time. And it said why: the host was shared with two Kind clusters and three unrelated containers, and *"an earlier run of this work recorded 35s times that turned out to be quiet-host luck rather than a real improvement."*

It defined success — a median under 40 seconds over three clean runs, with the cluster verification passing afterwards — and then it defined an acceptable failure: *"'No lever found' is an acceptable and valuable outcome."*

Its last line is this chapter's epigraph.

Three things in that brief cost nothing and are almost never written down. A single metric, defined precisely enough that two people measuring it get the same answer. A list of what has already been proven not to work. And explicit permission to come back empty-handed. The third matters most. A team that cannot report *no lever found* will, eventually, report a lever.

Where the fifty seconds went

Chapter 7 tells the story of the instrument itself: a read-only script, `scripts/phase-timings.sh`, that reconstructs a cluster's boot from timestamps the platform already records — when each Kubernetes object was created, when each VM changed phase, when the guest kernel printed its

first line, when k3s announced it had started — without logging in to anything.

Its first job was the baseline. Re-measured in the same session, as the brief required, it was not 48.5 seconds but 50.0: runs of 52.7, 50.0, and 46.6. Here is where those seconds went, median of three clean runs:

Phase	Seconds	Share
<code>kubect1 apply</code> → VM instance created (Cluster API controllers)	6.6	14%
Launcher pod created → QEMU running	10.0	21%
Firmware, kernel, initramfs	3.7	8%
systemd and cloud-init → k3s start	5.5	12%
Control-plane k3s start → first <code>200</code> from <code>/readyz</code>	7.0	15%
Worker joins, after the API is already answering	11.8	25%
Median time-to-ready	50.0	52.7 / 50.0 / 46.6

Two things in that table changed the plan immediately.

The first was how wrong the estimate had been. The rough split gave the control plane's k3s startup seventeen seconds. The instrument measured seven. The theory both image experiments had been built on — *the slow part is k3s starting up* — had bet on a row worth fifteen percent of the total.

The second was the top two rows. **Thirty-six percent of the wall clock — 16.6 seconds — elapsed before the guest kernel printed its first character.** No change to the image could touch any of it, because none of it happens inside the image.

Two rows looked suspicious. Ten seconds between a launcher pod being created and QEMU running, eight of which the first version of the table labelled "*containerDisk unpack.*" And nearly twelve seconds for a worker to join a control plane that was already answering. The team went after the worker first. It took two wrong turns.

Dead end three: the retry that was doing work

The worker's k3s agent does not wait politely for the control plane. It starts, tries to reach the API, fails, waits, and tries again. The phase data showed the agent starting at 26.7 seconds, the API first answering at 34.8, and the agent's next attempt landing at 36.7 — because it retries on a cycle of roughly five seconds. That looked like 1.9 seconds of pure waste: the API was ready, and the agent was asleep.

The fix looked obvious. Have the worker wait until the API's unauthenticated `/readyz` endpoint answers, *then* start the agent, so that its first attempt succeeds.

It made the cluster slower. The median went to **53.1 seconds — a 3.1-second regression.** In the repository's words: "*The premise was wrong.*"

The worker's own log explained why. During those "wasted" retry cycles, the agent was starting its container runtime and preparing the node — real work, overlapping with the control plane's boot. The retry window was not idle; it was parallel. Holding the agent back until the API answered turned two things happening at once into two things happening in sequence. The change was reverted.

| *From the outside, a retry loop looks like waste. Often it is the only part of the system working in parallel.*

Dead end four: no route to host

The worker's log held something else: a connection error, three times, each attempt taking about 3.1 seconds, before the agent fell back to the control plane's stable virtual IP, `172.18.255.215`, and succeeded.

```
connect: no route to host
```

The explanation is a genuinely strange property of running VMs inside Kubernetes. When KubeVirt connects a VM to the network in bridge mode, the guest receives its pod's IP address *and that address's/24 subnet*. So the guest believes every other pod in the range is on the same local wire, and tries to find them the way machines on a shared wire do: by broadcasting an ARP request. But pods on this cluster's network are not on a shared wire. Traffic between them is routed. The broadcast finds nothing, and the connection fails with *no route to host*.

Why was the worker trying to reach a pod address at all? Because when the agent first contacts the control plane, the control plane hands back a list of addresses where its API can be reached, and on that list was the control-plane VM's own pod IP. There was also a stale `10.0.2.2:6443` — the address of QEMU's internal gateway — left over from bake time in the warm image's datastore.

Routing the pod range through the gateway fixed the errors completely. Six logged errors per run became zero, and the agent made one clean connection.

The timing benefit was nothing. **The median was 54.3 seconds.**

The log showed why. Between the line where k3s announced that it was starting the kubelet — `Running kubelet --address=0.0.0.0` — and the kubelet's first output, there was a **ten-second silence**. The failed dials had been happening *inside* that silence, not adding to it. Remove them, and the silence was still ten seconds long.

It was a real bug, and the fix was kept for correctness. It was not the slow thing.

| *A real bug is not necessarily the expensive one. The instrument is what tells you which is which.*

Found on the way

The same investigation turned up a third bug, one that cost nothing today and could have cost a great deal later. The worker's bootstrap asked for k3s's *agent* mode through an environment variable, which the installer baked into the image silently ignored: it read only positional

arguments. So every worker started the k3s *server*, which failed about 0.8 seconds later and aborted the install script before it wrote the file that marks a successful bootstrap. The cluster came up anyway — but only because the provider had been configured not to check for that file. Moving one word to the right place on the command line fixed it.

Ten point zero zero zero

The ten-second silence had one more trick. Across runs where the dials behaved completely differently, it measured 10.019 seconds and 10.017 seconds. A gap that stable, whatever the network did, looked like a fixed timer somewhere inside k3s that had nothing to do with the dials.

Turning on the agent's debug logging named it exactly:

```
Dial error from server 10.244.0.186:6443@UNCHECKED after 10.000162522s: i/o timeout
```

The kubelet's first log line landed seven milliseconds after that timeout expired.

It was not a timer unrelated to the dials. It *was* the dial: a ten-second TCP connection timeout, reached by two different routes. Before the routing fix, the agent spent the window failing repeatedly. After it, a single connection hung silently until the timeout ran out. Either way, the kubelet did not start until the agent's connection attempt gave up.

And the address it was trying to reach, `10.244.0.186`, was the control-plane VM's pod IP. A two-line test made the problem undeniable:

```
node -> CP VM pod IP:6443    = 200, connect 0.0003s
pod  -> CP VM pod IP:6443    = timeout
```

A KubeVirt VM's pod address is reachable from the node it runs on, but not from other pods. And the worker VM, as far as the network is concerned, is just another pod. It could never reach that address. It waited out the full ten seconds, then fell back to the virtual IP it had been configured with all along.

The fix is one line, added to the commands the control plane runs before k3s starts:

```
- "echo 'advertise-address: 172.18.255.215' >> /etc/rancher/k3s/config.yaml"
```

It tells the control plane to advertise the virtual IP — an address every VM can already reach — instead of its own pod address.

	Before	After
Median time-to-ready	50.0 s	42.7 s
Runs	52.7 / 50.0 / 46.6	42.7 / 42.8 / 41.2
Run-to-run spread	6.1 s	1.6 s
Dial timeouts per run	4-6	0
Worker join after API ready	11.8 s	5.1 s

Before trusting it, the team checked that nothing had broken. Inside the target cluster, the built-in `kubernetes` service now pointed at the virtual IP, and a pod could still reach `kubernetes.default` with a `200`.

The spread row matters as much as the median. Run-to-run variation fell from 6.1 seconds to 1.6, because the ten-second timeout had been its largest source. What had looked like a noisy host was, mostly, the bug.

The noise was the bug.

One percent of a core

That left the other suspicious row: ten seconds from a VM's launcher pod being created to QEMU actually running, eight of them attributed in the first table to *"containerDisk unpack."*

That label was a guess, and it was wrong.

Every VM in KubeVirt runs inside a launcher pod, and that pod carries small helper containers alongside the one that runs the VM. The instrument showed when each of them started:

Container	Job	CPU limit	Started at
<code>guest-console-log</code>	copies the VM's serial console to the pod log	15m	+3.0 s
<code>volumesystemdisk</code>	holds the VM's disk image for the launcher	10m	+9.0 s
<code>compute</code>	runs QEMU	none	+9.0 s

The limits are the clue. In Kubernetes CPU units, `10m` means ten millicores: one percent of one core. Linux enforces that as a quota — roughly one millisecond of CPU time in every hundred. Even a program that does almost nothing (this container's process literally runs with a `--no-op` flag) needs CPU time to load its libraries and page itself into memory. At one millisecond per hundred, that takes dozens of scheduling periods. The disk container could not start until it had them, and the VM could not start until the disk container had.

Nobody guessed that. It was found by elimination, and every rival explanation was measured and discarded first:

Hypothesis	Test	Result
The 1.17 GB image is slow to unpack	Mount the same image as an image volume	Container started in +0.0 s
Two VMs starting at once contend	Start one VM on its own	9.0 s — <i>slower</i> than 8.0 s
The node's container runtime is slow	Start a trivial two-container pod	1.0 s in total
Mounting image volumes is slow	As above	Fast

What gave it away was that the eight seconds was **deterministic**: 8.0 seconds on every VM, on every run, whether the host was busy or not. Load produces variation. A quota produces a constant.

Nobody on the team had chosen ten millicores. It came with the platform. Changing it meant raising a ceiling in KubeVirt's own configuration — `supportContainerResources` — from one percent of a core to a full core:

```
{
  "type": "container-disk",
  "resources": {
    "requests": {
      "cpu": "100m",
      "memory": "40M"
    },
    "limits": {
      "cpu": "1",
      "memory": "100M"
    }
  }
}
```

The limit is the part that removed the eight seconds, and a limit is a ceiling, not a reservation. These containers do their work at startup and then sit idle, so raising the ceiling costs nothing once the VM is running. (The same change sets a modest request of a tenth of a core per helper container.)

	Before	After
containerDisk phase	8.0 s	1.0 s
Pod → QEMU running	10.0 s	4.0 s
Median time-to-ready	42.7 s	34.5 s
Runs	42.7 / 42.8 / 41.2	34.5 / 34.2 / 34.7
Run-to-run spread	1.6 s	0.5 s

There is one catch, and it belongs in the ledger. This setting lives in the running cluster's KubeVirt configuration, not in a manifest in Git. Rebuild the management cluster and it is gone. The fix was wrapped in a script, `scripts/configure-kubevirt-perf.sh`, and wired into the management cluster's setup so that a rebuild would reapply it. A week later, after a rebuild from nothing, it was found missing anyway and had to be applied by hand. That story is Chapter 19.

Where it landed

baseline	50.0 s	(52.7 / 50.0 / 46.6, spread 6.1 s)
+ advertise-address	42.7 s	(42.7 / 42.8 / 41.2, spread 1.6 s)
+ supportContainerRes	34.5 s	(34.5 / 34.2 / 34.7, spread 0.5 s)
	-15.5 s	(-31%); the gate was < 40 s

Fifteen and a half seconds off every cold build, and a pipeline that now finishes within half a second of itself.

Neither fix touched the image. The repository's own summary is the sentence this chapter is named for: *"Every attempt to make the image smarter (`:preinit`, `:warm`) failed, while the two things that worked were a bad advertised address and a CPU quota."* Add the parallel boot from June, and the pattern holds across the whole summer: every change that moved the number was about how the pieces were orchestrated and configured, not about what was baked into them.

There is a rhyme between June and September worth pausing on. In June a stopwatch said 35 seconds, and it was luck. In September it said 34.5, and it was true. The difference was not the number. It was a spread of half a second across three runs, and two causes you can point to in a log.

The instrument also said what is left, by measured size: about six seconds of Cluster API controllers doing their work before any VM exists, which no change inside the guest can reach; seven seconds for the control plane's k3s to answer; five and a half seconds of systemd and cloud-init. That last phase is the one a leaner, systemd-free image would attack — and a perfect result there is worth about three seconds. It is the fourth lever, not the first.

The caveat that limits all of this

Every number above comes with a condition, and the repository records it under a heading that does not soften it: *"Measurement caveat that limits all of the above."*

At the start of this work, run-to-run spread on this host was about six seconds. With three runs per candidate, any change smaller than roughly five seconds disappears into that noise. The two fixes that worked were each worth seven to eight seconds — just large enough to see. A three-second improvement would have been invisible, and a three-second regression could have shipped unnoticed. Anything finer needs seven or more runs, a quiet host, or both.

The fixes changed that condition too. With the spread down to half a second, a one-second regression is now visible in three runs. Removing the bug did not only make the cluster faster. It made the cluster *measurable*.

Postscript: the ten seconds that weren't

On 16 September, a week later, everything underneath the platform was upgraded at once — Kubernetes on the two host clusters from 1.35 to 1.37, k3s inside the VMs from 1.31 to 1.37 — and both host clusters were destroyed and rebuilt from nothing. Afterwards, the same instrument, three clean runs, reported **23.9, 24.1, and 23.7 seconds**. Another ten seconds, and steadier.

It was wrong.

As Chapter 7 tells, the instrument had been stopping its clock when the control plane turned `Ready` — fooled by a ghost node carried inside the golden image — and reporting the control plane's time as the total for both nodes. Measured again on 17 September, naming each node instead of counting them, the upgraded platform builds both nodes in **33.7 seconds** (30.4, 35.1, 33.7). Within the noise of three runs, that is the same as the 34.5 before the upgrade.

This chapter's rule — a faster number is a claim until the phase table says why — turned out to apply to the phase table too. This time the explanation for the good news was the stopwatch.

* * *

MERIDIAN · THE FOLLOWING TUESDAY

Vikram put another sheet on the table. Three lines this time: `34.5`, `34.2`, `34.7`.

Anita looked at the spread before she looked at the median. "Half a second."

"Half a second."

"What did it cost?"

"A script to measure it. Two lines of configuration."

"Hardware?"

"None. The second line raises a ceiling. It doesn't reserve anything."

"And the standby?"

"Still one deep. But the second team in the minute now waits thirty-four seconds instead of fifty."

"And June?"

"June is written down, with its numbers. Nobody here pays for it twice."

She opened the board summary on her laptop, found the line that said `< 40 s`, and deleted it. Then she typed the replacement herself: *34.5 s median, three runs, spread 0.5 s — measured 9 September*.

"Put a date on every number," she said. "A number without a date and a spread is a slogan."




Figure 8.1 — Two one-line changes: the median falls from 50.0 s to 34.5 s, and the spread from 6.1 s to 0.5 s.

The ledger

- **Spent:** two golden-image experiments, each a full bake cycle to disprove; one instrumentation script; two further hypotheses — gating the worker, rerouting the pod network — that did not survive measurement.
- **Changed:** one line in the control plane's startup commands; one CPU ceiling on KubeVirt's helper containers.
- **Hardware added:** none.
- **Saved:** 15.5 seconds on every cold cluster build — 50.0 s to 34.5 s, -31%. Run-to-run spread from 6.1 s to 0.5 s.
- **Proved:** four dead ends, recorded with their numbers so that nobody pays for them twice.
- **Still owed:** a setting that lives in cluster state rather than in Git, and two timing scripts that can stop the clock on a ghost node (Chapter 7).

Ask your team

1. What is our one number, exactly how is it measured, and when was it last measured on this hardware? Ask to see the timing output, not the wiki page.
2. What have we already proven does not work, and where is it written down? If the answer is "people remember," you will pay for it again.
3. Which of our performance-critical settings did nobody here choose? Ten millicores came with the platform. Find yours.

Open the repo

- `make phase-timings ARGS="--runs 3"` — runs `scripts/phase-timings.sh` and prints the per-phase table, using this chapter's three-run protocol.
- `docs/fable-cluster-speed-prompt.md` — the brief, verbatim.
- `docs/sub-60s-cluster-strategy.md` — "What failed, and how it was fixed" for `:preinit`; "Phase-level measurement" for the budget, both dead ends, and both fixes.
- `03-target-cluster/target-cluster-warm.yaml` — the `advertise-address` line in the control plane's `preK3sCommands`.
- `make kubevirt-perf` — applies the CPU ceiling; the header of `scripts/configure-kubevirt-perf.sh` explains it.
- `git show c7ad792` — the commit that landed both fixes, with the measurements in its message.
- **Without the repo:** Appendix G.2 quotes the `advertise-address` line with its measurement comment; G.3 quotes the `supportContainerResources` patch.

Chapter 9 — Warmth Is Mandatory

Warmth is mandatory. No persona may enter the fleet without a warm-path guarantee.

— Architecture Constraint 5, `Sovereign_Cloud_Agentic_Strategy.md`

MERIDIAN · A THURSDAY IN SEPTEMBER · THE CTO'S OFFICE

The prospect had asked the AI console a simple question, and for thirty seconds the answer did not come. Then the request timed out, on the big screen, in front of their head of infrastructure.

Vikram asked again. The second answer arrived in four seconds and was correct.

"Why was the first one slow and the second one fast?" Anita asked afterwards.

"The first question had to wake the model up. It loads into memory on first use. After that it stays loaded, and everything is quick."

"Like the clusters," she said. "Build one from nothing, it takes half a minute. Hand over one that's already running, it's instant."

"Same idea."

"Then it has the same economics." She pulled a pad toward her. "Keeping things warm costs hardware that sits idle. Being cold costs moments like that one. I want both prices, for both things, on one page. And I want to know who is keeping them warm right now."

* * *

BUILD LOG · JUNE–SEPTEMBER 2026

Two kinds of cold

This platform has two things that are slow to start and fast once running, and the whole of this chapter is about the difference.

The first is a **cluster**. Chapter 7 measured a cold build — from `kubectl apply` to both nodes `Ready` — at 33.7 seconds after the September upgrade, and Chapter 8 tells how it came down from fifty. Every one of those seconds is intrinsic to building a cluster from nothing: controllers creating virtual machines, the machines booting, k3s starting, a worker joining. Chapter 8's lesson was that a smarter golden image cannot remove them. Only one thing can: doing the work *before anyone asks*.


The second is a **model**. An agent's first request loads its language model into memory. The strategy document wrote the cost into its fifth architecture constraint, as measured when the platform's model ran on the CPU: *"Cold-start for llama3.2 on CPU is ~28–30s; unwarmed agents time out the first real request."*

Different machinery, same trade: capacity held in reserve against time lost to a cold start.

A standby, not a pool

The design for warm clusters is in `docs/warm-pool-strategy.md`, and it begins by rejecting the obvious answer. A classic pool keeps several pre-built machines and hands them out. This platform runs exactly one target cluster at a time, so the document asks for something smaller: a *size-1 pre-deployed standby*. Keep one cluster built and idle. "Deploy" means *claim* it. "Delete" means tear it down and build the next standby in the background, off the user's clock.

The lifecycle is a label on the cluster object, `pool.local/state`:



One constraint makes it tractable, and it is easy to miss. A KubeVirt VM's pod address is assigned when its launcher pod starts, and the document records that a changing address was what broke the datastore in the warm-image experiment. *"A pooled VM that stays alive keeps its pod IP"* — so its datastore identity and its certificates stay valid. That is also why the document rejects the other obvious answer, snapshot and restore: a restore *"re-provisions the PVC and reboots, so you pay the 17s boot again and get a new IP."*

A second level is described and deliberately deferred: a pool of *paused* VMs, frozen by KubeVirt with their memory and addresses intact, resumable in under a second. It would allow more than one warm cluster, at the price of a real controller and, in the document's words, *"ongoing maintenance."*

One hundred ninety-four milliseconds

BUILD LOG · JUNE 2026

The standby was built on 27 June as a background loop inside the UI's Go backend. Every few seconds it looks for a target cluster. If none exists, it builds one through the same Cluster API path as any other deploy and labels it `WARM` once it is ready. When a user clicks Deploy and a warm standby exists, the handler relabels it `CLAIMED`, refreshes the credentials the agents use to reach it, and streams a short progress report.

Measured end to end, with a warm standby in place: **194 to 195 milliseconds** from the click to done, with both nodes already `Ready`. The background rebuild after a delete took about 60 seconds, and the cycle — claim, delete, rebuild, claim — repeated cleanly. The acceptance gate had been under forty seconds; the claim beat it by a factor of about two hundred.

One bug is worth recording, because it is Chapter 7's lesson in miniature. The first version reset the progress log from a background thread while the user's browser was already subscribing to it. The browser received the *previous* build's log and an immediate "done": a fast, confident, false success. Only an end-to-end test caught it.

Then came the price. The standby uses the full profile: four cores and 8 GiB of memory for the control plane, four cores and 6 GiB for the worker. That is **eight cores and fourteen gigabytes held idle, all day, on purpose**, on a host that had about 21 GiB free. The document says it plainly: *"Idle resource cost. One cluster always running."*

Within a week, that cost produced a decision. The pool had started out enabled, which meant a standby cluster was built every time the UI started and rebuilt after every cleanup — whether or not anyone intended to deploy. That surprised its own user. On 1 July it became opt-in, and the start script now says so: *"Opt-in: export POOL_ENABLED=true before running this script to have the backend auto-build/rebuild a standby cluster in the background."* Warmth that nobody asked for is not a feature. It is a machine working when no one is watching.

The standby today

BUILD LOG · 17 SEPTEMBER 2026

Checked against the running platform on 17 September, the fastest path in the platform turns out to be the least exercised one.

The UI backend reports no standby: `{"state": "none"}`. The pool is opt-in, and the script that starts the UI leaves it off by default. It also lives only as long as that backend process runs, so the command-line deploy path has never benefited from it.

Its standby manifest still points at the image tagged `ubuntu-noble-k3s:latest`, and that image was built on **3 March**. It was not rebuilt during the September upgrade in Chapter 19. Turning the pool on today would build its standby from a six-month-old image, under a manifest that now declares a much newer k3s. Nobody has tried it.

Its readiness check counts *two Ready nodes*, the same pattern that fooled the stopwatch in Chapter 7. A standby built from the warm image would be declared ready while its worker was still booting, because that image carries the ghost node. Whether the March image carries anything similar has not been checked.

And the 194-millisecond claim has not been measured since June.

None of this means the design is wrong. It means a warm path is only as good as the last time someone used it. The image aged, the platform moved underneath it, and a number from June stayed on the page.

Warming the model

BUILD LOG · APRIL–SEPTEMBER 2026

The model's warmth has a history of its own, told in three tools.

The first was a script, `06-symposium/ollama-warm.sh` : send a one-token request and ask the model server to keep the model loaded for twenty-four hours. The second was the Phase 0 deliverable from Chapter 16: a heartbeat schedule that pings `cluster2-agent` every four minutes. Its comment explains why: *"Cold-start on CPU is ~28-30s; this schedule removes that latency from the first real KubeUI chat request."* Chapter 14 found that schedule had fired twice and stopped. The third tool is the one actually doing the job: the host's Ollama service runs with `OLLAMA_KEEP_ALIVE=-1` , which keeps a loaded model loaded indefinitely. That setting lives in the host's service configuration, not in the repository.

The GPU changed the price

BUILD LOG · 17 SEPTEMBER 2026

The constraint's thirty seconds was a CPU number. The host now has an NVIDIA RTX 4050 laptop GPU with 6 GB of memory, and Ollama places each of the fleet's models entirely on it. So the cold start was measured again on 17 September, using the load time Ollama reports for each request — first with a model unloaded, then with it loaded:

Model	Cold load	Warm load
qwen2.5:7b	2.57 s, and 2.38 s on a second cold start	0.14 s, 0.17 s
llama3.2	2.33 s	0.17 s

About two and a half seconds cold, under two-tenths of a second warm. The difference is still roughly fifteen-fold, and still visible in an interactive chat. But a two-and-a-half-second load no longer times out a request by itself. The hardware changed under the rule, and the rule's number went stale.

The same measurement showed something the rule never anticipated. **The GPU holds one of these models at a time.** Every cold load evicted the other model. `OLLAMA_KEEP_ALIVE=-1` does not mean *forever*. It means *until something else needs the room*.

That matters because the fleet uses both. `cluster2-agent` and `mesh-sre-agent` run on `qwen2.5:7b` ; `target-cluster-agent` runs on `llama3.2` . And it had already happened that morning, unnoticed. At 00:55 UTC, `qwen2.5:7b` was the loaded model. At 01:06, the two-agent pipeline lab from Chapter 16 ran its reviewer on `llama3.2` . By the time anyone checked, `llama3.2` was loaded and `qwen2.5:7b` — the model two of the three agents depend on — was not. A test run had quietly cooled the production agents. (It was restored for this chapter.)

Warmth on shared hardware is not a property of a model. It is a queue.

Two inventories, both one deep

Put the two kinds of warmth side by side and they turn out to have the same shape:

	A cluster	A model
Cold	33.7 s, both nodes	2.3–2.6 s on the GPU (28–30 s on CPU)
Warm	467 ms median (18 Sept, 3 runs; was 194 ms in June)	0.14–0.17 s
What warmth holds	8 cores and 14 GiB of memory, idle	most of a 6 GB GPU
How deep	one standby	one model at a time
What keeps it warm today	a background loop that is off by default	a host setting outside the repository

For a CxO, warmth is inventory. It trades capacity you pay for up front against time you lose later, and it has the same three failure modes as any inventory.

It can be the wrong size. One standby covers the first request in a minute, not the second. One model slot covers one model, not a fleet that uses two.

It can spoil. A standby built from a March image, and a heartbeat that stopped in September, look like warmth on paper and are not.

It can be taken by someone else. A test run that loads a different model empties the shelf the production agents were counting on.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Vikram's page had two columns, one for clusters and one for models, and a price at the bottom of each.

Anita read it twice. "The model costs two and a half seconds cold now, not thirty."

"On this GPU. And only if it's the model that was already loaded. The GPU holds one at a time, and our agents use two."

"Then we use one," she said. "The one that uses its tools properly. One model for the whole fleet, until the hardware says otherwise." She moved to the other column. "And the cluster standby?"

"Built from a March image, and switched off."

"Then it isn't warm. It's a cold path with a design document." She wrote across the bottom of the page. "Rebuild the image. Turn it on. And every week, someone claims it, deletes it, and times the rebuild — naming both nodes."

"Every week?"

"Anything we keep warm, we exercise," Anita said. "Otherwise we are paying for idle hardware to hold a promise nobody has checked."

Turned on, and timed

BUILD LOG · 18 SEPTEMBER 2026

The mandate was carried out the next day, and it found the same ghost that Chapter 7 found in the timing scripts, waiting in the same shape.

`targetNodesReady`, the function both the on-demand deploy path and the pool controller use to decide "are both nodes up," counted Ready nodes without checking their names. A standby built from the `:warm` image would have been declared ready on the control plane plus the leftover bake-VM ghost node, before the real worker ever joined — exactly the bug that made Chapter 7's stopwatch stop early, now confirmed live in `ui/backend/handlers/pool.go`, not just in the standalone scripts. Fixed the same way Chapter 7 recommends: count only nodes named `<cluster>-cp-*` or `<cluster>-workers-*`.

With that fixed, `POOL_STANDBY_MANIFEST` moved from `target-cluster-parallel.yaml` (the March image, k3s v1.31.4) to the warm manifest — and the pool's build function gained the one step it was missing to use it safely: seeding the fixed CA and token secrets before applying, the same order the on-demand deploy path already used, so KThrees adopts the baked material instead of falling back to a slower path.

Then the cycle Anita asked for — claim, delete, rebuild, claim — ran three times, timed from the `POST /api/v1/cluster/deploy` request to the stream's closing `done` event:

Run	Time
1	1.059 s
2	0.467 s
3	0.465 s

Median: 467 milliseconds. Not 194. The first run is reported rather than discarded — an unexplained outlier, kept on the same principle this book applies to every other number in it. Runs 2 and 3 agree closely, and the time is not free: of the roughly 465 milliseconds, about 265 is `clusterctl get kubeconfig` alone; relabeling the cluster and syncing the in-cluster kubeconfig Secret split most of the rest.

Two and a half times the June figure, and still far faster than building from nothing: the claim now beats the 33.7-second cold build by a factor of about 72 — not the factor of two hundred the June number implied, but still the same conclusion Anita drew from the wrong number: handing over something already running beats building it, by nearly two orders of magnitude. The design held. The number on the page had simply gone stale, in exactly the way this chapter's own middle section warned that warmth does.

The ledger

- **Built (June):** a one-deep cluster standby, claimed in 194–195 ms instead of built from scratch, rebuilt in about 60 seconds in the background.
- **Cost of that warmth:** one full cluster idle — eight cores and 14 GiB of memory.
- **Today (17 September):** no standby running; the pool is opt-in and off by default; its manifest points at an image built on 3 March; the claim time has not been re-measured since June.
- **Fixed and re-measured (18 September):** the pool's node-readiness check now counts by name, not by count (the same ghost-node bug Chapter 7 found, confirmed live in `pool.go`); its standby manifest switched to the warm image, with CA/token seeding added to the build path. Three claim-delete-rebuild-claim cycles: 1.059 s, 0.467 s, 0.465 s — **median 467 ms**, not 194. Still about 72x faster than a cold build. The pool remains opt-in — this fixes what happens when it's on, not the default.
- **Model warmth, measured on the GPU (17 September):** 2.3–2.6 s cold, 0.14–0.17 s warm — against the 28–30 s CPU figure in the rule.
- **Found:** the GPU holds one model at a time, and a lab run evicted the model two of the three agents use.
- **Held by:** a host setting outside the repository, not the heartbeat the rule names, which stopped after two runs.

Ask your team

1. What do we keep warm, what does that capacity cost us, and what would being cold cost instead?
2. When did we last actually use our fast path — not measure it once, but use it?
3. What shared resource do our warm things compete for, and whose work gets evicted first?

Open the repo

- `docs/warm-pool-strategy.md` — the three design levels, and the stable-IP constraint.
- `ui/backend/handlers/pool.go` — the standby controller; `POOL_ENABLED`, `POOL_STANDBY_MANIFEST`, and `GET /api/v1/cluster/pool-status`.
- `docs/superpowers/specs/2026-06-26-warm-pool-standby-design.md` — the acceptance gate.
- `06-symposium/ollama-warm.sh` and `06-symposium/schedules/ollama-warm.yaml` — the first two ways the model was kept warm.
- `curl -s localhost:11434/api/ps` on the host — which model is warm right now; the `load_duration` field of any `/api/generate` response shows whether a request paid for a cold start.
- `POOL_ENABLED=true POOL_STANDBY_MANIFEST=03-target-cluster/target-cluster-warm.yaml` — the corrected pool config (18 September); `git log --oneline -- ui/backend/handlers/pool.go run-ui.sh` for the fix commits.

PART III

The Fabric

The mesh. The confession, then the fix.

Chapter 10 — Hardcoded IPs and Other Confessions

`05-istio/` is left untouched. It stays as the "before" picture.

— `07-istio-advanced/README.md`

MERIDIAN · A TUESDAY IN SEPTEMBER · THE EXECUTIVE BRIEFING CENTRE

The prospect was a regional bank, and their architect asked the question every bank asks: if one site goes down, what happens to the other?

Vikram had the answer on the wall screen. Two clusters, drawn as boxes. Four live probes between them, all green. "And the traffic between the sites is encrypted?" the architect asked. "It all goes through the service mesh," Vikram said, and pressed a button that switched off a service in the first cluster. One probe went red. A moment later an amber badge lit across the top of the diagram: ⚡ **FAILOVER ACTIVE.**

The architect nodded and wrote something down.

Afterwards, Anita stayed in the room. "Which service took over?"

"None of them." Vikram brought the code up on his laptop. "The badge comes on when a service here is down *and* a call to a different service over there still works. The two aren't connected."

"So the screen showed two true facts, and a word that joined them."

"Yes."

"And the word is what the bank wrote down." She looked at the diagram a while longer. "How much of the rest of this picture is like that?"

"I don't know yet."

"Find out. Every box, every arrow, every word on it. I don't want it fixed this week — I want a list. Some of it we will keep; a before picture is worth having. But I want to know which parts are the before picture, and which parts are just wrong."

* * *

BUILD LOG · 12 MARCH 2026

Eight steps, one commit

Shortly after midnight on 12 March, a single commit added the platform's whole cross-cluster story: an eight-step interactive demo script, a mind-map diagram of the traffic path, a "Live Demo" tab in the web console with real probes and failure controls, and six more console panels for the service mesh. Thirty-six files; 5,362 lines added.

The problem it took on was genuinely hard. `cluster1` is a Kind cluster on the laptop. The target cluster is k3s running inside two KubeVirt virtual machines, which run inside pods of a *second* Kind cluster, `cluster2`. For a program in `cluster1` to call a program in the target cluster, a request has to leave one cluster, find the right virtual machine inside another, and get in.

And there is no shared address plan to lean on. `cluster1` and `cluster2` both number their pods from the same range, `10.244.0.0/24` — checked on both clusters on 17 September. The target cluster's control-plane machine holds the address `10.244.0.115`. Typed into `cluster1`, that address names a slot in `cluster1`'s *own* pod network, not a machine in the other cluster.

So the demo built a path by hand, and prints it at the end:

```
sleep pod → ztunnel → ServiceEntry → MetalLB 172.18.255.216
→ cluster2 proxy Svc → virt-launcher pod:30080
→ masquerade NAT → k3s VM → nginx
```

In plain language, three pieces of plumbing:

- **A door on every machine.** In the target cluster, `nginx` is published on port 30080 of every node — a *NodePort*.
- **A fixed address on the shared network.** In `cluster2`, a Service is given the address `172.18.255.216`, typed into an annotation for MetalLB to honour. Behind it, instead of letting Kubernetes find the backends, the script writes the destination in itself: the current address of the virtual machine's pod, port 30080.
- **An entry in the mesh's address book.** In `cluster1`, an Istio *ServiceEntry* tells the mesh that the name `nginx.target-cluster.global` lives at `172.18.255.216:8000`. A matching entry in the target cluster points back at `172.18.255.200`, `cluster1`'s `httpbin`.

The eight steps then tell a story with the shape of every good demo. Show the two clusters. Deploy the services. Prove each works locally. Show that the other cluster's service cannot be reached by name — "*✗ Connection failed — service not discoverable.*" Create the *ServiceEntries*. And then, in the script's own words: "*Now the magic: cluster1 calls target-cluster and vice versa.*" Step 7 adds a traffic policy, a connection pool and a circuit breaker. Step 8 shows the logs.

Fail, fix, succeed. It ran, and because it ran, nobody questioned it.

It was also, in six distinct places, not what it said it was. In August, when the platform built the proper answer, it kept this directory untouched, on purpose, as the "before" picture. On 17 September the path was re-plumbed by hand on the rebuilt clusters, using the script's own YAML, to see which of its claims still held; everything created was measured and then removed. What follows is the list Anita asked for.

The diagram was wrong before it was committed

The mind map is this chapter's figure, and it is a good diagram: clear boxes, the forward path in one colour, the reverse path in another, and a panel of key concepts. One of those concepts reads:

KubeVirt Masquerade — VM gets static IP 10.0.2.2. NAT via virt-launcher pod. All ports forwarded (DNAT).

Masquerade is one of two common ways to connect a virtual machine to a Kubernetes network. The machine sits behind its pod like a laptop behind a home router: it has a private address, `10.0.2.2`, and the pod translates traffic in and out. The other way, *bridge*, has the machine take the pod's address as its own, with nothing translated.

The same commit that added the diagram switched the target cluster from one to the other:

```
      interfaces:
-         - masquerade: {}
+         - bridge: {}
```

It also deleted the firewall rule that translated traffic for `10.0.2.2`. So the hop the diagram draws had already been removed before the diagram reached the repository. It is still absent today. The target cluster's control-plane node reports its internal address as `10.244.0.115`, exactly the address of the pod that hosts it. There is no `10.0.2.2`, and no translation step. The script's narration says "masquerade NAT" anyway, once at the start of the demo and again in the closing summary.

Nothing failed, which is exactly why nobody noticed: a wrong diagram of a working system raises no error. And diagrams travel. The copy of the mind map kept with the briefing material is the same file, byte for byte.

Step six did not need step five

Look again at the demo's logic. Step 4 tries the *name*, and fails. Step 5 teaches the mesh the name. Step 6 succeeds. The audience draws the obvious conclusion: the ServiceEntry made the call possible.

But Step 6 does not call the name. It calls the address:

```
curl -s "http://${TARGET_LB}:8000"
```

`TARGET_LB` is `172.18.255.216`. Any pod in cluster1 can reach that address with or without a ServiceEntry, because it is simply an address on the shared network.

BUILD LOG · 17 SEPTEMBER 2026

On the re-plumbed path, a pod enrolled in cluster1's mesh called the address five times *before any ServiceEntry existed*. All five returned `200`, in 1.1–1.7 milliseconds. Then the demo's ServiceEntry was applied, verbatim. Calls to the address carried on returning `200`, in 1.3–1.6 ms. Calls to the *name* — which a moment earlier had failed with "could not resolve host" — now succeeded too, in 1.7–2.1 ms. The mesh answered the lookup with `240.240.0.1`, an address Istio assigns to such entries by itself.

So the name works, at least on today's mesh. The demo never used it. The step that looked like the cause of the success sat beside the success, not inside it.

A policy nothing read

Step 7 applies a *DestinationRule* to the remote nginx. Among its settings: one request per connection, and eject a backend after three server errors in a row. The script prints *"Traffic policy applied: connection pooling + circuit breaker"*, and then proves it with one request that returns 200.

That request would have returned 200 without the policy. And in *ambient* mode — the mesh design used here, which puts no proxy inside the application's pods — the component on each node, ztunnel, handles connections, not individual HTTP requests. Request-level rules like these need a separate layer-7 proxy, a *waypoint*, and the demo never deployed one. When the "after" picture was built in August, the repository's own comparison table filled in the before column with a single word:

L7 policy — none (ztunnel is L4-only)

The same document, describing a different kind of policy, warns that a rule enforced at the connection layer *"will silently ignore HTTP path and method rules."* Silence is the failure mode again, just as with the diagram. An unenforced policy looks exactly like an enforced one until the day it is needed.

The one leg without a lock

The re-plumbed run also captured what cluster1's ztunnel recorded for the cross-cluster calls. Here is one of those log lines, trimmed, beside an ordinary call between two workloads inside cluster1's mesh:

```
cross-cluster  src.workload="sleep" src.namespace="ch10-before" src.cluster="cluster1"
               dst.addr=172.18.255.216:8000 dst.service="nginx.target-cluster.global"
               dst.workload="nginx-target-cluster" dst.namespace="ch10-before"
               dst.cluster="cluster1"

inside cluster1 src.workload="waypoint-..." src.namespace="demo-apps"
               src.identity="spiffe://cluster.local/ns/demo-apps/sa/waypoint"
               dst.addr=10.244.0.25:15008 dst.service="echo-v1.demo-apps.svc.cluster.local"
               dst.identity="spiffe://cluster.local/ns/demo-apps/sa/echo"
```

Two differences matter.

The call inside cluster1 carries an **identity** at both ends. That is a cryptographic name for each workload, and it is what mutual TLS authenticates. The call also travels to port 15008, the mesh's encrypted tunnel. The cross-cluster call has no identity at either end and uses no tunnel. The

project's own notes say it plainly: *"mTLS terminates at cluster edge; cross-MetalLB traffic is plain HTTP."* The one leg of the journey that crosses the shared network between clusters is the one leg with no encryption and no identity.

The second difference is quieter: `dst.cluster="cluster1"` , `dst.namespace="ch10-before"` . A `ServiceEntry` is a local record, so the mesh files the target cluster's nginx as a workload *in cluster1*, in the caller's own namespace. The mesh's own logs put the remote service in the wrong cluster.

For a board, this is the sentence to carry out of the chapter: when someone says traffic between sites is "in the mesh" or "encrypted in transit", ask which legs they mean.

A word that joined two facts

The console's Live Demo tab, added in the same commit, runs four real probes every five seconds: each cluster calls its own service, and each calls across to the other. Its failure controls scale a service down to zero. The amber badge from Meridian's briefing is computed in

`ui/backend/handlers/cross_cluster.go` :

```
httpbinAlive := ok(probes[0])
nginxAlive := ok(probes[2])
// Failover: a local service is down but the cross-cluster path is still up
failoverActive := (!httpbinAlive && ok(probes[1])) || (!nginxAlive && ok(probes[3]))
```

`probes[1]` is cluster1 calling the *target cluster's nginx*. So when cluster1's httpbin is scaled to zero, and an unrelated nginx in another cluster still answers, the console logs *"Cross-cluster FAILOVER now active — remote service handling traffic."* No traffic meant for httpbin goes anywhere. There is no second httpbin to go to, and nothing reroutes. The probes are real and the facts they report are true; only the word joining them is false. That code is unchanged since March, and the tab is still in the console today, kept as the legacy demo.

The six mesh panels from the same commit went further. None of them imported anything except React and an icon library: no API client, no data source of any kind. The pod addresses they displayed, such as `10.42.0.12` and `10.244.0.60` , and the services named `target-cluster-nginx` , were written into the page. They were retired on 21 August, and the comment left where they used to be is the most honest sentence in this chapter:

none made an API call, and several hardcoded invented pod names and IPs for a target-cluster mesh that no longer exists.

An address with an expiry date

The last confession is the one in the chapter's title. The proxy in cluster2 reaches the virtual machine through its pod's address, typed in by the script. The Service has no selector, so no controller will ever correct that address. The pod it names was created when that morning's target cluster was built. A rebuild — 33.7 seconds, as Chapter 7 measured — replaces the pod, and the

address goes with it. Traffic reaches the pod at all only because it arrives from outside the cluster and is forwarded by the node; as Chapter 8 found, other pods cannot dial a virtual machine's pod address directly.

In fairness, the script looks the address up again each time it runs. But it does not remove what it creates: it prints five cleanup commands and exits. Whatever is left behind holds a snapshot of an address that expires with the cluster.

Three smaller things came with it:

- **A comment the code does not keep.** The selection loop is introduced with *"We select the CP specifically,"* then takes the first matching virtual machine, control plane or worker. It picks the control plane only because `cp` sorts before `workers`.
- **A warning sent nowhere.** Writing endpoints by hand uses an API that Kubernetes deprecated in version 1.33. On today's 1.37, `kubectl` says so — *"v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice"* — and the script's `2>/dev/null` throws the message away.
- **An address that has to be defended.** `172.18.255.216` has had its own row in the project's IP allocation table since the March commit. In August, when cluster2 needed a new address, the pool was widened rather than reuse `.216`, because *"double-booking it would make the two demos mutually exclusive."* A hardcoded address is not a line of configuration. It is a reservation that somebody has to keep honouring.

Why the before picture stayed

BUILD LOG · 21 AUGUST 2026

The proper answer arrived in August as `07-istio-advanced/`: five acts, each of which *"ends with assertions, not screenshots."* Its README opens with a table that is, in effect, this chapter's summary:

	<code>05-istio/</code> (before)	<code>07-istio-advanced/</code> (after)
Istio	1.28.0, ambient, one cluster	1.30.3, ambient, two clusters
L7 policy	none (ztunnel is L4-only)	waypoint + <code>AuthorizationPolicy</code> on SPIFFE identity
Cross-cluster	<code>ServiceEntry</code> + hardcoded MetalLB IP + NodePort proxy	<code>istio.io/global: "true"</code>

A single label on a Service replaces the address book, the fixed address and the hand-written proxy. Chapter 11 tells that story, including the day a waypoint silently broke it. (The "after" column has since moved to Istio 1.31; Chapter 19.)

It would have been easy to delete `05-istio/` that same day. Instead the README records that it *"is left untouched. It stays as the 'before' picture."* A company keeps last year's accounts for the same reason: an improvement you cannot compare against is only a claim.

There is one honest limit, and the README states it too. Act 3 connects **cluster1 and cluster2** — not the target cluster. The target cluster is *"k3s nested in KubeVirt VMs (pod IPs unreachable from cluster1 without a NodePort proxy) and the warm pool destroys and rebuilds it, which would break mesh membership each time."* On 17 September the target cluster had no service mesh at all: four namespaces, none of them Istio's. The virtual-machine clusters are the product this book is about. The only cross-cluster path they have ever had is the one in this chapter.

* * *

MERIDIAN · THE FOLLOWING WEEK

Vikram's list had six lines. Anita read them in order and stopped at the fourth.

"Plain HTTP between the sites."

"Between the clusters, on our own network. Inside each cluster, it's encrypted. The crossing isn't."

"You told the bank it all went through the mesh."

"I told the bank what the diagram said."

She didn't cross anything out. "Here is what we do. The before picture stays. It is the best evidence we have that the new design is better. But every screen that shows it gets a label: *hand-plumbed; not encrypted between clusters; not failover*. The badge loses the word. The diagram comes out of the sales folder until it matches the machines. And the VM clusters get their own line on the roadmap, because the new design does not reach them yet."

"And the bank?"

"We call them before they find it," Anita said. "I would rather show a customer our confession than have them discover it."

The ledger

- **Built (12 March 2026):** cross-cluster calls in both directions between a Kind cluster and the VM-hosted target cluster, over a path plumbed by hand. The eight-step demo, the mind map, a live console tab and six mesh panels all arrived in one commit: 36 files, 5,362 lines.
- **Measured (17 September, path re-plumbed by hand; Kubernetes 1.37, Istio 1.31 on cluster1):** 200 in 1.1–1.7 ms by address *before* any ServiceEntry existed; 1.7–2.1 ms by name after it.
- **Found:** a diagram describing a translation step that its own commit removed; a demo step that did not depend on the step before it; a request-level policy given to a layer that sees only connections; a cross-cluster leg with no identity and no encryption; a failover badge joining two unrelated probes; six panels with no data source; a deprecation warning discarded.
- **Kept on purpose:** the whole before picture, as the baseline Chapter 11 is measured against.
- **Still open:** the VM-hosted target cluster has no service mesh, and no after picture.
- **Cleaned up:** everything created for the measurement was deleted; .216 was released.

Ask your team

1. Which step of our best demo would still succeed if we skipped the step before it?
2. When a dashboard says *failover*, *encrypted* or *healthy*, what exactly does the code behind that word measure — and when did someone last read it?
3. Which addresses have we typed in by hand, and what happens to each of them the next time we rebuild?

Open the repo

- `05-istio/cross-cluster-demo.sh` — the eight steps. Step 5 (the ServiceEntries) begins at line 460; Step 6's call by address is at line 537.
- `05-istio/cross-cluster-mindmap.html` — the chapter figure. A byte-identical copy sits in the author's briefing material (Appendix H), so correcting one leaves the other wrong.
- `git show 44aa28a -- 03-target-cluster/target-cluster.yaml` — `masquerade` becomes `bridge` in the same commit as the diagram.
- `ui/backend/handlers/cross_cluster.go` (line 151) — the failover rule;
`ui/frontend/src/components/visual/CrossClusterDemo.tsx` — the badge.
- `ui/frontend/src/pages/VisualDashboard.tsx` — the retirement note for the six panels.
- `07-istio-advanced/README.md` — the before/after table, and the Act 3 note on why the target cluster is left out.
- `kubectl --context kind-cluster1 logs -n istio-system -l app=ztunnel | grep "connection complete"` — check any connection for `src.identity` and `dst.identity`.

Chapter 11 — One Label

Marking a Service global is the ENTIRE cross-cluster config.

— `07-istio-advanced/act3-multicluster/03-global-service.yaml`

MERIDIAN · A MONDAY IN SEPTEMBER · THE PLATFORM TEAM'S ROOM

Vikram had cleared the whiteboard of the old diagram and written two words where it used to be: *one label*.

"That's the replacement," he said. "No fixed addresses, no hand-written proxy, no address book. We mark a service as global, and the mesh in each site learns where the copies in the other site are. If a site loses its copies, calls go to the other one."

Anita had the confessions list from the week before in her hand. "The last design had a demo that worked too."

"This one is different. Every step ends in a check that reads the answer back from the system. Not a screenshot."

"Good. Then show me one of those checks going red."

"Red?"

"Take the service down in both sites. Nowhere left to fail over to. If the check is real, it fails." She put the list down. "I've been shown plenty of green. I want to see what red looks like before I believe the green."

* * *

Decisions you can only make once

BUILD LOG · 21 AUGUST 2026

The replacement arrived as `07-istio-advanced/`: five acts, each of which, in its README's words, *"ends with assertions, not screenshots."* This chapter covers the first three. Acts 4 and 5 are Chapter 12.

Before any act, the README opens with a section called *"Order is not optional."* It is the most important page for a board, because it is about decisions that can only be taken on the first day.

Joining two clusters into one mesh needs four things fixed when the mesh is installed: a mesh name, a name for each cluster, a network name for each, and a flag that switches on the multi-network design. It also needs both clusters to trust one **root certificate** — the single authority from which every workload's identity is issued. Istio reads that root only when it starts. The install script says what happens if it starts without one:

if istiod starts without it, it self-signs its own root, and the only repair is to reinstall, which reissues every workload certificate in the mesh.

So the root is generated first, and both meshes are installed ready for two clusters, even though the second cluster isn't demonstrated until Act 3. The root is created locally, with no outside certificate authority involved: *"that is the sovereign-cloud story, and it also means the demo works air-gapped."*

For a CxO the point is general. Some architecture choices are cheap on day one and a full rebuild on day four hundred. Whether you will ever run more than one site is one of them. On 17 September both clusters still reported the same root fingerprint, and each control plane reported its peer as `synced`.

The first refusal

Act 1 puts a front door on cluster1 using the Kubernetes **Gateway API**. The door is declared as a single object, a `Gateway`. There is no installer and no hand-built proxy. Istio reads the object and creates the proxy Deployment and its Service itself. On 17 September both existed, as `edge-istio`, serving `172.18.255.201`.

The act's real subject is *who may do what*. Three namespaces play three roles. `gateway-infra` belongs to the platform team, which owns the Gateway and its certificate. `demo-apps` belongs to an application team, which owns the routes that send traffic to its services. `rogue-tenant` is labelled untrusted. The Gateway accepts routes only from namespaces labelled as applications.

Then the rogue tenant tries to take over the application team's hostname. It writes a route claiming `echo.demo.istio.local` and pointing it at a service of its choosing. The Kubernetes API accepts the object: the tenant is allowed to write objects in its own namespace. The Gateway refuses to attach it. This is the route's status, read back from the cluster on 17 September:

```
Accepted=False NotAllowedByListeners: hostnames matched parent hostname
"*.demo.istio.local", but namespace "rogue-tenant" is not allowed by the parent
```

Traffic to the hostname carried on reaching the application team's service.

That is governance in the platform, not in a wiki page. Permission to write configuration is not permission to receive traffic, and the refusal is written down where anyone can read it. The rest of Act 1 is ordinary good practice, and it was verified again on 17 September: HTTPS with a locally rooted certificate, `200`; a 90/10 canary release, which measured 8% on 60 requests; and a header that pins internal users to the new version, which held on 10 of 10 requests.

The second refusal

Act 1 has one more beat, and it did not appear in the August commit's list of verified results.

A **ListenerSet** lets an application team add its own entrance to the platform team's Gateway — here, a listener for `tenant-a.demo.istio.local` on port 8443 — without editing the Gateway itself. The act's ListenerSet borrows the platform team's certificate, which lives in `gateway-infra`. On 17 September its status read:

```
Accepted=False certificateRef echo-tls/gateway-infra not accessible to a Gateway
in namespace "demo-apps" (missing a ReferenceGrant?)
```

The Gateway had refused the demo's own tenant, for the same reason it refused the rogue one. Borrowing something from another team's namespace requires that team to sign a permission slip first, a `ReferenceGrant`. Nothing in the repository contains one.

The refusal is correct; the demo is what's wrong. And nobody noticed, because this step prints a status instead of asserting one. It then prints *"Gateway now reports attached ListenerSets in its status"*, followed by the Gateway's listener names — `http https`, with nothing attached — and moves on. The verification script does not check it. Both refusals come from the same rule. Only one of them was ever checked.

Identity, not address

Act 2 is the fix for the policy in Chapter 10 that nothing enforced.

In ambient mode, the component on each node (ztunnel) works with connections, not HTTP requests. For request-level rules, Act 2 adds a **waypoint**: one proxy per namespace, declared, like the front door, as a Gateway object. No proxy is placed inside any application pod. On 17 September each echo pod ran exactly one container, `echo`.

Two client deployments are identical in every respect except their Kubernetes service account, `trusted` or `untrusted`. The mesh turns that account into a cryptographic identity: `spiffe://cluster.local/ns/demo-apps/sa/trusted`. The policy names identities, paths and methods, not IP addresses. Measured on 17 September:

Caller	Request	Result
<code>trusted</code>	<code>GET /secure</code>	<code>200</code>
<code>untrusted</code>	<code>GET /secure</code>	<code>403</code> — <i>"RBAC: access denied"</i>
<code>trusted</code>	<code>POST /secure</code>	<code>403</code>
<code>untrusted</code>	<code>GET /public</code>	<code>200</code>

The application behind the waypoint also learned who was calling, without doing any cryptography itself. It received a header naming both the waypoint and the original caller:

```
x-forwarded-client-cert=By=spiffe://cluster.local/ns/demo-
apps/sa/waypoint;URI=spiffe://cluster.local/ns/demo-apps/sa/trusted
```

The act also switches the trusted client's service account to `untrusted`, and access drops to `403` with no change to the policy and no restart of the application. That beat was not re-run for this chapter.

The README flags the one mistake that undoes all of this. If the policy is attached to the workloads rather than to the waypoint, it is enforced by ztunnel, which cannot see paths or methods — and it *"will silently ignore HTTP path and method rules."* That is the same shape of failure as Chapter 10's circuit breaker: a rule that looks enforced and isn't.

For a board, this is the counterpart to Chapter 10's plaintext link between clusters. Access is decided by a verifiable identity for the caller, not by where the caller happens to sit on the network. And it can be revoked without touching the application.

One label

Act 3 joins cluster1 and cluster2. It needs only three things, because the hard decisions were taken at install time.

- **An east-west gateway on each side** — the doorway through which one cluster's mesh reaches the other's, over the mesh's encrypted tunnel on port 15008. cluster1's is at `172.18.255.202` ; cluster2's is at `.221` , because `.216` still belongs to Chapter 10.
- **A credential each way**, so each control plane can read the other cluster's list of services. The README records a trap here: on this setup, unless the command is told the real API server address, the credential embeds `127.0.0.1` , and *"the peer cluster then sits at `STATUS=timeout` while everything else looks healthy."* The act waits for both directions to report `syncd` before it continues.
- **One label** on the Service, in both clusters:

```
metadata:
  name: echo
  labels:
    istio.io/global: "true"
```

Set that against Chapter 10: a ServiceEntry in each direction, a MetalLB address typed into an annotation, a proxy Service with a pod address written in by hand, and a NodePort on every node. All of it, replaced by one line.

Then comes the moment the act calls *"THE MONEY SHOT."* Every copy of echo in cluster1 is scaled to zero, with no configuration change anywhere. The same service name is called twenty more times, and the act prints the tally.

Morning and evening

BUILD LOG · 21 AUGUST 2026

The commit that added the acts landed at 06:36 and reported the tally: *"Act 3 success=20 failed=0 — cluster1 scaled to zero, traffic failed over to cluster2 with no config change."* The verification script: 22 passed, 0 failed.

At 18:34 the same day, a second commit recorded a problem. Act 2 enrolls the whole namespace in its waypoint. A waypoint, it said, resolves only *local* endpoints, so once the namespace is enrolled, scaling the local copies to zero gives "*503 on every request with the waypoint in path.*" The evidence was the waypoint's own endpoint list, which held only cluster1's pod addresses. The fix was to take the global service out of the waypoint's path during failover. The console's failover button was built to do exactly that, and to put it back afterwards.

Put those two commits side by side. The morning's run came straight after Act 2, with the namespace already enrolled. The evening says that configuration fails every request. Both cannot be describing working failover — unless the morning's tally was measuring something other than success.

What the counter counted

BUILD LOG · 17 SEPTEMBER 2026

Here is how Act 3 decides that a call succeeded:

```
call() { $K1 exec -n "$NS_APPS" deploy/client-trusted -c curl -- \
    curl -s --max-time 10 http://echo/ 2>/dev/null || true; }
...
b=$(call)
if [[ -n "$b" ]]; then okc=$((okc+1)); else failc=$((failc+1)); fi
```

A call counts as a success if it returns *any text at all*. The HTTP status is never examined.

On 17 September Anita's test was run for real. Echo was scaled to zero in **both** clusters, leaving nothing anywhere to fail over to, and the service was called twenty times. Every call returned **503**, with a nineteen-byte body: "*no healthy upstream.*" Under Act 3's rule, nineteen bytes count as success. The act would have printed:

```
success=20 failed=0
```

— the same line, character for character, that it prints for a perfect failover. The check could not go red.

The verification script didn't cover the gap either. Its Act 3 section asserts four things: that both east-west gateways are programmed, that a credential is present, and that the label is present. It checks that the label *exists*, not that it *does* anything. None of its twenty-two checks sends a request to a scaled-down service.

A correct counter did exist. The console's failover action, written that evening, counts only status **200**, and it measures how long convergence takes instead of assuming it. It is in a different file, and `run.sh` was never updated to match.

What the morning's twenty calls actually returned cannot be recovered now. The evening's measurement says `503` ; the morning's counter would have printed the same line either way.

What is actually true

Measured on 17 September by HTTP status, on Istio 1.31 and Kubernetes 1.37, with every change put back afterwards:

Failover works — even with the waypoint in the path. The namespace was left enrolled in the waypoint, exactly as the acts leave it, with no opt-out label. cluster1's copies of echo were scaled to zero. All 30 of the next calls returned `200` , and every one was answered by a pod in cluster2. The first call after the last local pod disappeared was already one of them. The waypoint's endpoint list explains why. In August it held only local pod addresses. Now it also holds cluster2's east-west gateway:

```
envoy://inner_connect_originate/172.18.255.221:15008    HEALTHY    inbound-vip|80|http|echo.demo-apps.svc.cluster.local
```

The problem recorded that August evening no longer reproduces. The opt-out the console applies is now unnecessary, though harmless. What changed is not established. The Istio upgrade of Chapter 19 is the obvious candidate.

The other August gotcha still reproduces. Act 2 also leaves a route that splits in-mesh traffic between the `echo-v1` and `echo-v2` Services — and those Services are not global. With that route re-applied and cluster1's copies scaled to zero, 20 of 20 calls returned `503` . After the route was deleted, the second call returned `200` , 0.15 seconds later, and the next 20 all did. A route that points at local-only services overrides global routing. Act 3 deletes the route for exactly this reason.

Healthy traffic does not prefer home. The Service's manifest says *"Local endpoints are preferred; remote ones are used when no healthy local endpoint exists."* After its baseline calls, the act prints *"traffic served (locality-aware routing prefers in-cluster endpoints)"* — without comparing any numbers. Measured with everything healthy, 60 calls split exactly **30 to cluster1 and 30 to cluster2**. The act's own local-or-remote classifier recognises only one of cluster2's three pod names. That day every remote call happened to land on a different one, so the act would have counted all of them as local.

On one laptop, the price of that split is a fraction of a millisecond: 2.67 ms mean for calls served locally, 2.96 ms for calls served across. Between two real sites, it is a network hop for half of every service's traffic, on a day when nothing is wrong. That is a latency and data-transfer bill, set by a default nobody chose.

What one label bought

It is worth being precise about the verdict, because both halves are true.

The label is real. One line replaced every piece of Chapter 10's plumbing. Failover works and was measured by status code, and the connection between clusters carries identities through an encrypted tunnel. The front door refuses tenants it should refuse — including, by accident, the demo's own.

But the simpler the configuration became, the more the risk moved into the checks. Chapter 10's demo could not fail because its steps didn't depend on each other. This one could not fail because its counter could not tell a response from an error page. The replacement was far better engineering, and it inherited the older demo's habit of printing green.

One more limit carries over from Chapter 10 unchanged. Act 3 joins the two Kind clusters. The target cluster — the virtual-machine clusters this book is about — is not in the mesh, for the reasons given there.

* * *

MERIDIAN · TWO WEEKS LATER

Vikram put two terminal screenshots side by side on the wall screen. Both ended with the same line: `success=20 failed=0`.

"The left one is failover working," he said. "The right one is both sites down."

Anita looked for a long time. "So the check I asked to see go red can't."

"It couldn't. It counts replies, and an error page is a reply. The failover itself works — we measured it properly. Every one of thirty calls came back from the other site."

"And when nothing is broken?"

"Half our traffic crosses to the other site anyway. Nobody set that. It's the default."

"That one isn't an engineering question," Anita said. "That's a bill." She took a pen to the list from the week before and added a line at the bottom. "New rule. No check goes on a dashboard of mine until someone has watched it fail on purpose. And the traffic split comes to me with a cost attached, and we choose it."

Failover working

cluster1 scaled to 0 · cluster2 serving

```
$ ./act3-multicluster/run.sh
scaling echo-v1 → 0 replicas (cluster1)
20 requests via client-trusted...
```

success=20 failed=0

✓ Traffic failed over to cluster2

Total outage

both clusters scaled to 0 · nothing serving

```
$ ./act3-multicluster/run.sh
scaling echo-v1 → 0 (cluster1 + cluster2)
20 requests via client-trusted...
```

success=20 failed=0

✓ Traffic failed over to cluster2

identical

Why both print the same line — the check never looks at the status code:

```
b=$(call); if [[ -n "$b" ]]; then okc=$((okc+1)); else failc=$((failc+1)); fi
```

Any non-empty response body counts as a success. A check that cannot fail is not a check.

Source: 07-istio-advanced/act3-multicluster/run.sh (quoted in Appendix G.4).

Figure 11.3 — One check, two opposite realities, the same output.

The ledger

- **Built (21 August 2026):** a Gateway API front door with a team ownership split; request-level authorization on cryptographic identity, with no proxy in any application pod; one mesh across two clusters under a single local root certificate; a cross-cluster service declared with one label.
- **Re-verified (17 September; Istio 1.31.0, Kubernetes 1.37):** verification script 22 passed, 0 failed. Rogue route refused. Identity rules 200 / 403 / 403 / 200. Caller identity delivered to the application.
- **Measured by HTTP status:** cluster1 scaled to zero → 30 of 30 200 , all from cluster2. Both clusters at zero → 20 of 20 503 , which Act 3's counter reports as **success=20 failed=0** .
- **Found:** the ListenerSet beat refused for a missing ReferenceGrant , and never asserted; healthy traffic split 30/30 across clusters, despite the manifest's comment; a local/remote classifier that sees one of three remote pods; a verification section that checks configuration, not behaviour; one August gotcha that no longer reproduces, and one that still does.
- **Corrected elsewhere:** Chapter 19 had cited **success=20 failed=0** as proof of failover. It now cites the status-code measurement.
- **Cleaned up:** every replica count, label and route changed for measurement was put back, and checked.

Ask your team

1. When did we last watch this check fail — deliberately?
2. If our primary site and our backup were both down, which of our dashboards would still be green?
3. How much of our traffic crosses between sites when nothing is wrong, and who decided that?

Open the repo

- `07-istio-advanced/README.md` — the before/after table, "Order is not optional", and "Gotchas found while building this".
- `07-istio-advanced/01-install/gen-mesh-ca.sh` and `istio-values-cluster1.yaml` — the day-one decisions.
- `07-istio-advanced/act1-gateway/03-gateway.yaml` , `05-rogue-route.yaml` , `06-listenerset.yaml` — both refusals.
- `07-istio-advanced/act2-waypoint/01-authz.yaml` — the identity policy, and the comment about attaching it to the waypoint.
- `07-istio-advanced/act3-multicluster/run.sh` — `call()` and the tally; `03-global-service.yaml` — the label.
- `ui/backend/handlers/istio_actions.go` — `HandleIstioFailover` : the counter that checks for 200 , with convergence measured.
- `07-istio-advanced/verify.sh` , section 5 — what the Act 3 checks do and do not assert.
- `istiocctl proxy-config endpoints deploy/waypoint.demo-apps --context kind-cluster1 | grep echo` — whether the waypoint can see the other cluster.
- **Without the repo:** Appendix G.4 quotes both counters side by side — the act's non-empty-body check, and the console's status-code check.

Chapter 12 — The Model Is Just Another Service

Model traffic now carries mesh policy + telemetry, same as any service.

— `07-istio-advanced/act4-ai-gateway/run.sh`

MERIDIAN · A WEDNESDAY IN SEPTEMBER · THE BOARDROOM, THE EVENING BEFORE THE BOARD MEETS

Anita had the board pack open at the page headed *Artificial intelligence*. It asked two questions, the same two the board had asked every quarter since the first pilot: *Who can use it?* and *What does it cost?*

"This quarter we answer them with a screen, not a paragraph," she said.

Vikram put the console on the wall. "Every model call goes through an AI gateway now. Same as any other service in the mesh — it gets policy, and it gets measured. And the visual mode shows the whole estate, live."

"Then tomorrow I can say every model call passes one control point, and the board can watch the traffic on that screen."

"That's the design."

Anita looked at the screen for a while. It was full of green. "Before I say it to the board," she said, "show me yesterday's model calls on it. And show me who was allowed to make them."

* * *

A gateway for models

BUILD LOG · 21 AUGUST 2026

Act 4 of `07-istio-advanced/` is the one its README calls Istio's headline feature for this release: an **AI gateway**.

It uses the same Gateway API object as Chapter 11's front door, with one field changed.

`gatewayClassName: istio-agentgateway` swaps the general-purpose proxy for **agentgateway**, "*a proxy purpose-built for AI/agent traffic*." The class is experimental and has to be switched on when Istio is installed, one more day-one decision. The gateway takes the address `172.18.255.203`, and one route sends everything it receives to the models.

The models are not in the cluster. They run on the laptop's own GPU, under Ollama, and the cluster reaches them through a shim: a Service with a hand-written endpoint pointing at the host, `172.18.0.1`. This is Chapter 10's pattern, used here on purpose, because keeping models off the cluster "*keeps ~4GB of RAM per replica free*."

For a board, the reason to put a gateway in front of models is the reason to put one in front of anything: a single place to decide who may call which model, to count what they use, and to switch providers without touching every caller. On sovereign infrastructure that matters twice over, because the models, the GPU and the data never leave the building.

The act ends by asserting the result in one line, which is this chapter's epigraph: *"Model traffic now carries mesh policy + telemetry, same as any service."* The console's AI Gateway panel says the same, and it fetches the model list *through* the gateway, *"so it proves the data path, not just the config."*

That is three claims: a data path, policy, and telemetry. They were tested separately.

The data path is real

BUILD LOG · 17 SEPTEMBER 2026

Through the gateway, the model list came back with all 14 of the host's models, and a chat completion returned `ok`.

The more useful question is what the extra hop costs. With `qwen2.5:7b` already loaded, the same short request was sent seven times directly to Ollama and seven times through the gateway, alternating between the two:

Path	Median	Range
Direct to host Ollama	188.4 ms	173.0–216.4 ms
Through the AI gateway	179.7 ms	164.9–198.7 ms

The gateway adds nothing measurable. The small difference in its favour is noise, not speed.

Streaming — the way the agents actually receive answers, a token at a time — passed through intact. A request to count to twenty arrived as 51 chunks by both paths, with the first byte after 0.163 seconds direct and 0.171 seconds through the gateway.

So the first claim holds, and it is good news. Putting a control point in front of the models costs nothing a user would notice.

A door with no lock

On 17 September there was no authorization policy of any kind in the gateway's namespace. With no credential, from the laptop:

Request through the gateway	Result
<code>GET /api/version</code>	<code>200</code>
<code>GET /api/ps</code> — which model is loaded right now	<code>200</code>
<code>POST /api/show</code> — a model's full configuration	<code>200</code>

The route forwards every path, so Ollama's administrative endpoints — pulling new models, deleting models, creating them — are behind the same route. Those were deliberately not called.

The gateway is also not a mesh workload in the way the echo services are. Its pod carries the label `istio.io/dataplane-mode=none`, and the hop from the gateway to the model is plain HTTP to the host, which no mesh can enrol. Compare the agent platform's own serving endpoints: Chapter 11's verification script cannot even check them without first fetching a bearer token.

The gateway did not create this exposure. Ollama on the host listens on every interface (`*:11434`), so anything that can reach the host can reach the models without going near the gateway. But the gateway did not close the exposure either. A control point is where a policy *could* live, and on 17 September none did.

Counted under another name

Chapter 11's verification script and the console's Observability panel both judge "is telemetry flowing?" the same way: the number of `istio_requests_total` series in Prometheus. On 17 September the panel reported 16 series, labelled `source: live`.

About forty requests had just been sent through the AI gateway. Not one of those sixteen series mentioned the gateway or the model backend. Every series came from Act 1's front door or Act 2's waypoint — including, as it happens, the `503 s` from Chapter 11's failover experiments. The per-node mesh component recorded nothing either: no connection metrics had the gateway as their source.

Kiali, the mesh map from Act 5, agreed. Its graph for the last hour had **0 nodes and 0 edges** for `ai-gateway`. For `demo-apps`, it had 10 nodes and 14 edges.

The gateway does keep its own count. Prometheus was scraping it successfully, and a family of `agentgateway_*` metrics recorded requests by method, status, route and backend. The problem is what they omit: there is **no model name and no token count**. Those are the two dimensions that turn a request count into a cost. And nothing reads these metrics — not the console panel, not Kiali, not the verification script.

So the model traffic is measured, but in a metric nobody displays, and without the two labels that would answer the board's second question.

Nobody uses it

Then there is the question of whose traffic this is.

The platform's three serving agents run on cluster2, and each names its model endpoint in its own configuration:

```
baseUrl: http://host-ollama.symposium-system.svc.cluster.local:11434/v1
```

That is a second shim, on cluster2, going straight to the host. The AI gateway is on cluster1. None of the three agents sends a single request through it. The comment in the gateway's route file was carefully worded: the endpoint the agents speak "*is reachable through the mesh.*" Reachable, yes. Reached, no.

The act's own check even gets in the agents' way. By default it asks a different model, `qwen3.5:4b`. The GPU in this laptop holds one model at a time (Chapter 9). On 17 September the act's single completion took 8.2 seconds, reasoned through 556 characters of hidden thinking before replying `ok`, and evicted `qwen2.5:7b` — the model two of the agents use. Putting it back cost a cold load of 3.30 seconds, which the next agent request would have paid.

For a CxO, this is the chapter's central finding. The governance layer for AI exists, and it works. It stands *beside* the AI fleet, not in front of it.

An honest limitation about the wrong thing

Act 4 has a second tier, and it comes with the most candid passage in the repository:

HONEST LIMITATION, worth saying out loud on stage: the EPP schedules on KV-cache utilization and queue depth, which vLLM and Triton export and Ollama does NOT. What follows demonstrates the routing API surface, not genuine load-aware scheduling.

Tier 2 is the Gateway API **Inference Extension**. It groups model servers into a pool and picks a server for each request based on how busy each one is. The console repeats the limitation in an amber box. It is true: Ollama does not publish those load figures.

But Tier 2 never ran. The Inference Extension was not installed on either cluster. Its installer script and its manifest are referenced by nothing — not the act, not the Makefile, not the demo runner. The act checks whether the extension exists, finds it missing, prints the limitation, and moves on.

And the manifest could not have been applied. Checked against the schemas in the extension version its own installer pins (v1.0.1):

- **The pool** is declared at version `v1`, but uses field names from the older `v1alpha2` (`targetPortNumber`, `extensionRef`). Version `v1` requires `targetPorts` and `endpointPickerRef`.
- **The objective** is declared in an API group that does not exist in that release, and uses two fields, `modelName` and `criticality`, that do not exist in the one that does. Its real fields are `poolRef` and `priority`.
- **The pool's selector** matches no pods: the model backend is a shim Service, not a set of pods.
- **The endpoint picker** it names, `ollama-epp`, is not deployed anywhere.

The limitation is real, but it is a limit the demo would only have reached after four more basic problems. Honesty about a limit you would hit later can hide the ones you hit first.

The stage

BUILD LOG · 21 AUGUST AND 17 SEPTEMBER 2026

Act 5 installs **Kiali**, the mesh map, and **Prometheus**, which stores its metrics, on cluster1. The act states a principle this book would sign: *"A Kiali screenshot is not evidence; the metric count is."* The console's Observability panel follows it, showing the series count rather than just whether the pods are up — because *"a Ready pod alone does not mean telemetry is flowing."*

Three facts about that stage belong in front of a board:

- **The map is open to anyone.** Kiali at `172.18.255.204` reported its authentication strategy as `anonymous`. Anyone who can reach the address can browse every namespace, workload and identity in the mesh, and a map of your systems is itself sensitive. This is the default in Istio's sample installation, which is what the act uses.
- **The history is temporary.** Prometheus is set to keep fifteen days of data, on storage that disappears whenever its pod restarts. On 17 September its history began on 16 September at 07:29 UTC, the moment of Chapter 19's rebuild.
- **The stage is downloaded, not owned.** The act fetches both installations at install time from Istio's public GitHub branch for the release, not from a pinned, local copy. The mesh's root certificate was created so the demo *"works air-gapped"*. Its observability does not.

The screen that says live

The console's visual mode was rebuilt on 21 August. It retired six panels that made no API calls (Chapter 10) and put up a comment saying *"Everything below reads live cluster state."* Most of it does. The mesh overview reads whether both clusters share a root certificate. The AI Gateway panel lists models through the gateway. The Observability panel counts series. Each reports where its data came from.

One view was kept from before, with a note that it *"is real."* It is the topology view: the network map, drawn as a force-directed graph. On 17 September its API returned 21 nodes, labelled `source: live`.

Fourteen of those nodes do not exist:

- In cluster1, the `mc-demo` namespace: `httpbin`, its pod, `sleep`, the load balancer on `.200`, and the ServiceEntry for `nginx.target-cluster.global`.
- In cluster2, the proxy on `.216`.
- In the target cluster, the whole Istio installation and sample application: `istiod`, two ztunnels, `nginx`, its NodePort, its pod, `sleep`, and the ServiceEntry back to cluster1.

That is Chapter 10's March picture. The graph showed thirteen of the fourteen as **healthy**. The one exception is a ztunnel marked `error` with the note *"Not Ready"*, typed into the source code. The worker virtual machine appeared as `degraded`, also typed in. The real worker was `Running` and ready.

The mechanism is simple, and it is worth understanding because it is so common. The handler starts from a fixed drawing and tries three live lookups. It labels the result `live` if none of those lookups *returns an error*:

- **Cluster2's node list** succeeded, and correctly updated one node.
- **The virtual machines** were found, but they are matched to the drawing by exact name: `target-cluster-cp` and `target-cluster-worker`. The real machines are `target-cluster-cp-s9j8h` and `target-cluster-workers-rz9xt-m5x6z`. Nothing matched, so nothing changed.
- **The target cluster** was queried for Services in a namespace called `sample`, which no longer exists. In Kubernetes, listing an empty or missing namespace returns an empty list, not an error. The lookup "succeeded", and the target cluster was marked healthy.

The only live number the target side received was `ztunnelPods: 0/0 ready`, attached to a node drawn as healthy.

The same picture shows up somewhere else, too. The briefing given to one of the platform's AI agents still tells it that cluster1 runs *"Istio ambient mesh, httpbin/sleep."*

For a CxO, the lesson is general. *Live* on a screen is a claim about the plumbing, not the picture. What a status label needs to tell you is not whether its queries ran, but whether the thing drawn was found.

* * *

MERIDIAN · LATER THAT EVENING

Vikram's answers fitted on one side of a page. Anita read them out loud, one at a time.

"Who can use the models: anyone who can reach the network. What it costs: the gateway counts requests, but not which model and not tokens. How much of our AI traffic goes through the gateway: none of it — the agents go around. And the screen." She looked up. "Fourteen of the things on that map don't exist."

"Thirteen of them are drawn green."

She was quiet for a moment. "The design is right," she said at last. "A model *is* just another service. So it gets the same three things every service gets from us: an identity, a lock, and a meter. The meter counts tokens by model, because that's the bill. The agents go through the gateway, or we stop calling it the AI control point. And nothing goes on a board screen unless every box on it traces back to a lookup that could have come back empty."

"And tomorrow?"

"Tomorrow I tell them the truth," Anita said. "We have built the control point. Our AI doesn't use it yet. I'd rather present a gap we've measured than a screen we haven't."




Figure 12.2 — A view labelled "live", drawing fourteen things that do not exist.

The ledger

- **Built (21 August 2026):** an AI gateway on cluster1 in front of the host GPU's models; a load-aware second tier designed and documented, with its limitation stated; Kiali and Prometheus; a visual mode rebuilt around live data.
- **Measured (17 September):** no measurable latency cost through the gateway (179.7 ms against 188.4 ms, median, warm); streaming intact (51 chunks, first byte 0.17 s).
- **Found:** no authentication or authorization on the model gateway; model calls absent from the mesh's request metric and from Kiali (0 nodes); the gateway's own metric has no model or token labels, and nothing reads it; none of the three agents routes through the gateway; the act's check evicts the agents' model (3.30 s to reload); Tier 2 never installed, and its manifest invalid against its own pinned version; Kiali open anonymously; Prometheus history on temporary storage; a topology view labelled `live` in which 14 of 21 nodes do not exist.
- **Restored:** `qwen2.5:7b` reloaded onto the GPU and confirmed. No cluster objects were created or changed.

Ask your team

1. What share of our AI calls passes through the control point we show the board — and which calls go around it?
2. Can we say, per model, how many tokens we used last month, and from which system?
3. On our operations screens, what does "live" actually mean — and what would the screen show if the thing it draws were deleted?

Open the repo

- `07-istio-advanced/act4-ai-gateway/01-agentgateway.yaml` , `00-host-ollama.yaml` , and `run.sh` — the gateway, the shim, the claim, and the honest limitation.
- `07-istio-advanced/act4-ai-gateway/02-inference-pool.yaml` and `install-inference-ext.sh` — compare against the v1.0.1 CRD schemas.
- `06-symposium/cluster2-agent.yaml` (and the other agent manifests) — `baseUrl` , where the agents' model traffic really goes.
- `07-istio-advanced/act5-observability/run.sh` — the stage, and where it is downloaded from.
- `ui/backend/handlers/mesh.go` — the fixed drawing, `vmNodeID` , and `applyTargetOverlay` .
- `ui/backend/handlers/istio.go` — `HandleIstioAIGateway` and `HandleIstioObservability` , the panels that do read live state.
- `curl -s localhost:8080/api/v1/mesh/topology` with the console backend running — the graph, node by node.
- Prometheus query `agentgateway_requests_total` — the metric that counts model traffic, and the labels it lacks.

PART IV

The Brain

Agentic operations.

Chapter 13 — An Agent as a First-Class Object

Every `AgentRun` remains queryable as a first-class Kubernetes object.

— Architecture Constraint 6, `Sovereign_Cloud_Agentic_Strategy.md`

MERIDIAN · A WEDNESDAY IN AUGUST · THE CTO'S OFFICE

The vendor's last slide said *AI SRE copilot — connected in minutes*, and the account executive let it sit on the screen while she waited for questions.

Anita Rao had three, written on a pad before the meeting started.

"Where does the model run?"

"In our cloud. Your cluster connects out to us through a lightweight agent. Nothing inbound."

"What did it do last Tuesday — exactly?"

"Every action is in the activity feed in our console, retained according to your plan, and exportable."

"If we cancel, what do we keep?"

"Your exports."

Anita thanked her, and waited for the door to close. Then she turned to Vikram Iyer, who had said nothing for forty minutes. "Same three questions. Ours."

He put a terminal on the wall screen and typed.

```
$ kubectl get agents.symposium.ai -n symposium-system
```

NAME	PHASE	ACTIVE AGENTS	TOTAL RUNS	AGE
cluster2-agent	Serving	1		12h
mesh-sre-agent	Serving	1		12h
target-cluster-agent	Serving	1		12h

"Where does the model run? On a machine in this building. What did it do last Tuesday? Every conversation creates an object in the cluster — the task, the answer, how many commands it ran, how many tokens it used, which account it ran as. You list them with the same command you'd use to list anything else. And if we cancel —" he shrugged. "There's nothing to cancel. It's our cluster's API."

Anita looked at the table for a while.

"So the agent isn't a service we call," she said. "It's a workload we run."

"Under the same permission system as everything else we run."

"Then I'm not choosing an AI vendor," she said. "I'm choosing rules."

* * *

BUILD LOG · FEBRUARY–SEPTEMBER 2026

From a chat window to an object

The platform's AI tab did not start as a control plane. In February 2026 it was a chat window backed by kagent, talking over an agent-to-agent JSON-RPC protocol. In April the first Symposium personas appeared — a cost analyzer and an incident responder, written as demo agents. By late June, the tab's chat proxy had been switched to Symposium and spoke the OpenAI Chat Completions format, the request shape nearly every AI client library already understands. By August the AI tab no longer rendered its own chat at all; it embedded Symposium's console directly.

Each step moved the agent further from being a feature of the UI and closer to being a thing in the cluster. That is the idea this chapter is about, and the strategy document states the ambition behind it in one sentence:

"You are not building 'Claude for Kubernetes.' You are building a sovereign, air-gap-capable, VM-native, mesh-aware agentic control plane — a category that barely exists in the open-source world today."

The same document names the three assumptions this platform breaks. Most agentic Kubernetes work assumes pod-native workloads on a single cluster with cloud models. Here, virtual machines are first-class; there are *"two control planes"*, and an incident on one is not an incident on the other; and *"sovereignty is a hard constraint. External LLM egress is often forbidden."* An agent that operates this platform has to live inside the same boundary as the platform.

Five nouns

Symposium describes agents the way Kubernetes describes everything else: as custom resources, reconciled by controllers. Five matter for this part of the book.

- **Agent** — a persona. Which model, where to reach it, which tools, which policy, and a standing briefing about the environment.
- **AgentRun** — one execution of that persona. It behaves like a Job: it is created, it runs a pod, it finishes, and its result stays behind.
- **SkillPack** — a bundle of tools, delivered as a sidecar container with its own permissions.
- **SymposiumPolicy** — the rules an execution is admitted under. Chapter 14 is about these.
- **SymposiumSchedule** — a trigger that creates runs on a timer.

Here is `cluster2-agent`, the agent that operates the management cluster, as it sits in the repository:

```

apiVersion: symposium.ai/v1alpha1
kind: Agent
metadata:
  name: cluster2-agent
  namespace: symposium-system
spec:
  agents:
    default:
      model: qwen2.5:7b
      baseURL: http://host-ollama.symposium-system.svc.cluster.local:11434/v1
  authRefs:
    - provider: ollama
      secret: llm-credentials
  skills:
    - skillPackRef: web-endpoint    # serve it as an HTTP endpoint
    - skillPackRef: k8s-ops        # give its runs kubect1
  policyRef: sandbox-restricted
  memory:
    enabled: false
    systemPrompt: |
      You operate "cluster2", the management cluster of a KubeVirt
      multi-cluster platform. ...

```

Read it as a CxO would read a contract. The model is a 7-billion-parameter model served by Ollama on hardware inside the building, reached through an in-cluster address. The tools are named. The policy is named. The briefing is text anyone can review in a pull request. None of it is a setting in someone else's console.

The object the documentation described

The chapter would be dishonest without the correction that sits at the center of the repository's Symposium labs. The strategy document's own table of these nouns lists the persona object as a `SymposiumInstance`, and for months the repository kept its agents in exactly that form.

In July, while building the labs, the team listed the controllers actually running inside Symposium's controller manager: `Agent`, `AgentRun`, `Ensemble`, `MCPServer`, `SkillPack`, `SymposiumPolicy`, `SymposiumSchedule`. No `SymposiumInstance`. The resource definition existed and accepted objects, but nothing reconciled them. The labs index calls it *"the one correction that matters most"*: the objects the repository had been editing were *"inert documentation of intent."* The agents worked only because a separate `Agent` object with the same name happened to exist beside each one.

It got stranger. The resource definition existed on that cluster only because an older release had installed it; the current chart does not ship it at all. And typing the obvious command, `kubect1 get agent`, did not return Symposium agents — it resolved to a definition left behind by kagent, the

platform's previous agent framework, because two products had each defined something called *agent*. The labs have used the fully qualified name, `agents.symposium.ai`, ever since.

A name is not an identity. In Kubernetes, the API group is part of the name — and in a platform that has changed vendors, it is the part that matters.

One execution, one object

An `AgentRun` is where the idea becomes concrete, and lab 02 is built around one. On 17 September it was run again on the current versions — Symposium 0.10.75 on Kubernetes 1.37 — to check that its lessons still hold.

The first lesson is that the schema flatters you. It suggests that naming an agent and giving it a task is enough. Asked to validate exactly that, the cluster refused, just as the lab recorded two releases earlier:

```
The AgentRun "book-ch13-dryrun" is invalid:
* spec.agentId: Required value
* spec.model: Required value
* spec.sessionKey: Required value
```

That check was a server-side dry run; it created nothing. When a controller creates a run on your behalf — for a chat message, a schedule, an ensemble — it fills these fields in. When you write a run yourself, you supply them.

With the fields supplied, lab 02's run asks the agent to list the cluster's nodes and virtual machines and summarize their health. It finished in about thirty seconds, and everything worth knowing about it was recorded on the object and the pod it created:

Phase	Succeeded
Containers in the pod	agent , ipc-bridge , skill-k8s-ops
Identity it ran as	symposium-run-lab-agentrun-health
Tool calls	4
Tokens	7,418 in, 226 out
Model time	21.4 s

Its answer was a small table — nodes , KubeVirt VMs  — and a sentence: "*The cluster appears healthy with no issues reported for nodes or KubeVirt VirtualMachines.*" It was true: one node, `Ready` ; two virtual machines, `Running` . It also contained no detail the model could not have guessed, which is exactly the problem Chapter 15 takes up. For this chapter, the point is different. The record exists, and anyone with read access to the namespace can inspect it.

The lab's history holds the more instructive run. The first time it ran, back on release 0.10.38, the `skills` list was left out on the assumption that the run would inherit it from the agent. It did not. The pod came up with two containers instead of three, and the model's six attempts to run `kubectl` were written to a channel no tool was listening on:

```
04:54:05 tool_call [1]: execute_command args={"command":"kubectl get nodes; ..."}
04:54:05 Wrote exec request 1782968045241295119: kubectl get nodes; ...
... (6 tool calls, all silently dropped) ...
04:58:25 LLM call succeeded (tokens: in=11457 out=752, tool_calls=6)
```

After 262 seconds and 12,209 tokens, the model concluded: *"Given the persistent issues with command execution, let's assume there is a connectivity or availability problem with the skill sidecar."* The lab's verdict is one line: *"It was right."*

How a model gets hands

That failed run shows how tools actually reach a model, and it is worth understanding, because it decides what an agent can touch.

The model never runs a command itself. It emits a structured request — `execute_command`, with a command string. The `ipc-bridge` container writes that request to a shared channel. A separate sidecar container, supplied by a `SkillPack`, picks it up, runs `kubectl` with its own credentials, and writes the output back. Only then does the model see anything real.

So what an agent can do is not decided by what it is told. It is decided by which sidecar containers its run receives, and by the permissions those containers hold. Remove the sidecar and the model is talking into an empty room, however confident it sounds. Narrow the sidecar's permissions and the model's reach narrows with them.

That design paid for itself twice in this repository, and cost twice.

It paid when permissions turned out to be too narrow. The built-in `k8s-ops` tool pack could read nodes and pods but not virtual machines — its bundled permissions include no `kubevirt.io` group. The fix was an additive grant, reviewed in a file (`06-symposium/agent-kubevirt-rbac.yaml`), not an instruction to the model. The service-mesh agent needed the same for Istio's resources.

It cost when the rules for delivering sidecars changed between releases. On 0.10.38, runs created through Symposium's API received no tools at all, so the repository added a small admission webhook to attach the right tool pack to each agent's runs. On 0.10.47, the API began copying the agent's own skills list onto the runs it created. The agents listed only the serving skill at the time, so their runs now carried a non-empty skills list with no tool in it — and the webhook, which only filled *empty* lists, stood aside. Every tool call failed with *"the Kubernetes sidecar container is not currently running."* The fix was to list the tool packs on the `Agent` itself, which is why `cluster2-agent` names `k8s-ops` today. The webhook is still deployed, and now does nothing.

An agent that speaks the standard protocol

The first skill in `cluster2-agent`'s list, `web-endpoint`, is not a tool. It turns the agent into a service. The skill pack is marked as requiring a server, so listing it makes the controller create a long-running run in `server` mode and a Deployment named `cluster2-agent-web-endpoint-server`. The agent gets its own access key, stored in a Secret the platform generates, and an endpoint that answers the same requests as a commercial model API.

Asked what models it offers, on the current release, it answers:

```
{ "object": "list", "data": [ { "id": "qwen2.5:7b", "object": "model", "owned_by": "symposium" } ] }
```

The integration consequence is the one the lab emphasizes: *"anything that already speaks the OpenAI protocol ... can use a Symposium agent as a drop-in backend."* An existing client library, pointed at a different address, is now talking to a model running inside the building.

The lab also records the catch, and it is the right catch to know: *"You're talking to an agent, not a bare LLM."* Asked to count from one to five, the agent did not print five digits. It wrote them to a file called `/tmp/count.txt` in its sandbox and reported that it had done so. An endpoint that looks like a model API is attached to something that acts.

Behind the endpoint, every chat message becomes a short-lived `AgentRun` of its own. A conversation is not a stream of text that disappears. It is a series of objects.

Three windows onto the same objects

The platform's web UI has three modes, and each is a different window onto the same cluster.

SRE mode is the operator's dashboard: nodes, pods, virtual machines, events, the cluster-deployment flow, and a web terminal.

AI mode embeds Symposium's own console. The backend runs a small reverse proxy on a separate port, 8081, which hides the vendor's top bar and pre-loads the login token and the right namespace. Both details came from experience. The console's single-page app uses absolute paths that would collide with the UI's own routes, which is why it needs its own port. And the console defaults to a namespace where nothing lives, so without the pre-loaded setting its agent, run, and schedule panels simply appear empty. The proxy is also unauthenticated — acceptable for a demo platform, and recorded as such in its own comments.

Visual mode is the service-mesh view, and its command center still carries the platform's original chat panel. That panel talks to agents through the backend: a two-minute limit per request, a streaming reader sized for large responses, and each agent's key read from its Secret. It can also do one thing more than chat. If an agent's answer contains a structured proposal for an action, the backend turns it into something a person can act on, and it will only ever execute actions on an explicit allowlist. Today that allowlist has one entry:

```
// Keeping this tight: any new agent-driven action must be added explicitly.
var allowedActions = map[string]bool{
    "deploy_target_cluster": true,
}
```

One honest defect belongs here. The agent picker in that panel still asks the cluster for `SymposiumInstance` objects, which no longer exist, gets an error, and falls back to listing only the default agent. The service-mesh agent is running and serving; the picker cannot see it.

Why it matters that it is an object

Return to Anita's three questions, because the architecture answers each of them without a vendor in the room.

Where does the model run? On an Ollama server on the platform's own host, reached through an in-cluster address. No prompt, no tool output, and no answer leaves the boundary on the way.

What did it do last Tuesday? Every execution is an `AgentRun`: the task it was given, the answer it produced, its phase, its token and tool-call counts, its duration, the pod it ran in, and — since the upgrade described in Chapter 15 — a service account created for that run alone. Access to those records is governed by the same role-based permissions as every other object in the cluster, and they can be created by the same GitOps pipelines. There is no second audit system to buy, learn, or trust.

If we cancel, what do we keep? Everything. There is no separate service to cancel.

That is what *first-class object* buys: the agent inherits the governance machinery the organization already runs, instead of arriving with its own.

It does not buy immortality, and this is where the build log has to be exact. The strategy document's sixth constraint says the trail is append-only: *"No design may erase, short-circuit, or replace that trail."* Nothing enforces that. The labs delete their runs when they clean up. The repository's own repair script for serving agents works by deleting their runs so the controller will recreate them. And on 16 September, when both host clusters were deleted and rebuilt for the upgrade in Chapter 19, every run record that existed that morning — fifty-three of them, the oldest nearly four weeks old — was deleted with the cluster that held it.

An object is auditable for exactly as long as it exists. The same constraint already names the answer — durable storage off the cluster that *"mirrors `AgentRun`s, it does not replace them"* — and it has not been built.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Anita's approval for the agent pilot ran to four lines.

No external model. No vendor copilot. Every agent run stays a cluster object, and every run is copied off the cluster before anything is allowed to delete it. Access to the run records is reviewed like access to production.

Vikram read the third line twice. "The copy doesn't exist yet."

"Then the pilot starts when it does," she said. "An agent we can't list is an agent we can't answer for."

The ledger

- **Built:** three serving agents as `Agent` objects — `cluster2-agent`, `target-cluster-agent`, and `mesh-sre-agent` — each with a Deployment, an endpoint in the standard chat-completions format, and its own access key.
- **Measured on 17 September, current release:** one hand-written run — three containers, a per-run identity, 4 tool calls, 7,418 input and 226 output tokens, 21.4 seconds of model time.
- **Learned the hard way:** the persona object the documentation described had no controller; runs a controller creates for you and runs you write yourself get different defaults; and tools reach a run only through its skills list, whose delivery rules changed between releases.
- **Kept:** every run as a queryable object, governed by existing permissions — for exactly as long as nobody deletes it.
- **Lost:** fifty-three run records in the 16 September rebuild.
- **Still owed:** an off-cluster copy of the run records; authentication on the embedded console; an agent picker that can see the service-mesh agent; a strategy document that still names the old persona object.

Ask your team

1. Where does our AI model run, and what crosses our boundary when an agent answers a question? Ask for the network path, not the vendor's assurance.
2. Can we list every action an agent took last week — with the identity it used and what it cost — without asking a vendor? If not, who can?
3. What can our agents touch, and is that decided by a permission system or by a prompt? Only one of those is a control.

Open the repo

- `06-symposium/cluster2-agent.yaml` — an agent, declared: model, skills, policy, and briefing, with the history of skills delivery in its comments.
- `06-symposium/labs/01-chat-serving/` — the standard chat endpoint, including a 30-line client.
- `06-symposium/labs/02-agentrun/` — one execution as an object: the required fields, and the run whose tools never arrived.
- `kubectl get agents.symposium.ai,agentruns.symposium.ai -n symposium-system` — the live records, by their fully qualified names.

- `ui/backend/handlers/ai.go` (the chat proxy and its action allowlist) and `ui/backend/handlers/dashboard_proxy.go` (the embedded console).
- `06-symposium/labs/README.md` — "the one correction that matters most."
- **Without the repo:** Appendix G.6 quotes the `cluster2-agent` Agent object whole, briefing included.

Chapter 14 — Rules the Model Cannot Break

If refusal must be guaranteed, enforce it below the model ... not in the prompt.

— `CLAUDE.md`, the `mesh-sre-agent` section

MERIDIAN · A THURSDAY IN AUGUST · THE CTO'S OFFICE

Anita Rao had the strategy document open at a section called *Architecture Constraints*. Eight numbered rules, each one sensible.

"Local model first. Pinned endpoint. Address discipline. An egress allowlist. Keep the model warm. Append-only audit. Tenant isolation. A sandbox for anything that takes free-form chat." She scrolled back to the top. "These are good rules."

"They're the contract," Vikram said. "Every change to the agent layer is supposed to satisfy them."

"Then I want one more column." She turned the screen toward him. "Next to each rule, write what happens if the agent ignores it."

"The agent doesn't decide most of these."

"Then write down who does. A person. A script. The API server. The permission system. The network." She paused. "And if the honest answer is *we told the model*, write that down too."

"That's not going to be a flattering column."

"It isn't supposed to be," Anita said. "I'd rather read it from you than from an incident report."

* * *

BUILD LOG · JULY–SEPTEMBER 2026

Where a rule can live

A rule for an AI agent can live in several places, and they are not equally strong. From weakest to strongest:

1. **The prompt.** The model is asked.
2. **A document.** People are asked.
3. **A script.** The right thing happens, if someone runs it.
4. **Admission.** The API server refuses to create the object.
5. **Permissions.** The credentials cannot do it.
6. **The network.** The packets never arrive.
7. **Absence.** The tool, the credential, or the route does not exist.

A model can ignore the first. People can ignore the second and third. The rest do not care what anyone intended. This chapter is the column Anita asked for: each of the platform's eight rules, and the layer that actually enforces it, measured on the running platform rather than read from its documentation.

Most of the measurements were taken on 17 September 2026, on Symposium 0.10.75 and Kubernetes 1.37, specifically for this chapter. Where they contradict what the repository says, the chapter says so.

The rule that lived in the prompt

Start at the bottom of the ladder, because the repository has already measured what happens there.

The `mesh-sre-agent` from Chapter 15 was deliberately given no path to metrics — no route to Prometheus or Kiali. Its briefing opens with a rule meant to stop it from pretending otherwise:

"RULE 1 — METRICS ARE UNAVAILABLE. ... If asked for latency, p50/p95, request rates, error rates, throughput or saturation, do NOT run a command and do NOT estimate: reply exactly 'Metrics are not reachable from this agent ...' Never invent a number."

Measured on `qwen2.5:7b` across repeated attempts, the repository records that it *"does not reliably comply: it ignores 'do not run a command', and in one run answered 'we can make an educated guess based on the availability of related metrics'."* It never actually produced an invented number in testing. But moving the rule to the very top of the prompt did not fix the compliance, and the conclusion written into the repository is the sentence this chapter takes as its epigraph: *"the guard is a nudge, not a control."*

The same repository gives every agent the same closing instruction in its briefing: *"Read-only unless the user explicitly approves a change."* Hold on to that sentence. It comes back.

What held

Some rules genuinely live below the model, and they held when tested.

Admission refused a self-granted exception. The platform's policy for `cluster2-agent` denies a web-fetching tool called `fetch_url`. Lab 04 tries to create a run that explicitly re-allows it. On 17 September, asked to validate that run without creating it, the API server refused, exactly as the lab recorded in July:

```
Error from server (Forbidden): admission webhook "vagentpod.symposium.ai"
denied the request: tool "fetch_url" is denied by policy
```

That webhook is configured to *fail closed*: if it cannot be reached, the request is rejected rather than waved through. Its sibling, which validates tool packs, is configured the same way, and on 16 September it showed what that means in practice: while the platform was being reinstalled and the

webhook was not yet serving, the install was refused until it came up. An inconvenience during setup; the right behavior for a guard.

Permissions kept agents read-only outside their own namespace. An agent run's identity was checked directly. Across the cluster, it cannot delete deployments or virtual machines, and it cannot create workloads in other namespaces. The extra grants the repository adds, for virtual machines and for the service mesh, are read-only by design and say so in their headers.

The network closed the metrics path. This one needed a test, because the lab written in July said the opposite (more on that below). Two short-lived pods were started in the agents' namespace: one carrying the same labels as a real agent run — `sympozium.ai/role=agent` and `app.kubernetes.io/part-of=sympozium`, confirmed against a live run pod the same morning — and an identical control pod with neither label. Both tried the same four destinations:

Destination	Agent-labeled pod	Unlabeled control
Ollama, through the in-cluster Service (port 11434)	200	200
Kiali on <code>cluster1</code> (port 20001)	no connection	302
<code>example.com</code> over HTTP (port 80)	no connection	200
<code>example.com</code> over HTTPS (port 443)	200	200

The control pod reached everything; the agent-labeled pod could not reach Kiali or plain HTTP. Network policy is enforced on this cluster, and the metrics endpoints really are out of an agent's reach. Whatever the model says about latency, it did not read it from Prometheus.

The UI will only execute one action. However an agent phrases a proposal, the web backend acts only on an allowlist written in code, and that list has one entry: deploying a target cluster. A model cannot talk its way onto that list.

Those four are real controls. Now look at the last row of the network table again.

What did not hold

HTTPS goes anywhere. The agent-labeled pod reached `example.com` on port 443. The reason is in the repository's own network policy, `sympozium-allow-ollama`. It was written to open two paths agents genuinely need — Ollama, and the Kubernetes API — and it opens them *by port*, with no destination:

```
to= ports=[{"port":11434,"protocol":"TCP"}]
to= ports=[{"port":443,"protocol":"TCP"}, {"port":6443,"protocol":"TCP"}]
```

Port 443 was opened for the Kubernetes API, and a port number cannot tell the Kubernetes API from the rest of the internet.

A denied tool ran anyway. Lab 04's other run creates no exception at all. It simply asks the agent to fetch `example.com` with the tool its policy denies. Re-run on 17 September, the agent made one tool call, and the answer described the real page's contents. The deny rule only blocks a run that *explicitly asks* to re-allow the tool; a run that says nothing about tools — which is how nearly every run is written — can use it freely. Put that together with open HTTPS and the result is straightforward: an agent on this platform can send a request to an arbitrary web server.

Notice the irony in the epigraph's source. The same note that says to enforce refusal below the model names policy tool gating as one way to do it. On this version, tool gating is the control that only works when nobody needs it.

Inside their own namespace, agents can write. The built-in `k8s-ops` tool pack grants two sets of permissions. Across the cluster, they are read-only. Inside the namespace where the run executes, they are not: `create`, `update`, `patch`, and `delete` on pods, deployments, jobs, network policies, roles, role bindings — and secrets. Checked against a real run identity from a chat with `cluster2-agent`:

```
delete pods -n symposium-system      yes
get secrets -n symposium-system      yes
```

That namespace is where the agent platform itself lives. It holds every agent's API key, the model credentials, and one more Secret worth naming: `target-cluster-kubeconfig`, whose user is `target-cluster-admin`, for the target cluster's API on port 6443 — which the network policy above leaves open. A `cluster2-agent` run's identity is permitted to read it.

`cluster2-agent`'s briefing says its tools talk to its own cluster only, and every briefing says *read-only unless the user explicitly approves a change*. Outside the namespace, permissions make that true. Inside it, the only thing standing between a run and those writes and those credentials is the sentence in the prompt — the layer this chapter began by measuring.

The scorecard

Here is the column, rule by rule.

#	Constraint	What actually enforces it today	Verdict
1	Local-LLM-first	The agent's configuration; not the network, which allows HTTPS anywhere	Partly
2	Pinned model endpoint	An in-cluster Service defined in the repository	Held — at a different address than the rule names
3	Address-pool discipline	MetalLB within each cluster; only a table and a warning between the two	Partly
4	Egress allowlist	Network policy, by port rather than destination	Partly
5	Warmth is mandatory	A setting on the host, outside the repository	Partly, and by accident
6	Audit is append-only	Nothing	Not held
7	Tenant isolation	Nothing inside the shared namespace	Not held
8	Sandbox free-form chat	A policy reference that never applies	Not held

A few rows need their evidence.

Rule 2 names the Ollama address `172.18.0.1` and warns against a hostname that does not resolve on this kind of cluster. The agents no longer dial that address at all. The repository records that agent pods *"couldn't dial that IP directly under kindnet's NetworkPolicy enforcement,"* so it added an in-cluster Service whose endpoint points at the host. The rule's intent — one fixed, local model endpoint — held. Its text is out of date.

Rule 3 requires every load-balancer address change to land in three places at once. Within a single cluster, MetalLB itself refuses to hand the same address to two services. But `cluster1` and `cluster2` advertise on the same network segment, and nothing stops their pools from overlapping except a table in the repository and a warning never to run one cluster's install script on the other. The pool the rule describes, `.211-.220`, has already been widened to `.211-.225`.

Rule 5 says no persona enters the fleet without a warm-path guarantee, and names the mechanism: a heartbeat schedule. The repository ships one, set to run every four minutes. On 17 September its status recorded two runs in total, the last at 07:44 UTC the previous day, with a "next run" time that had passed seventeen hours earlier. The schedule is not firing. The model stays warm anyway, because the host's Ollama service runs with `OLLAMA_KEEP_ALIVE=-1` and never unloads a model once it is loaded. That setting lives in the host's service configuration, not in the repository (whose Ollama notes suggest 24 hours). It keeps only models that have already been loaded: on 17 September, `qwen2.5:7b` was resident, and `llama3.2`, the target-cluster agent's model, was not.

Rule 6 is Chapter 13's finding: the audit trail is deletable, was deleted by routine scripts, and vanished entirely in the rebuild.

Rule 7 says one namespace per tenant and one identity per agent, with no cross-tenant bindings. Every agent shares one namespace. Identities are now per run — better than the rule asked — but Chapter 15's fix bound mesh and VM read access to every service account in that namespace, and the write permissions above let a run of any agent equipped with `k8s-ops` — `cluster2-agent` or `mesh-sre-agent` — read every other agent's key.

Rule 8 puts free-form chat under a restricted network policy allowing DNS, the Ollama host, and the Kubernetes API only. `cluster2-agent` does reference that policy. But the policy's network rules are scoped to pods labeled as sandboxed, and every agent in the repository runs with its sandbox disabled, so no agent pod carries the label. On 17 September, a real run pod carried `symposium.ai/role=agent` and no sandbox label at all. The restricted rules exist; they select nothing.

A correction

Lab 04, written in July, concluded that network policy was *"inert on this platform"* for two reasons. The first — the missing sandbox label — is still true. The second was that Kind's default network plugin *"has no policy engine"* and never enforces policy. On 17 September, on the rebuilt cluster, that was false: the probe pods above show enforcement plainly, and the repository's own shim Service for Ollama exists *because* a policy blocked a direct connection.

Whether the lab's second finding was wrong in July, or became wrong later, the record cannot say — the July cluster no longer exists. What it shows is the chapter's thesis turned on its authors. "Network policy is not enforced" was a claim about a control, made without a test that could have failed. The test takes two pods and a minute.

Rules that cannot move below the model

Not every rule can be pushed down the ladder, and it is worth being clear about which.

"Never invent a number" is one. The network can guarantee that the agent cannot read real metrics — and here it does. It cannot stop a language model from writing a plausible latency figure out of nothing. No permission, admission rule, or firewall reaches inside the text a model generates. For rules like that, the control is not enforcement but verification: the ladder of evidence in Chapter 15, and a process that treats an agent's answer as a claim.

Everything else in this chapter's failures *can* move down, and the moves are unglamorous:

- Replace port-only egress with destinations: the API server's address, the target cluster's address, the model host.
- Run agents somewhere their credentials cannot read the platform's own secrets, or strip the tool pack's write permissions inside that namespace.
- Enable the sandbox, so the restricted policy actually selects agent pods.
- Replace the heartbeat that stopped with something whose failure someone notices, and write the host's keep-alive setting into the repository.

- Copy run records off the cluster before anything can delete them.

None of these needs a better model. All of them need someone to decide that a rule is not real until something other than the model enforces it.

* * *

MERIDIAN · TWO WEEKS LATER

Vikram brought the column back on one page, with the verdicts in the right-hand margin: one rule held, four held partly, three did not.

Anita read the page twice, then asked the question she always asked first. "Which of these lets an agent do something we can't undo?"

"Seven. Inside its own namespace, an agent's run can write, and it can read every secret there — including an administrator credential for the target cluster."

"And what stops it today?"

"A sentence in its briefing."

She drew a box around row seven. "That one first. Then the network rule. Then the sandbox." She slid the page back. "And strike *Architecture Constraints* off the top of the document."

"And call it what?"

"*Intentions*," she said, "until each one has something under it that the model can't argue with. If the model can break it, it isn't a rule. It's a request."

The ledger

- **Tested on 17 September, current release:** network enforcement with two probe pods; both of lab 04's runs; the real permissions of an agent run's identity; the warm-up schedule's status; the resident models; how the admission webhooks fail.
- **Held below the model:** admission refused an explicit tool override, and fails closed; permissions kept agents read-only outside their namespace; the network blocked metrics endpoints and plain HTTP; the UI backend executes one allowlisted action.
- **Did not hold:** a denied tool was used anyway; HTTPS reached an arbitrary host; agent runs can write in their own namespace and read every secret there, including a target-cluster administrator credential; the chat sandbox's policy selects no pods; the warm-up schedule stopped firing seventeen hours before the check.
- **Held by accident:** model warmth, kept by an unrecorded host setting rather than the mechanism the rule names.
- **Corrected:** lab 04's claim that network policy is inert on this cluster.
- **Cannot be enforced below the model:** "never invent a number." That one needs verification, not a control.

Ask your team

1. For each rule we have given our AI agents, what enforces it if the model ignores it? Count how many answers are "the prompt."
2. What can an agent's credentials read, and which of those secrets unlock something bigger than the agent?
3. When a guardrail fails — a webhook down, a policy matching nothing — does it fail open or closed, and would anyone notice?

Open the repo

- `Sovereign_Cloud_Agentic_Strategy.md` — "Architecture Constraints", and §4 for the reasoning behind them.
- `06-symposium/labs/04-policy/` — both runs; re-run `run-allowed.yaml` with `--dry-run=server` to see the refusal without creating anything.
- `06-symposium/policies/sandbox-restricted.yaml` beside `06-symposium/agent-egress-ollama.yaml` — the destination-specific allowlist that never applies, and the port-only rule that does.
- `kubectl auth can-i --list -n symposium-system --as=system:serviceaccount:symposium-system:symposium-run-<run-name>` — what a real agent run can actually do.
- `06-symposium/mesh-sre-agent.yaml` — RULE 1, and the header explaining why the metrics path is closed.
- `ui/backend/handlers/ai.go` — `allowedActions`.
- **Without the repo:** Appendix G.5 quotes `agent-egress-ollama.yaml` whole — the port-only rules are visible in it.

Chapter 15 — Trust, but Verify

An LLM's answer is a claim, not evidence.

— a comment in `06-symposium/demo-mesh-sre.sh`

MERIDIAN · A THURSDAY IN SEPTEMBER · THE CTO'S OFFICE

Vikram had printed two answers and folded the top of each page over, so that the name of the agent that wrote it was hidden.

"Same question to two agents," he said. "Two different clusters. List the nodes, and report exactly what the command printed."

Anita read the first.

The ``kubectl get nodes -o name`` command printed:

```
node/cluster2-control-plane
```

This indicates that there is one node named ``cluster2-control-plane`` in the cluster.

Then the second.

After retrying, I was able to run the ``kubectl get nodes -o name`` command. Here is the output:

```
target-node-1
```

```
target-node-2
```

This is the list of nodes in the target cluster.

"One of those is invented," Vikram said.

Anita did not take long. She tapped the second page. "Not this one. It says it retried. It's being careful."

"That's the invented one. There is no `target-node-1`. There never was."

She looked at the page again, the way she had looked at the stopwatch printout a few weeks earlier.

"Operations wants an agent on the pager by October," she said. "When it wakes someone at three in the morning and tells them which node is failing, how do they know it's true?"

"From the answer? They don't. The answer is the one part of it we can't use as evidence."

"Then what is?"

"Everything around the answer."

Two answers, one of them invented

Both answers were collected on the same day, during the platform rebuild described in Chapter 19, while checking that every agent could still reach what it was meant to reach. The prompt was identical — *"Use your kubectl tool: kubectl get nodes -o name. Report exactly what it printed."* — sent to two agents with two different jobs.

The first came from `cluster2-agent`, which operates the management cluster. It was correct: that cluster has exactly one node, `cluster2-control-plane`.

The second came from `target-cluster-agent`, whose tools reach inside the k3s cluster that runs as virtual machines. That cluster had two nodes, and their real names were `target-cluster-cp-9x5sr` and `target-cluster-workers-x79t2-tlzn9`. The agent got the count right and every name wrong. Its answer also described a process — *"After retrying, I was able to run..."* — for which there is no evidence at all.

The first thing checked was not the model. It was the path the model's tool would have used. This agent reaches the target cluster through credentials stored as a Kubernetes Secret, and those credentials go stale whenever the target cluster is rebuilt. The repository already warns about exactly this, because when it happens the tool's TLS handshake fails *silently* and the model is left to fill the gap. So the certificate authority recorded in the Secret was compared with the one in the live kubeconfig. They matched. Then a throwaway pod, with no model anywhere near it, mounted the same Secret and ran the same command:

NAME	STATUS	ROLES	AGE	VERSION
target-cluster-cp-9x5sr	Ready	control-plane	33m	v1.37.0+k3s1
target-cluster-workers-x79t2-tlzn9	Ready	<none>	33m	v1.37.0+k3s1

The path worked, and the credentials were good. The invented answer was not caused by any plumbing that could be checked from outside. Whether the model called its tool at all in that run was never established.

What is on record is a comment in the agent's own manifest, from an earlier day: *"Verified: a chat run returned both real target node names, matching a kubeconfig-mounting probe pod byte for byte."* Same agent, same kind of question. Truth one day, invention another.

Earlier on 16 September, `cluster2-agent` — the one that got it right — had been asked a slightly different question: run `kubectl get nodes` and *"report the node name back to me in one short sentence."* Its entire reply was `ronics-01`. There is no such node anywhere on the platform. Asked again, this time for the exact output, it answered correctly.

One pair of runs proves nothing about phrasing. It does prove something about confidence: a wrong answer and a right one arrive in the same voice.

Where the truth goes missing

A Symposium agent does not run `kubectl` itself. The model asks a sidecar container to run the command, the sidecar returns the output, and the model writes its answer from that output. That gives the truth three places to go missing, and from the outside, all three look like an answer.

- **The tool never runs.** On the older 0.10.38 release, chat runs were created with no tool sidecar at all, so the model had nothing to call. Its reply in that state — *"It seems there might be an issue with the skill sidecar..."* — was the model giving up, not a real sidecar crash, and it took an investigation to tell the difference.
- **The tool runs and fails quietly.** Stale credentials make the TLS handshake fail without a clear error, and the repository records what the smaller model did next: *"llama3.2 then confabulates plausible namespace output."* That is why `mesh-sre-agent` runs `qwen2.5:7b` instead, and the repository says so in as many words: not llama3.2, because *"it confabulates kubectl output when a tool call fails silently."*
- **The tool runs, and the model mishandles the output.** The demo script lists the failure modes it has seen on this exact setup: empty replies, and tool calls leaked into the answer as raw text instead of being executed. It notes a quieter one too — the model *"often mistakes tool output for user-supplied input and answers 'it looks like you provided...'"* — and puts the underlying reliability bluntly: *"tool-call reliability on a 7B model is roughly a coin flip."*

None of this is lying. A language model's job is to produce the most plausible continuation of what it has been given. When the evidence does not arrive, a plausible cluster is what it produces. The failure is not in the model's character. It is in any process that treats the model's output as the evidence.

* * *

BUILD LOG · AUGUST 2026

Seven scenes, no edits

In August the team built a demo for `mesh-sre-agent`, an agent that answers one question — *is the service mesh healthy?* — from live cluster state. Demos of AI agents are often the least trustworthy artifacts in a project: recorded until the model behaves, cut to the good take. This one was written to be the opposite, and its rules sit in the header of `06-symposium/demo-mesh-sre.sh`:

"Nothing here is scripted output: every kubectl result and every agent answer is captured live at run time. If a call fails, the failure is what gets recorded — do not hand-edit the transcript to make a nicer demo."

The script does allow retries, and the rule for them is the most important design decision in the file. A question is asked again only when the reply is *structurally* broken: empty, a transport error, or a tool call leaked as raw text. A reply that is well-formed but unhelpful is kept.

"This is a retry, NOT a filter for answers we dislike... A well-formed but unhelpful answer is kept as-is — the demo shows what the agent actually said."

The seven scenes run in a deliberate order: what the agent is; that it is an ordinary object in the cluster; the permissions gap that had to be closed before it could see the mesh at all; two live questions; a scene titled *Trust, but verify* that puts ground truth next to the agent's answer; and a closing scene on what the agent is *not* — including the admission that its instruction to refuse questions about metrics is "*a nudge and not a control*."

The recording from 29 August is in the repository as a 51-second video. Here is what it actually shows.

In scene 4, the agent was asked whether the mesh's control plane was healthy. It answered:

It seems there was an issue with the `kubectl` command. Let's try the command again to ensure we are using it correctly. Could you please provide me with the specific `kubectl` command you were trying to execute? If you're not sure, I can provide a common example based on typical usage.

That is a failure, on camera, in the showcase recording — and it stayed in, because it was well-formed. The script's rule held against the script's own author.

In scene 5, the agent was told to run one specific command and show its output. It answered:

It seems like the response you provided is related to Istio East-West Gateway, which appears to be functioning since the "PROGRAMMED" column shows "True" and it has been operational for 8 days. Is there a specific action or information you need regarding this gateway? ...

There is the quirk the script's comments warned about — "*the response you provided*" — the model treating its own tool output as something the user had pasted in. Scene 6 then put the ground truth beside it, fetched straight from the Kubernetes API:

NAMESPACE	NAME	CLASS	ADDRESS	PROGRAMMED	AGE
istio-system	istio-eastwest	istio-east-west	172.18.255.221	True	8d

Did the tool really run? Scene 6's narration tells the viewer to check "*the pod-name suffixes and the gateway address*," because those are unguessable. But the answer in scene 5 contains neither. The gateway's name does not count as proof either: the agent's briefing names `istio-eastwest` outright, so the model could have produced it from its instructions without looking at the cluster. The one detail in the answer that could only have come from the live API is the age — *8 days*.

That is enough. But the proof sat in a different detail from the one the narration, written before the run, said to look for.

The honest scorecard of the showcase demo, then, is one failed answer and one correct answer with a quirk. Roughly a coin flip — just as the script's own comments predicted.

An unguessable detail proves that a tool ran. "Unguessable" has to be judged against what the agent was told, not against what looks obscure.

Scoring the work

A demo shows a sample of two. The labs that close the repository's Symposium learning series were built to answer the question a demo cannot: *does the agent do useful work, and how reliably?*

Lab 10 plants a fault. It deploys a workload whose container image tag does not exist — `nginx:this-tag-does-not-exist-9z9z9` — so that its pods sit in `ImagePullBackOff`, with the reason written plainly in their events. Then it gives `cluster2-agent` only a symptom: a workload in the namespace is unhealthy; find it, determine the root cause from the actual cluster state, propose a fix. *"Do not guess — inspect the cluster."*

The scoring is deliberately strict and entirely mechanical. A run passes only if its answer names the broken workload *and* the image-pull cause. An agent that says *"a pod is unhealthy, try restarting it"* without identifying the image *"is a fail — that's the difference between useful triage and noise."*

Three runs, on `qwen2.5:7b`:

Run	Verdict	Tool calls	Input tokens	Output tokens	Seconds
1	PASS	1	6,228	416	14
2	PASS	2	7,136	837	29
3	PASS	4	10,390	520	25

Three for three. And the cost of one investigation, in the lab's own words: *"~6–10k input tokens and 14–29s — the number to multiply by your alert volume when sizing this."* On this platform, that multiplication does not land on an invoice. The model runs on hardware the organization owns, so tokens are paid for in CPU and GPU time on that hardware. The arithmetic still decides how much hardware to buy.

The table rewards a second look. The same task, on the same model, took one tool call on the first run and four on the third. Run 3 dug deepest and named the exact image tag that does not exist. Run 1 passed with a single call; the part of its answer the lab recorded names the pod and its `ErrImagePull` status, and infers from that status that the image could not be pulled. That is true — but the lab's own reading guide treats two or more calls, a `get` followed by a `describe` or `events`, as the sign of a real investigation. The scorer is a pattern match, and it gives both runs the same PASS.

The scorer verifies the agent. Nothing yet verifies the scorer.

The success that looked at nothing

Lab 09 asks a different question: *which model?* It runs the same three real tasks — count the pods in a namespace, report node readiness, list the virtual machines — on two local models, with the model as the only variable. Every figure but one comes from the platform's own record of each run; only tokens per second is derived.

Model	Runs succeeded	Avg tool calls	Avg duration	Avg tokens/s (whole run)	Avg tokens
llama3.2	3 / 3	0.7	22.4 s	4.0	4,146
qwen2.5:7b	3 / 3	1.3	15.1 s	11.0	7,355

On the summary line, both models are perfect: three runs, three successes. The per-run table underneath says otherwise. One of llama3.2's runs — the one about nodes — is marked `Succeeded` with **zero tool calls**. It answered a question about the live cluster without running a single command. The lab's README calls it what it is: *"it answered the node question without running kubectl, so that 'success' is untrustworthy."* The same prompt on `qwen2.5:7b` made one tool call.

`Succeeded`, it turns out, meant *the run finished*. It did not mean *the answer is true*, and nothing in that field could have said so.

The fix costs almost nothing, which is what makes it worth a paragraph. The tool-call count is recorded on every run by the platform, not by the model. An alarm on *succeeded with zero tool calls* is one query. It would not catch every invention — the node list at the start of this chapter might have come from a run that called its tool and ignored the output — but it catches the cheapest kind for free.

The rest of the table is a trade-off, not a verdict. `qwen2.5:7b` used its tools more, finished sooner, and produced output about 2.7 times as fast over the whole run — and consumed about 77% more tokens doing it. More rigor, more cost. That is a decision to make per workload, with the numbers in front of you.

* * *

BUILD LOG · 16 SEPTEMBER 2026

The day the check was wrong

Everything so far has been about checking the model. This section is about checking the check.

On 16 September, the agent platform itself — Symposium — was upgraded from chart version 0.10.57 to 0.10.75. The repository's standing instruction for confirming that `mesh-sre-agent` could still see the mesh was a single deterministic command, with no model involved, run as the agent's service account:

```
kubect1 get gateways.gateway.networking.k8s.io -A \
--as=system:serviceaccount:symposium-system:symposium-agent
```

It returned the east-west gateway. Access confirmed.

Then the agent itself was asked to list the gateways. It answered, in part:

```
It appears that the service account 'sympozium-run-mesh-sre-agent-web-x76fk' does not have the
necessary permissions to list resources of type 'gateways' in the API group
'gateway.networking.k8s.io' at the cluster scope. This error suggests a permissions issue.
```

This time the model was right and the deterministic check was wrong. The agent had reported a real error, accurately — and in doing so it named an account the check had never tested.

The explanation came from reading a live run's pod rather than any document. Somewhere between chart versions 0.10.58 and 0.10.75, Symposium stopped running every agent task under one shared service account. Each run now gets its own, named `sympozium-run-<run-name>`, created with the run, given its own narrowly scoped role, and deleted when the run finishes. The repository's extra permissions for reading the mesh had been granted to the old shared account by name, so every new run arrived under a name that grant had never heard of. Meanwhile the verification command kept asking about the shared account, which still had its access, and kept answering yes.

The check had been right when it was written. A week earlier, after the previous upgrade, the repository recorded that the account had been *"verified via the actual `spec.serviceAccountName` on live pods before and after."* A verification is true on the day it runs.

The fix rebound both permission grants — for the mesh and for the virtual machines — from one account name to the group Kubernetes automatically assigns to every service account in the namespace, so each new per-run account inherits read access the moment it exists. It was then verified the way it should have been verified in the first place: by reading the account name off a live run's pod, and checking *as that account*.

```
kubectl get gateways.gateway.networking.k8s.io -A \
--as=system:serviceaccount:sympozium-system:sympozium-run-mesh-sre-agent-web-twt55
```

Then the agent was asked again:

```
It appears that you have an Istio East-West Gateway service running in the 'istio-system' namespace
with the name 'istio-eastwest'. The service is accessible at '172.18.255.221' and has been
configured to be active since 26 days.
```

The address is not in the agent's briefing, and the age — twenty-six days — matched the cluster exactly. The tool ran.

The fix has a cost, and it belongs in the ledger. Per-run accounts are a *better* security design than one shared account: each run holds only what that run needs. Binding the namespace's whole service-account group restored the agent's reach, and at the same time gave read access to the mesh and virtual-machine configuration to every service account in that namespace, agent or not. It is read-only, and this is a demo platform. It is still a trade someone should make on purpose rather than inherit from an upgrade-day fix.

One more loose end, recorded because this chapter would be hypocritical without it. The demo script's own verification scenes still check as `sympozium-agent`. They pass, because the group binding covers that name too, so nothing on screen would look wrong. But they no longer test the identity the agent actually runs as. For now, the scene called *Trust, but verify* is verifying the wrong thing.

A ladder of evidence

Taken together, this chapter's incidents sort into five levels of evidence, from weakest to strongest. Each rung costs more than the one below it, and each catches something the rung below cannot.

1. **The answer.** A claim, nothing more. Tone carries no information: of the two node lists, the invented one sounded more careful.
2. **The run's own record.** Did the agent call a tool at all? A success with zero tool calls is a red flag, and the platform records it for free.
3. **A detail it could not have guessed.** An age, an address, a generated suffix — judged against what the agent was *told*. The gateway's name was in its briefing; its age was not.
4. **Independent ground truth, as the identity the agent really runs as.** The same query, run by something that is not the model, using an account name read off a live pod — not remembered from a runbook.
5. **A score against a known answer, repeated, with the cost recorded.** A planted fault, a strict pass rule, several runs, tokens and seconds for each — and then a look at what the scorer cannot see.

A typical demo stops at the first rung. Rungs 2 and 4 can run automatically on every answer that matters. Rung 5 is how you decide whether to deploy an agent at all, and how you notice when a model update, a prompt change, or a platform upgrade quietly makes it worse.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Anita had written the rules for the October pilot herself, on one page.

The agent may triage alerts. It may not act on them. Every answer it gives reaches the on-call engineer with the raw tool output attached underneath — not summarized, attached. A run that reports success with no tool calls pages a human with a warning instead of an answer. Once a week, the triage scorecard runs against a planted fault, and its pass rate goes on the operations dashboard next to the cost of one investigation. And the access check reads the agent's identity off a live pod, every time it runs.

Vikram read it twice. "That's five rules for one agent."

"It's one rule," Anita said, "written out five ways."

She took the page back and wrote a line across the top.

No receipt, no claim.

The ledger

- **Measured:** 3 of 3 correct triages of a planted fault on `qwen2.5:7b`, at 6,228–10,390 input tokens and 14–29 seconds each, using anywhere from 1 to 4 tool calls for the same task.
- **Compared:** `qwen2.5:7b` against `llama3.2` on the same three tasks — 1.3 against 0.7 tool calls per run, 15.1 against 22.4 seconds, 11.0 against 4.0 tokens per second over the whole run, 7,355 against 4,146 tokens.
- **Caught:** a run marked `Succeeded` that never ran a tool; two invented node lists; a showcase demo answer asking the user which command they meant; a verification command testing an identity the agent had stopped using.
- **Fixed:** the agent's read access, rebound from one account name to the namespace's service-account group — restoring its reach, and widening who else holds that read access.
- **Still owed:** a demo whose own verification scenes check the old identity; a scorer that cannot tell a one-call pass from a four-call pass; lab figures captured on an older release and not yet re-run; and samples of three.

Ask your team

1. When an agent reports a fact, what evidence travels with it? Ask to see the tool output next to the answer, not the answer on its own.
2. How would we know that a "successful" agent run looked at anything at all? Is there an alarm on zero tool calls?
3. When did we last check an agent's access as the identity it runs as today? Not the one written in the runbook.

Open the repo

- `06-symposium/demo-mesh-sre.sh` — the seven scenes and their retry rule. `SCENE_DIR=<dir>` writes a transcript; `scripts/render-demo-video.sh <dir> <out.mp4>` renders it. The August recording is `docs/demo/mesh-sre-agent-demo.mp4`.
- `bash 06-symposium/labs/10-sre-triage/run-triage.sh` — the scored triage. Add `LAB_MODEL=llama3.2` to run it on the other model.
- `bash 06-symposium/labs/09-metrics/run-metrics.sh` — the model comparison, read from each run's `status.tokenUsage`.
- `kubectl get pod <run-pod> -n symposium-system -o jsonpath='{.spec.serviceAccountName}'` — the identity to verify against. `06-symposium/agent-istio-rbac.yaml` explains the group binding and its history.
- `CLAUDE.md`, the `target-cluster-agent` pitfall — the probe pod that mounts the agent's kubeconfig Secret and runs `kubectl` with no model in the loop.

Chapter 16 — The Fleet

Execute Phase 0–1 well and you have a credible demo. Execute Phase 2–3 well and you have a product.

— `Sovereign_Cloud_Agentic_Strategy.md`, §7

MERIDIAN · A TUESDAY IN SEPTEMBER · THE CTO'S OFFICE

The catalog ran to eleven pages, and Anita Rao had read it the way she read a hiring plan: job titles down the left margin, and a question mark next to most of them.

VM Lifecycle Warden. Golden Image Curator. MetalLB IP-Space Accountant. Workload Onboarder. Chaos & Resilience Tester. RBAC Auditor. Idle VM Reaper. Right-Sizer. Incident Postmortem Drafter.

"Twenty-nine," she said. "Twenty-nine agents."

"Grouped by where they work," Vikram said. "Three planes — the management cluster, the tenant clusters, the mesh between them — and four shared teams: security, cost, developer experience, observability. The roadmap says most of them are a day or two of configuration and prompt writing each."

"Nobody I hire is productive in a day or two." She turned to the last page. "How many of these work today?"

"Three. Sort of."

"Sort of?"

"They answer questions. None of them acts on anything."

Anita capped her pen. "Then before we talk about twenty-nine, tell me what it would take for one of them to be allowed to act. Not what it does. What it has to prove before it gets the job."

* * *

BUILD LOG · APRIL–SEPTEMBER 2026

Twenty-nine job titles

The strategy document organizes its agents the way the platform is organized. An *infrastructure plane* on the management cluster: virtual machines, disk images, cluster provisioning, load-balancer addresses. A *tenant plane* inside the target clusters: workloads and their health. A *mesh plane* across clusters: service discovery, mutual TLS, traffic. Over all three sits a *control tower* — the UI's AI mode, drawn as "*Orchestrator + chat surface + approval gate*" — and beneath them a shared layer for security and compliance, cost, developer experience, and observability.

Counted section by section, the catalog lists twenty-nine personas: five infrastructure agents, four tenant agents, three mesh agents, five for security, four for cost, five for developer experience, and three for observability. Each entry reads like a job description — what it does, which tools it needs,

what triggers it, a sample task, and the value it delivers. The first one sets the tone:

"Find all VMIs not in `Running` phase. For each, check: `DataVolume` ready? CPU/memory pressure on node? `virt-launcher` pod scheduled? Report root cause and, if safe and idle > 10 min, restart the VMI."

Read the catalog for its verbs and a pattern appears. The VM warden *restarts*. The cluster provisioner *drives the rollout*. The idle-VM reaper *proposes a stop*. The onboarder deploys; the `ServiceEntry` generator generates; the CVE watcher *opens issues*. The value the catalog promises comes overwhelmingly from agents that change things.

Two sentences set the catalog's technical assumptions. The first: "*Every agent is one `SymposiumInstance` ; every persona pack bundles the instances that belong together.*" The second, from the build-versus-reuse section: "*build three custom skill packs (`kubevirt-ops` , `capi-ops` , `istio-ops`) and maybe 15–20 bespoke prompts. Everything else you can reuse.*"

Chapter 13 has already told what happened to the first noun: `SymposiumInstance` was never reconciled, and the chart no longer ships it. The second never existed; the document's own April status note records that "*there is no `PersonaPack` CRD.*" The three custom tool packs were never built.

What exists on 17 September

Checked against the running platform on 17 September:

Closest catalog persona	What exists	State on 17 September
Infrastructure-plane agents (VM warden, disk-image custodian, and others)	<code>cluster2-agent</code> , a generalist	Serving; read-only outside its namespace
Target-Cluster SRE	<code>target-cluster-agent</code>	Serving; its tools use the target cluster's administrator credential — read-only only by instruction
Service-Mesh Diagnostician	<code>mesh-sre-agent</code> (August)	Serving; read-only; no path to metrics
Idle VM Reaper	<code>cost-analyzer</code> (April demo)	Cannot be applied: written for a resource type the chart no longer ships
Incident responder	<code>incident-responder</code> (April demo)	Cannot be applied: same reason
Security and policy agents	An IaC security agent (March) — a separate Go service with policy rules and a local model, not a <code>Symposium</code> agent	Not redeployed after the September rebuild
The warm-up "keep-alive agent"	The <code>ollama-warm</code> schedule	Stopped firing (Chapter 14)

Asked to validate the two April personas without creating anything, the cluster refused both outright: "*no matches for kind '`SymposiumInstance`'.*" They were the first agents on the platform designed to *act* — one to stop an idle virtual machine on request, the other to diagnose and

remediate a stopped worker — and they are now documents describing something that cannot be installed.

One more row belongs in the table, because the strategy document's reuse plan depends on it. The document recommends starting from Symposium's `platform-team` bundle for the SRE and security personas. That bundle is installed. So are seven other example teams that ship with the chart, one of them seven personas strong. All eight are listed as ensembles in the agents' namespace, and all eight are switched off.

Three agents are running. Each is closest to one entry in the catalog but broader than it, and none is designed to act. Two are held to reading outside their own namespace by the permissions Chapter 14 examined. The third deserves a sentence of its own: `target-cluster-agent` reaches its cluster through the kubeconfig stored for it, and that kubeconfig belongs to `target-cluster-admin`. On the tenant cluster, the only thing keeping that agent read-only is its briefing.

The roadmap against the calendar

The roadmap was written with a status snapshot dated 24 April, and it gave its phases durations: one week for Phase 0, two to three for Phase 1, two for Phase 2, two to three for Phase 3. Phases 0 through 3 — the credible demo and the product, in the epigraph's terms — came to seven to nine weeks.

On 17 September, twenty-one weeks later:

- **Phase 0 — tighten the foundation.** Delivered on paper. The manifests are bundled into one install. The two network policies ship — and, as Chapter 14 found, the sandbox policy selects no pods. The warm-up schedule ships, and has stopped.
- **Phase 1 — infrastructure-plane personas.** Not built as personas; `cluster2-agent` covers the plane as a read-only generalist.
- **Phase 2 — mesh plane.** The Diagnostician exists, since August. The ServiceEntry generator does not, and neither does the planned hook that would run a diagnosis when someone clicks an edge in the mesh view.
- **Phase 3 — tenant plane and cost.** `target-cluster-agent` exists. The onboarder, the generalized reaper, the right-sizer, and the per-persona model selector do not.
- **Phase 4 — security and developer experience.** Not built as agents.
- **Phase 5 — the orchestrator.** Not built. The mechanism it would stand on works, as the next section shows.

It would be easy to call this a slipped roadmap and move on. The repository's own history says something more useful about where the weeks went. The Symposium labs in July discovered that the persona object had no controller, that tools did not reach runs, and that policy guardrails were weaker than their schema implied. The agent platform was upgraded three times — to 0.10.47, 0.10.57, and 0.10.75 — and each upgrade broke something that had to be diagnosed. The cluster speed investigation ran in June and again in September. The five-act service-mesh demonstration was built in August. And in September the whole platform was rebuilt on new Kubernetes versions.

None of that was on the roadmap. All of it was underneath it. A persona that takes "*~1–2 days of YAML + prompt engineering*" takes one or two days only on a platform that behaves the way its documentation says — and on this platform, finding out how it actually behaved was most of the work.

A pipeline of two

Phase 5 imagines "*a 'planner' persona whose tools are the other agents*": one chat window, a whole fleet behind it. Symposium's closest building block is the `Ensemble`, which lets agents hand work to one another. Lab 06 builds the smallest useful version: an analyst answers a question, and a reviewer receives the analyst's answer with instructions to "*verify claims, flag anything unsupported, and produce the final improved answer.*" The analyst runs on `qwen2.5:7b`, the reviewer on `llama3.2`.

It was run again on 17 September, on the current release, and the mechanics worked exactly as designed. Applying one object created two new `Agent` objects, each with its own memory Deployment. The kickoff prompt was delivered to the analyst. When the analyst's run finished, the reviewer's run started on its own, with the analyst's output passed along. About a hundred seconds from start to finish, with nothing triggered by hand.

The work did not go as well.

The analyst ran for 25.5 seconds, used 3,091 input and 859 output tokens, and called no tools. It produced a tidy, sectioned comparison of k3s clusters in KubeVirt VMs against bare Kind clusters — and one of its central claims was wrong. It described Kind clusters as "*leveraging the hypervisor's direct hardware access for containers.*" Kind runs Kubernetes nodes as ordinary containers on the host. There is no hypervisor involved.

That is exactly what a reviewer is for. The reviewer ran for 6.8 seconds, used 5,877 input and 126 output tokens, and made one tool call. Its entire output was a heading, "*Tool Usage,*" a sentence saying it would search the workflow's memory, and a short block of Python calling a memory-search function. No verdict. No correction. No final answer.

The lab had recorded the same outcome in July, on an earlier release: the reviewer "*ended with 'Please proceed with breaking down the key trade-offs...' — a hand-back instead of a refined final answer.*" Its conclusion holds word for word two releases later: "*the orchestration is deterministic, the output quality is not ... don't assume the last persona in a pipeline produces the polished deliverable every time.*"

Chapter 15 argued that an agent's answer is a claim, not evidence. A reviewer agent does not change that. Its output is another claim, from another small model, and a pipeline of agents adds each model's unreliability to the chain rather than a check on it. The planner in Phase 5 would be a longer chain.

A model for each job, on hardware you own

A sovereign fleet runs on the hardware the organization owns, so the size of the fleet is bounded by that hardware rather than by an invoice. Lab 07 asks the platform's model-fit service what fits.

On 17 September it described the host as an AMD Ryzen AI 9 HX 370 with an NVIDIA RTX 4050 laptop GPU — 6 GB of video memory — and about 30 GB of system memory, and it scored 1,365 candidate models against that hardware. Its top picks were 30- to 36-billion-parameter mixture-of-experts models, rated a "Good" fit at an estimated 26 to 32 tokens per second.

The two models the fleet actually uses are more revealing:

Model family, as rated	Fit on this host	Estimated speed
Llama 3.2, 3B Instruct (llama.cpp)	Perfect	~60 tokens/s
Qwen 2.5, 7B Instruct (llama.cpp)	Marginal	~36 tokens/s

The model this hardware suits best is `llama3.2`: the one that invented a list of nodes in Chapter 15, and the reviewer that handed back instead of reviewing above. The model that used its tools more reliably in lab 09, `qwen2.5:7b`, is the one rated a marginal fit. On this platform, the best fit for the hardware is not the best fit for the job, and a fleet chosen by a fit score alone would get worse.

The host has fourteen models downloaded. The fleet uses two. The per-persona model selector the roadmap planned for Phase 3 was never built, and every persona that needs a different model will compete for the same six gigabytes.

The auditors' job descriptions

Two entries in the catalog deserve a closer reading, because they are the agents a CxO would reasonably want first: the ones that watch the others.

The **RBAC Auditor** *"scans every `ClusterRoleBinding` + `RoleBinding` across both clusters. Flags wildcards, flags service accounts bound to `cluster-admin`, proposes least-privilege replacements."* The **Secret Hygiene Bot** *"flags in-cluster secrets older than 90 days, secrets mounted by more pods than necessary, and literal credentials in ConfigMaps."*

Now hold them against the most serious finding in this part of the book. Chapter 14 found that every run of `cluster2-agent` holds permission to read every Secret in its namespace — including an administrator credential for the target cluster. The role that grants it lists its resources and verbs explicitly: no wildcards. No agent is bound to `cluster-admin`. The Secret is not old, and it is mounted only by the agent it belongs to.

As written, the auditor's flagging rules would not have raised it, and neither would the hygiene bot's. Both job descriptions came from general best practice — sensible, and blind to the platform's actual risk.

The lesson generalizes. Write an auditor's job description against the incidents you have already had, then score it the way Chapter 15 scored the triage agent: plant a copy of the real finding, fix a strict pass rule, run it several times, and record what it costs. An auditor that cannot find the problem you already know about will not find the one you don't.

What a persona needs before it joins

Anita's question — what an agent has to prove before it gets the job — has an answer, and it is the rest of Part IV compressed into a checklist. The catalog's *one or two days of configuration and prompt writing* is only the first line.

1. **A job narrow enough to score.** A planted fault, a strict pass rule, several runs (Chapter 15).
2. **A model chosen for the job's behavior on this hardware** — not by fit score alone, and not by reputation (lab 07, lab 09).
3. **Tools delivered as named sidecars, with permissions checked against a live run's identity** — not against the documentation (Chapters 13 and 14).
4. **Every "must not" enforced below the model**, or explicitly declared unenforceable and verified instead (Chapter 14).
5. **Evidence attached to every claim that drives a decision** (Chapter 15).
6. **A known cost per task:** tokens, seconds, and the model it keeps loaded (Chapters 14 and 15).
7. **A record that outlives the cluster** (Chapter 13).
8. **For any agent that acts: an approval step between its decision and the change.** The strategy document draws one in its control tower. The platform's only working version today is a one-entry allowlist in the UI's backend.

Twenty-nine personas times eight requirements is a program, not a sprint — and it is the honest size of the fleet the catalog describes.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Anita's decision fit on one line of the catalog's cover page, written across the title.

Three, not twenty-nine.

Underneath, smaller: *Keep the three. Next two are the auditors — after their job descriptions include the finding from Chapter 14, and after each of them catches a planted copy of it, three times out of three.*

Nothing that acts until the approval step exists.

Vikram read it. "The auditors first. Not the ones that save money?"

"The ones that save money need permission to change things," she said. "Before anything gets that permission, I want something watching the permissions." She handed the catalog back. "Hire the auditors first. Then check that they'd have caught what we already know."

The ledger

- **Designed:** 29 personas across three planes and four shared groups; a six-phase roadmap; three custom tool packs.

- **Running on 17 September:** three agents — generalists, one per plane, none designed to act. Two are read-only outside their namespace by permission; the third holds an administrator credential for the tenant cluster and is read-only only by instruction.
- **Stranded:** two April personas written for a resource type the chart no longer ships; a March security service not redeployed after the rebuild; eight example ensembles installed and switched off.
- **Built of the three custom tool packs:** none.
- **Measured on 17 September:** a two-agent pipeline that orchestrated flawlessly in about 100 seconds (analyst 3,950 tokens, reviewer 6,003) and whose reviewer produced no review; hardware-fit ratings that favor the less reliable of the fleet's two models.
- **Roadmap:** Phases 0–3 estimated at seven to nine weeks from 24 April; at twenty-one weeks, Phase 0 delivered on paper and one persona from Phase 2 running.

Ask your team

1. For each AI agent we plan to add, what fault would we plant to score it, and what counts as a pass?
2. Which of our planned agents act rather than report — and what stands between an agent's decision and the change?
3. Would our planned auditors have caught the last real incident we had?

Open the repo

- `Sovereign_Cloud_Agentic_Strategy.md` — §2 (the three planes), §3 (the catalog), §5 (the roadmap), §6 (build versus reuse).
- `06-symposium/labs/06-ensemble/` — the two-agent pipeline, and its note on run-to-run variance.
- `06-symposium/labs/07-model-fit/run.sh` — what fits this hardware.
- `06-symposium/cost-analyzer.yaml`, `06-symposium/incident-responder.yaml`, and `docs/superpowers/specs/2026-04-21-symposium-sample-agents-design.md` — the April personas.
- `docs/superpowers/specs/2026-03-29-iac-security-agent-design.md` and `make security-deploy` — the security service.
- `kubectl get ensembles -n symposium-system` — the chart's example teams, installed and switched off.

PART V

The Dividend

Back to the boardroom.

Chapter 17 — The Same Cluster Twice

Safe only because exactly ONE target-cluster runs at a time.

— `03-target-cluster/target-cluster-warm.yaml`, the comment above the fast path

MERIDIAN · A TUESDAY IN OCTOBER · THE OPERATIONS FLOOR

The runbook for a sick cluster ran to forty-one pages. Vikram had written most of it, and he was not proud of it.

"Page fourteen," Anita said. "*If etcd reports a lost quorum, identify the surviving member.*" She turned the page. "And then eleven more pages of that."

"Every line of it was written after an incident."

"I believe you. That's what worries me." She put the binder down. "How long does it take to build a new one?"

"Thirty-four seconds."

"Then why is any of this forty-one pages?"

Vikram had been waiting for that question for about a year, and he still did not have a clean answer. "Because the cluster isn't the thing we're protecting. It's what's *in* it. And because—" he stopped, and started again. "Because the answer changes depending on which thirty-four seconds you mean."

"Explain that."

"Building one is thirty-four seconds. Getting rid of the old one is longer, and I don't actually know how much longer, because nobody has ever timed it properly. And rebuilding the whole platform underneath it took an afternoon." He shrugged. "One of those is on a slide. The other two aren't."

Anita picked up the binder again and weighed it in her hand.

"Time the whole thing," she said. "Not the flattering part. Destroy one and put it back, and start the clock when you type the destroy command, not when you start building. And do it five times, because the last time we published a number from one run I had to explain it twice."

* * *

BUILD LOG · 19 SEPTEMBER 2026

The cycle, timed

The book has measured a cluster being built — 33.7 seconds to both nodes `Ready` — and it has measured one being destroyed, once. It has never measured the thing an operator actually cares about: how long from deciding a cluster is beyond saving to having a working one serving traffic again.

So the boundaries were written down and committed *before* the first measured run. That ordering matters, because the estimate landed near a minute, and a boundary that can be moved after the data arrives is not a boundary. The clock starts when the destroy command is issued and stops on the first HTTP 200 served through the replacement cluster's Service — not on "pod Running," which is a claim, but on a request that was answered.

Five runs, after one discarded rehearsal, on a quiet host:

	Median	All five runs	Spread
Destroy to first served request	81.96 s	80.40 / 81.20 / 81.96 / 82.59 / 84.03	3.63 s
Delete command returns	10.13 s	9.16 / 10.13 / 10.13 / 10.13 / 10.14	0.98 s
Teardown — last object gone	45.73 s	44.76 / 45.58 / 45.73 / 45.97 / 46.00	1.24 s
Build — both nodes Ready	34.39 s	32.15 / 33.89 / 34.39 / 34.49 / 36.39	4.24 s
Build — control plane alone	25.60 s	25.30 / 25.46 / 25.60 / 25.67 / 26.08	0.78 s
Workload to first HTTP 200	1.34 s	1.07 / 1.08 / 1.34 / 1.41 / 1.65	0.58 s

The build reproduced. **34.39 seconds** against the **33.7** published in Chapter 7 is comfortably inside the spread of either measurement — a year of this book's arguments rests on that number, and it survived being measured again by a different instrument on a different day.

The teardown did not reproduce, and it is not close.

The number nobody had checked

Chapter 7 publishes teardown as **18.9 seconds**, and Appendix D is careful to label it *a single run*. Measured five times, it is **45.73 seconds**, with a spread of 1.24 seconds. Not a noisy figure that happened to land low once: a figure that is wrong by a factor of two and a half.

Because the published number counted API objects, and there are two defensible definitions of *gone*, one more run separated them:

Boundary	Time
The delete command returns	9.11 s
The qemu processes stop — the machines are actually off	35.60 s
The last API object disappears	45.33 s

Neither definition is 18.9 seconds. And the ordering is the interesting part: **the objects outlive the machines by about ten seconds**. Cluster API is still reconciling finalizers on things that have already stopped existing in any way that matters. If you want the honest headline, it is the second row — the machines are gone at 35.6 seconds — and the platform spends another ten seconds tidying up records of them.

This is the third figure in this book to be corrected, and the pattern is worth naming, because it is not about arithmetic. A control-plane-only time was published as a both-nodes time because the stopwatch counted a ghost. A 194-millisecond claim stood for three months because nobody re-ran it. Now a teardown measured once turns out to be less than half the real figure. Every one of them was a number that nobody had reason to doubt, in a book whose entire method is doubting numbers.

So: **destroy to serving, 82 seconds**. Bring-up is under a minute. The cycle is not. The slow half is the half nobody had put on a slide.

What the number buys

Four things follow from this, and this platform can show all four.

Recovery stops being a procedure and becomes a command. The industry's reference point is DORA's, which puts elite performers at restoring service in **under one hour**. Against that, 82 seconds of infrastructure is not a line item; it is a rounding error. Which is precisely the point, and the uncomfortable half of it: if replacing the infrastructure costs eighty seconds, then **everything left in your recovery time is data**, and that is the part this platform has never once tested. Chapter 18 is blunt about it — *"a restore that returns a running cluster with an empty database has restored nothing a business cares about."* The forty-one-page runbook does not disappear. It gets shorter and much more specific: it stops being about etcd quorum and starts being about the data.

Debugging gets a control group. This book already works this way without having named it. Chapter 8 records four hypotheses tested and discarded against rebuilt clusters. Chapter 11 records a failure that stopped reproducing after an upgrade, which is only a sentence you can write if you can rebuild the thing and look again. When a rebuild is cheap, the question *"is this my change or is this accumulated state?"* has an experiment instead of an argument. Palantir make the same case for their own infrastructure: *"Immutable nodes derived from instance templates simplify the debugging of any node issues encountered in production."*

Upgrades stop being events. Chapter 19 is this chapter's case study and it did not set out to be one. Kind has no in-place upgrade, so moving Kubernetes forward meant deleting both clusters and building them again — and the chapter's own line is *"the upgrade is the rebuild."* What was experienced as a limitation of a development tool turns out to be the operating model. The caveat is the one Chapter 19 also records: it deleted first and rebuilt after. A true blue/green — new beside old, traffic shifted, old deleted — needs two clusters alive at once, which is exactly what this platform cannot do, for reasons the next section is about.

Drills become a calendar entry. Chapter 9 already prescribed one: *"every week, someone claims it, deletes it, and times the rebuild — naming both nodes."* It was never run. This chapter is that drill, finally executed, and it found a published number was wrong by 2.4×. A rehearsal that costs a day gets scheduled once a year and quietly skipped. A rehearsal that costs eighty seconds has no excuse.

The trade nobody put on the slide

Now the other four — cluster-per-pull-request, many small clusters instead of a few large ones, cluster-per-tenant, and rotating clusters on a schedule to deny an attacker persistence. All four are good ideas. All four need the same thing: **several clusters, at once, with separate identities**. This platform can produce neither, and the reason is not an oversight. It is the same decision that bought the thirty-four seconds.

The fast path works by baking a fixed certificate authority set, a fixed join token and a fixed API address into the golden image, and pre-seeding matching secrets so the control plane adopts them. That removes the certificate work from the boot path — no `--cluster-reset`, no cert purge. The manifest says what it costs, in a comment sitting directly above the thing it enables:

| *Safe only because exactly ONE target-cluster runs at a time.*

The design document is blunter still: *"If you ever allow concurrent clusters, switch to per-build random identity."* And per-build identity is what the slower path does — the `:preinit` build, whose floor is around a hundred seconds. **The speed and the multiplicity are the same decision, taken in opposite directions**. You may have a thirty-four-second cluster, or you may have many clusters. Not both, on this design.

Chapter 5 noticed the consequence by accident: a kubeconfig found in `/tmp`, written for a cluster that had since been rebuilt six times, still worked. This chapter tested it on purpose. Run 1's administrator credential was kept. Five destroy-rebuild cycles later, it was presented to a cluster thirty-two seconds old — a different cluster, different node names, built from a different `kubectl` apply :

```
$ kubectl --kubeconfig /tmp/drill-kubeconfig-run1 get nodes
target-cluster-cp-lmrt2           Ready    control-plane   32s    v1.37.0+k3s1
target-cluster-workers-xq7mb-vbjr2 Ready    <none>          23s    v1.37.0+k3s1

$ kubectl --kubeconfig /tmp/drill-kubeconfig-run1 auth can-i '*' '*'
yes
```

It authenticated, as full cluster administrator.

Read that against the four ideas. **Cluster-per-tenant** is not merely unbuilt here; a credential issued to one tenant would open every other tenant's cluster, which is worse than the namespace-based isolation Chapter 14's scorecard already grades *not held*. **Rotation as a defence** returns a cryptographically identical cluster — the attacker's credential survives the rotation, so the rotation defends nothing. Palantir's 48-hour node lifetime works because *"compromising a single node is insufficient for an attacker to gain persistent access"* — but that is a claim about **nodes inside a long-lived cluster**, where the cluster's identity outlives the node and can revoke it. Rotating whole clusters that share one identity is a different operation that happens to look the same from the outside.

And even if the identity were fixed, the arithmetic would not budge. Chapter 20 measured one target cluster reserving **14.9 GiB** — roughly half this host. A second one does not fit. Many small clusters is not blocked by build time. It is blocked by memory, and by an address plan where every service address is pinned by hand in a table.

So the honest scorecard for the paradigm: four consequences demonstrated, four unavailable — and the four unavailable ones are unavailable *because of* the decision that produced the four demonstrated ones.

What replacement destroys

There is one more cost, and this book has paid it three times without writing it down as a cost.

When the cluster was rebuilt for the upgrade, **fifty-three agent run records went with it** — every record that existed that morning, the oldest nearly four weeks old (Chapter 13). Chapter 14's audit rule is graded *not held* with the observation that the trail "*vanished entirely in the rebuild.*" And when Chapter 11 tried to re-examine a finding from July, it could not: "*the July cluster no longer exists.*"

That last one is the sharpest, because it is the debugging argument running backwards. Cheap replacement lets you bisect against a known-good present. It also destroys the past you might have wanted to bisect *against*. A cluster you can rebuild in eighty seconds is a cluster whose evidence has a half-life of eighty seconds.

Which resolves the question of what backup means here, and it resolves it in two directions rather than one.

For the infrastructure, backup genuinely does become archival. If the cluster is reconstructible from a repository and an image, there is little point snapshotting its control plane. Chapter 18 makes the stronger version of this argument: if the repository is the source of truth, rebuilding from it "*is the only version of a disaster-recovery plan that is tested every day.*" Keep the artefacts cold, for the auditor and for the forensic case, and stop maintaining a restore path for a thing you can regenerate.

For everything else, the opposite. The run records, the logs, the audit trail, the application data — the cluster was never their archive, and in a replaceable world it very obviously is not. They need somewhere to go *before* the replacement, and a restore that someone has actually performed. Chapter 18's rule survives intact and gets sharper: *until the restore has happened, the backup is a hypothesis*. What changes is where the rule points. It stops pointing at etcd and starts pointing at the export.

That is the whole discipline in one sentence: **when the cluster stops being durable, the export becomes the archive.**

What eighty-two seconds does not mean

One caution before the ledger, because the number is quotable and a quotable number travels further than its caveats.

This is not a production recovery time. There is no detection in it, no decision, no data restore, no high availability, no storage to re-attach, no DNS to reprogram, no other tenants to consider, and no human deciding whether replacement is the right call. It is a lower bound on mechanical rebuild.

Nor is it provisioning from scratch. Every run recorded one line that settles this: the rebuilt cluster's `kube-system` namespace reports a creation timestamp of **16 September** — the day the image was baked — while its nodes were created that morning. The golden image carries a pre-initialised datastore. This is **restore-from-image**, and the image took five to eight minutes to bake. Without it already cached on the node, the honest number is minutes.

And *both nodes Ready* is still not two nodes of capacity. In all five runs the workload landed on the control plane, because the worker carries the taint Chapter 5 documented and nothing on this platform removes it. One correction to that chapter, though, from a test run for this one: a pod that explicitly tolerates the taint and pins to the worker **runs there perfectly well**. The worker is not broken. It is excluded. That is a scheduling decision that nobody made on purpose, which is a different and more fixable problem than the one Chapter 5 described.

The next chapter is the right place to take all of this apart, and it is deliberately unsentimental about exactly this kind of claim: *a reader who skips that list will mistake thirty-four seconds for readiness*.

* * *

MERIDIAN · THE FOLLOWING WEEK

Vikram brought the binder back, with a bookmark in it.

"Pages one to fourteen can go," he said. "That's everything about repairing a control plane. We don't repair control planes any more; we replace them, and I can show you the eighty-two seconds."

"And the rest?"

"The rest is pages about data, and those pages are now the whole runbook. They used to be an appendix." He put it down. "There's a second thing. Every cluster we rebuild comes back with the same certificates as the one it replaced. That's what makes it fast. It also means our rebuild isn't a rotation, and if we ever want one cluster per customer we have to give the speed back."

Anita thought about that for a while.

"Write that on the same page as the eighty-two seconds," she said. "Not in a footnote. Somebody is going to read that number and ask for a cluster per customer by Friday, and the answer is that we can have it, and here is the price."

The ledger

- **Measured (19 September, N=5, boundaries pre-registered before the first run):** destroy to first served request **81.96 s** median (80.40 / 81.20 / 81.96 / 82.59 / 84.03, spread 3.63 s). Build to both

nodes 34.39 s, reproducing Chapter 7's 33.7 s. Workload to first HTTP 200, 1.34 s. Zero image pulls in every run.

- **Corrected:** teardown, published at 18.9 s from a single run, is 45.73 s (N=5, spread 1.24 s). Split once: the machines stop at 35.60 s, the last API object goes at 45.33 s. The objects outlive the machines by about ten seconds.
- **Proved:** an administrator credential from run 1 opened a different cluster, thirty-two seconds old, five rebuild cycles later, with full cluster-admin.
- **Also corrected:** Chapter 5's worker is excluded by a taint, not broken — a pod that tolerates it runs there normally.
- **Carried from other chapters:** 33.7 s build and 18.9 s teardown (Chapter 7, the latter now superseded); 467 ms warm claim (Chapter 9); 14.9 GiB reserved per cluster (Chapter 20); one afternoon for a full platform rebuild (Chapter 19).
- **Demonstrated consequences:** recovery as a command, debugging with a control group, upgrades as rebuilds, drills as routine.
- **Unavailable on this platform, and why:** cluster-per-PR, many small clusters, cluster-per-tenant, rotation-as-defence — all four need concurrent clusters with separate identities, and the fixed identity that makes the 34-second build possible forbids exactly that. Per-build identity costs roughly a hundred seconds.
- **Not claimed:** any production recovery time; any data restore; that this generalises beyond one cluster on one host.

Ask your team

1. **How long does it take us to destroy something, and when did we last measure it?** Every organisation can quote its build time. The teardown in this book was wrong by a factor of two and a half because it had been measured once.
2. **If replacing the infrastructure took ninety seconds, what would be left in our recovery time?** That remainder is your real recovery plan. Everything else is a procedure for a machine you could have thrown away.
3. **When we rebuild something, does it come back with a new identity or the old one?** If it is the old one, the rebuild is not a rotation, and anything you were relying on it to revoke has not been revoked.

Open the repo

- `measurements/ch-recovery-drill.sh` — the drill, with its pre-flight gates;
`measurements/README.md` carries the pre-registered boundary table, written before the first measured run.
- `measurements/ch17-recovery-drill-20260919-163611.tsv` — all five runs, raw.
- `measurements/drill-workload.yaml` and `drill-workload-worker.yaml` — the served workload, and the worker-pinned variant that tolerates the taint.

- `03-target-cluster/target-cluster-warm.yaml` — the header comment: *"Safe only because exactly ONE target-cluster runs at a time."* Quoted in Appendix G.1.
- `docs/warm-pool-strategy.md` — *"If you ever allow concurrent clusters, switch to per-build random identity."*
- **Without the repo:** Appendix G.1 quotes the manifest that makes the trade; Appendix D carries the measurement protocol.

Chapter 18 — From the Lab to the Data Centre

That is not a data centre, and the book never pretends it is. What transfers from a laptop to a data centre is not scale. It is the pattern, the measurements, and — more often than the vendor brochures would suggest — the failures.

— the prologue of this book

MERIDIAN · A WEDNESDAY IN OCTOBER · THE FOURTH-FLOOR BOARDROOM

The pilot had done what a pilot is for. Anita Rao put one slide up, and it had two columns rather than a number.

"Everything on the left, we measured," she said. "Everything on the right, we would have to build. I am not asking the board to approve the right-hand column today. I am asking it to stop confusing the two."

The chief financial officer read it twice. "The left column is a laptop."

"One laptop. Twenty-nine gigabytes of memory, one consumer graphics card, one power supply, one of everything. It builds a Kubernetes cluster in thirty-four seconds and hands over a spare in under half a second, and both of those numbers are real."

"And in a data centre?"

"The same thirty-four seconds, probably. That part is the software, and the software does not care where it runs." She let that sit. "What changes is everything the laptop lets us skip. One of everything is fine until the one fails."

"Give me the list."

"That is the right-hand column. High availability. Storage that survives a disk. Backups we have restored from, not just taken. Addresses that work on a real network. Secrets that are not in Git." She paused. "Some of those we already know we are missing, because this year we went looking, and the looking is the part I would actually defend to you."

The CFO turned to Vikram. "How much of the right-hand column is new engineering?"

"Less than you would think. Most of it is choosing well-known components and wiring them in. The part that is genuinely new is the discipline, and we already have that. It is the only thing we built that I would call finished."

"Then write it down as a design," Anita said, "and mark every line of it clearly as something we have not done. The moment this chapter reads like a status report, it becomes the thing we spent the year arguing against."

* * *

A note on what this chapter is. Every other chapter in this book reports something that happened, with a date. This one does not. It is a design: an honest mapping of what the platform in the preceding sixteen chapters would need in order to run a production workload in a data centre. **Nothing in it was built, measured, or tested.** Where it draws on a real number, that number comes from an earlier chapter, and the chapter is named. Where it names a piece of software, that is a category illustrated by a well-known example, not a recommendation and not a verified fit.

The reason for the chapter is the question that follows every demonstration in this book: *yes, but would it work for real?* The useful answer is not "yes." It is a list.

What the laptop was allowed to skip

Start with what makes the lab a lab. Each of these is a real property of the platform described in this book, recorded in the chapter named.

- **One of everything.** One machine. The management cluster runs as a single Kind node; the target cluster has one control-plane machine and one worker. There is no second copy of anything, anywhere.
- **One failure domain.** A power cut, a kernel panic, or a full disk takes the entire platform — management cluster, workload cluster, models, agents, registry and console — at once.
- **Storage that is a directory.** Virtual machine disks live on the laptop's filesystem. There is no replication, no snapshot policy, and no backup. The platform has never restored anything.
- **A network with no network in it.** Everything shares one Docker bridge. Addresses are handed out by MetalLB in layer-2 mode from a hardcoded range, `172.18.255.200–225` (Appendix B). There is no router, no VLAN, no firewall, no external DNS.
- **Trust that was generated once and committed.** The warm path's certificate authorities are fixed, and their private keys are in version control, labelled demo-only (Chapter 3). The mesh's root of trust was generated on the laptop.
- **No identity.** The web console has no login. Neither does the agent console embedded in it. The one authenticated surface is the agent platform's own API, which does require a bearer token.
- **Capacity that is one deep.** One warm standby cluster, one model resident on the GPU at a time (Chapter 9).

None of this is a criticism of the platform. A lab that paid for all of it would have taught less, more slowly. But every item is a thing the data centre would have to supply, and a reader who skips this list will mistake thirty-four seconds for readiness.

High availability, and what it is actually for

The platform's control planes are single. Cluster API's own model is the first thing that changes: `KThreesControlPlane` takes a `replicas` field, and production practice is three, across three failure domains, so that the loss of one machine leaves a quorum. The same declaration that makes one

machine makes three; the cost is not complexity in the manifest, it is having three real places to put them.

Three points are worth a board's attention, because they are where the money and the mistakes are.

A failure domain is a physical claim, not a label. Three control-plane machines spread across three hypervisors in one rack, on one power feed, behind one top-of-rack switch, are three copies of one failure. The declaration will look correct. Chapter 5's lesson applies exactly: a status that says `Ready` is a claim, and the claim needs checking against what is actually true.

The management cluster becomes production infrastructure. In the lab it is scaffolding. In a data centre it manufactures and repairs workload clusters, so its own availability is now a dependency of everything it made. It needs the same treatment as the clusters it produces — which is the argument for keeping it small, boring, and rebuildable, rather than letting it accumulate.

Availability has to be tested by breaking things. This is the book's own rule, from Chapter 11: a check that has never failed is not a check. A high-availability design that has never lost a node on purpose, in front of someone, during working hours, is a diagram.

Storage, and the backup nobody has restored

Virtual machine disks are the platform's real state. In the lab they are files on one disk. In a data centre they need replicated block storage with snapshots — Rook/Ceph is the usual open-source answer for a Kubernetes-native deployment, and an existing enterprise SAN reached through its CSI driver is the usual answer for an organization that already has one. Either satisfies the requirement: a disk can fail without losing the machine that was using it.

Backup is a separate requirement, and the distinction matters to a board because the two are frequently conflated. Replication protects against a disk dying. It does not protect against a deletion, a bad upgrade, or a corrupted image being faithfully replicated three times. Velero is the common tool for backing up Kubernetes objects and, with the right driver, the volumes behind them.

The honest framing is the one the platform's own history supplies. Chapter 19 describes deleting both clusters and rebuilding them from the repository in an afternoon. That is a real disaster-recovery rehearsal, and most organizations have never run one. But note what it proved and what it did not: it proved the *platform* could be rebuilt. No workload data was in it. A restore that returns a running cluster with an empty database has restored nothing a business cares about.

So the design guidance is a sequence, not a product: choose replicated storage; back up objects and volumes; and then restore from that backup into a different cluster, on a schedule, as a test that is allowed to fail. Until the restore has happened, the backup is a hypothesis.

Networking that leaves the bridge

MetalLB transfers; the way this platform uses it does not. In the lab it runs in layer-2 mode on a Docker bridge with a hardcoded range, and the book's own address table is a discipline invented to stop two clusters advertising the same address (Appendix B). In a data centre the same component runs in BGP mode, peering with the real routers, announcing service addresses into the network the rest of the organization already uses. That is a conversation with the network team, and it is usually the longest lead-time item in the whole build.

Four more differences, each of which the lab simply does not have:

- **Segmentation.** Management traffic, workload traffic, storage replication and out-of-band console access belong on separate VLANs. The lab has one flat network, which is why nothing in this book has ever had to think about it.
- **DNS.** The platform resolves nothing externally; the mesh's cross-cluster names work because Istio was told about them. Production clusters need real names in the organization's DNS, with a plan for who creates records and how fast they change.
- **Certificates.** The mesh's root of trust was generated locally, and the warm path's certificate authorities are committed to the repository (Chapter 3). In a data centre both belong in the organization's own public-key infrastructure, with a defined lifetime and a rotation procedure that has been performed at least once. Automated issuance — cert-manager is already present on the management cluster — covers the leaf certificates but does not answer the question of who owns the root.
- **Firewalling.** Chapter 14 found the agent egress rule allows ports rather than destinations, so anything the agents can address on port 443 is reachable. On one laptop that is a finding. On a network with production systems on it, it is the finding.

Capacity planning, from the one measurement that transfers

This is the section where the lab genuinely earns something, because Chapter 20 measured the thing most capacity plans get wrong.

On the management cluster, 59 running pods had **reserved 22.42 GiB** of memory while the cluster was **using 7.45 GiB** — about a third. The target cluster's control-plane machine had been given 4 cores and 8 GiB, and was using **59 millicores, 1% of its processor allowance**, and 1,114 MiB of memory. Those are measurements, from 17 September 2026, on this platform.

The transferable lesson is not the numbers; it is which number governs. **Capacity is consumed by what is reserved, not by what is used.** The lab demonstrated the consequence directly: a second target cluster would not fit, because its machines could not reserve the memory, even though two-thirds of the memory already reserved sat idle.

For a data centre that produces this sizing method:

1. Size hosts from the **sum of reservations** of everything you intend to place on them, not from observed utilization.

2. Then treat the gap between reservation and use as the thing to manage down, because that gap is the owned-hardware form of the cloud waste in the prologue. Right-sizing requests is the work; it is unglamorous and it is where the margin is.
3. Reserve headroom deliberately: enough to lose a failure domain and still place everything that was running on it. That is an availability cost that appears on the capacity plan, and it is routinely forgotten.
4. Price warmth explicitly. Chapter 9's standby holds eight cores and 14 GiB idle, permanently, to make a claim take under half a second instead of thirty-four. That may be an excellent trade. It is never a free one, and it belongs on the plan as a line.

Secrets, identity, and the two things the lab has no answer for

Secrets. The platform commits certificate authority private keys and a fixed join token, marked demo-only (Chapter 3), and the agent platform's tokens live in cluster Secrets, which are encoded rather than encrypted. Neither transfers. The two common patterns are SOPS, which encrypts values inside the Git repository so the repository stays the source of truth, and External Secrets, which keeps material in a dedicated store — HashiCorp Vault, or a hardware security module — and syncs references into the cluster. The choice matters less than the property both provide and the platform currently lacks: *no plaintext secret in version control, and a rotation procedure that someone has actually run.*

Identity. Chapter 12 records that the model gateway shipped without a lock and Kiali runs with anonymous access; the web console has no login at all. In a data centre, every human-facing surface authenticates against the organization's identity provider through OIDC, and Kubernetes role-based access control is bound to real groups rather than to a shared credential. The agent platform makes this sharper rather than softer: Chapter 15 found that an upgrade silently changed which service account agent runs execute as, and a permission check written against the old identity kept reporting success. Identity is not a one-time configuration. It is something to re-verify after every upgrade, against the identity a running pod actually carries.

GitOps — the answer this book already argued for

Chapter 19 is the reason this section exists, and it makes the case better than any vendor's slide.

When both clusters were deleted and rebuilt, four things turned out to exist nowhere but in the running clusters: the Kind clusters themselves, KubeVirt and CDI, a hand-applied container image patch that no script recorded, and the KubeVirt CPU setting worth eight seconds on every machine start. One of those, the image patch, had been made once, by hand, months earlier, and was recovered only because someone thought to read the old cluster before deleting it. The chapter's conclusion is this book's central claim about sovereignty: *everything that lives only in a running cluster is not owned; it is on loan from that cluster's continued existence.*

GitOps is the mechanism that makes that conclusion enforceable rather than aspirational. A controller in the cluster — Flux and Argo CD are the two common implementations — continuously reconciles the cluster against a Git repository. The properties that matter here are three:

- **Drift becomes visible.** A hand-applied patch shows up as a difference between the cluster and the repository, rather than as institutional memory in one person's head.
- **The rebuild is the normal path, not an emergency one.** If the repository is the source of truth, rebuilding from it is what the system does continuously, which is the only version of a disaster-recovery plan that is tested every day.
- **Change has a review and an audit trail** — which is also what Constraint 6 of the platform's own agent strategy asks for, and what an agent that proposes changes will eventually need in order to be allowed to make them.

The design guidance is therefore blunt: **put the four unscripted things in Git first, before adding anything new.** The cluster definitions, the platform components and their versions, the tuning that is currently cluster state, and the base images. Chapter 19 already produced that list at no extra cost, by the expensive method of losing everything on purpose.

Observability that remembers

The lab's observability is real but has no memory. Prometheus holds the mesh's metrics with the sample add-on's retention and no durable storage; Kiali is anonymous; the recorded agent demonstrations are files on a laptop. Chapter 12 found the model traffic — the traffic the board asks about most — carried no model name and no token count, and was displayed nowhere.

For a data centre, three requirements follow, in this order of usefulness:

1. **Retention long enough to answer a question about last quarter**, which means durable storage behind the metrics system and a stated retention period, chosen against the audit obligations rather than the disk price.
2. **Logs and traces with the same lifetime**, so that an incident review can reach the evidence rather than the dashboard's memory of it.
3. **The AI-specific meter the platform still lacks:** tokens by model, by tenant. Chapter 20 measured energy per token on the GPU and found the electricity is a rounding error, which makes the meter a capacity and chargeback instrument rather than a power one — but it remains the closest thing to an invoice line that local AI has.

And one rule carried forward unchanged from Part III: a dashboard is not evidence. Every panel a decision rests on should be traceable to a check that can fail, and every such check should have been watched failing at least once.

Running outside the vendor's tested range, on purpose

Chapter 19 records a deliberate decision to run Istio on a Kubernetes version outside its tested range, written down next to the version number. That decision was correct for a pilot and is a live risk in a data centre, so it needs a policy rather than a habit.

Design guidance, and this is the part most organizations skip until an incident forces it:

- **Decide who carries the risk.** Either an internal team is willing to debug a mesh nobody upstream has tested this combination of, or the organization buys support and then lives inside the supported matrix. Both are defensible. Running unsupported while assuming someone else will help is not.
- **Record every exception where the next person will trip over it** — which the platform already does, in the mesh's README, and which is the cheapest control in this chapter.
- **Distinguish passing from supported.** The platform's own 22 of 22 checks passed on an unsupported combination. That proves the assertions held, not that the combination is tested. Chapter 19 says so explicitly, and the distinction is exactly what an auditor will ask about.
- **Budget for the upgrade treadmill.** A supported posture means tracking upstream release cadences and their support windows, and that is a standing cost in people, not a project.

Transfers, and does not

The two-column slide Anita put on the wall, filled in from the book's own chapters.

Transfers as-is	Does not transfer
The declarative model: a cluster is a written declaration, applied (Chapters 4, 5; Appendix G.1)	Single control planes and a single failure domain — everything in this book runs one of each
Build and teardown speed: 34.4 s and 45.7 s, which are properties of the software, not the hardware (Chapters 7 and 17)	The storage under it: local files, never replicated, never restored from
Cluster API and KubeVirt as the factory, including VMs as first-class objects (Chapter 4)	A flat Docker bridge with a hardcoded address range (Appendix B)
The service mesh's model: one label for cross-cluster failover, identity-based authorization with no sidecars (Chapter 11)	MetalLB in layer-2 mode; production wants BGP peering with real routers
Golden images and air-gapped boot — nothing fetched from the internet at start-up (Chapter 6)	Fixed certificate authorities and a join token committed to the repository (Chapter 3)
The warm-standby pattern, and its explicit price in idle capacity (Chapter 9)	One standby, one resident model — a depth of one is a lab constraint, not a design
Agents as Kubernetes objects, with policy, identity and tools declared (Chapters 13, 14)	Unauthenticated consoles, an unlocked model gateway, anonymous Kiali (Chapter 12)
Local models on owned hardware, and the measured energy per token (Chapters 3, 19)	One consumer GPU holding one model at a time (Chapter 9)
The capacity lesson: reservations govern, and the reservation-to-use gap is the margin (Chapter 20)	The absolute figures — 22.42 GiB reserved, 7.45 GiB used — which are one host's numbers, not a sizing model
The discipline: measure by name, assert on values, keep the before picture, record the dead ends (Chapters 7, 10, 11, 15, 18)	The assumption that anything here has been operated — no on-call, no SLA, no users, no auditor

Read the right-hand column as a work plan and it is long but ordinary: well-understood components, wired in by people who have done it before. Read the left-hand column as what a year of a single engineer and an AI assistant produced, and the interesting claim of this book survives the move to a data centre — not *we built a cloud on a laptop*, but *we found out what is true, cheaply, before spending the money*.

That is the last thing that transfers, and the only one that cannot be bought.

* * *

MERIDIAN · THE SAME AFTERNOON

The CFO had been writing while Vikram talked, and what he read back was a question rather than a note.

"If we approve the right-hand column, how long before someone shows me a screen that says everything is fine?"

"About a week," Anita said.

"And how long before I should believe it?"

She thought about the year: a worker marked `Ready` that could not run anything; a failover check that printed success during a total outage; a map labelled *live* drawing fourteen things that did not exist; a permission check that passed while the thing it checked had quietly moved.

"When someone has broken it in front of you," she said, "and the screen went red."

The ledger

- **Measured, and carried into this chapter:** build 33.7 s (Chapter 7, reproduced at 34.4 s in Chapter 17) and teardown 45.7 s (Chapter 17, superseding an 18.9 s single run); a warm claim in 467 ms (Chapter 9); 22.42 GiB reserved against 7.45 GiB used, and a control-plane machine using 1% of its processor allowance (Chapter 20); energy per generated token on the GPU (Chapter 20).
- **Designed, not built:** every requirement in this chapter — high availability, replicated storage, tested restores, BGP and segmentation, external DNS and PKI, secrets management, OIDC identity, GitOps reconciliation, observability retention, and a support policy. **None of it exists on the platform.**
- **Already known to be missing, from the book's own findings:** replicated storage, any restore test, real authentication on the consoles, destination-scoped egress, secrets outside version control, and four pieces of platform state that live only in a running cluster (Chapters 3, 12, 14, 18).
- **The cheapest item on the list:** writing the exception down next to the version number. The platform already does it.
- **Not claimed:** that any of this was tested, costed, or scheduled. This chapter is a list of requirements, and a list is not a plan.

Ask your team

1. Which parts of our platform exist only in a running system — and what would we have to reverse-engineer if we lost it tonight? Chapter 19 produced that list by accident; produce yours on purpose.
2. When did we last restore from a backup into a different cluster, and who watched it? A backup that has never been restored is a hypothesis with a budget line.
3. Where are we running outside a vendor's tested range, who decided it, and who carries the risk if it breaks at three in the morning?

Open the repo

There is nothing in the repository that implements this chapter, which is the point. What the repository does hold is the evidence behind each requirement:

- `03-target-cluster/target-cluster-warm.yaml` — `replicas: 1` on the `KThreesControlPlane`, and the fixed certificate authorities it depends on (quoted in Appendix G.1).
- `03-target-cluster/warm-ca/` and `scripts/gen-warm-ca.sh` — the committed demonstration keys, and the script that regenerates them.
- `06-symposium/agent-egress-ollama.yaml` — the port-only egress rule (quoted in Appendix G.5).
- `01-metallb/metallb-config.yaml` — layer-2 mode and the hardcoded address range.
- `CLAUDE.md`, the 2026-09-16 entries — the four things that existed only in a running cluster.
- **Without the repo:** Appendix G quotes the manifests above; Appendix B holds the address plan; Appendix C.3 catalogues the statuses that were wrong in both directions.

Chapter 19 — The Day We Upgraded Everything

Do not assume `latest` is safe just because a past jump was.

— `CLAUDE.md`, the repository's operating notes

MERIDIAN · A MONDAY IN SEPTEMBER · THE FOURTH-FLOOR BOARDROOM

The memo from the board's risk committee was one paragraph long, and Anita Rao had underlined its last sentence twice.

If the platform were lost, how long would it take to rebuild it, and from what?

She slid it across the table. "I couldn't answer that. Can you?"

"From the repository," Vikram said. "In theory."

"When did we last do it?"

"End to end? Never."

Anita waited.

"There's a second thing," he said. "The tools on our laptops are already two Kubernetes versions ahead of the clusters, which is outside what Kubernetes supports. And the clusters can't be upgraded in place. To move them forward, we delete them and build them again from nothing."

"So the upgrade is the rebuild."

"The upgrade is the rebuild."

"Then it's also the answer to the committee," Anita said. "Do both. Time it. And write down everything you needed that wasn't in the repository."

"One decision is yours first. The newest Kubernetes is 1.37. The newest release of the service mesh is only tested up to 1.36. We can stop at 1.36 and stay supported, or go to 1.37 and run the mesh outside its tested range."

"What does *outside its tested range* mean, in practice?"

"It means nobody at Istio ran their test suite on it. It doesn't mean it breaks. It means that if it breaks, we're the ones who find out."

She thought about it longer than he expected. "This is the pilot. Go to 1.37. Put the decision in writing, next to the version number, where the next person will trip over it."

* * *

Three things called Kubernetes

The instruction that started the day was four words long: *upgrade Kubernetes too*. The first finding was that this platform contains three different things called Kubernetes, at three different versions, with three very different upgrade costs.

- **The two host clusters**, `cluster1` and `cluster2`, ran Kubernetes 1.35.0 in Kind containers. Kind has no in-place upgrade: a new version means deleting the cluster and creating it again.
- **The target cluster** — the k3s cluster that runs as virtual machines — was pinned to k3s v1.31.4, six minor versions behind.
- **The web UI's backend** was built against the Kubernetes client library at v0.36.2.

The workstation's own `kubect1` was already at 1.37.0, and every command printed the same warning: *"version difference between client (1.37) and server (1.35) exceeds the supported minor version skew of +/-1."* The tools had moved on without the clusters.

The upgrade was carried out by an AI coding agent working in the repository under human direction — the same arrangement as the investigations in Chapters 8 and 15. Its first act was not to upgrade anything. It asked which of the three the instruction meant. The answer was all three.

That morning, the agent platform itself had already been upgraded, from Symposium 0.10.57 to 0.10.75 — and Chapter 15 tells what that broke. By early afternoon, the upgrade had become something the risk committee in the Meridian story would recognize: a full rebuild of the platform from its repository, the test most platforms never run until they have to.

Two decisions that belonged to a human

The newest versions were easy to find. Kind 0.33.0 shipped node images up to Kubernetes 1.37.0; upstream's stable release was 1.37.0; k3s v1.37.0+k3s1 had been published two days earlier. The conflict was one layer up.

The mesh ran Istio 1.30.3, and the repository's own README recorded its tested range: *"1.30 supports 1.32–1.36."* Before recommending anything, the agent checked whether a newer Istio would clear the way. Istio 1.31.0 was the newest release, and its announcement was unambiguous: *"Istio 1.31.0 is officially supported on Kubernetes versions 1.32 to 1.36."* No Istio release anywhere supported Kubernetes 1.37.

That made the version a decision rather than a lookup, and it went to a human, twice. The first time, the recommendation was to stop at 1.36.4 — the newest Kubernetes with a supported mesh — and the human chose 1.37.0. The second time, with the Istio finding spelled out, the options were to cap at 1.36.4 or proceed with the mesh *best-effort*. The human chose 1.37.0 again, knowingly.

That choice is now written where the next person will find it, in the mesh's README:

"Running here unsupported/best-effort, by deliberate choice, not because it was verified compatible."

KubeVirt and its storage importer, CDI, were left exactly where they were — KubeVirt 1.9.0 was already its latest release, and nobody had asked for either to change. Deleting two clusters is not the moment to pick up extra upgrades.

The agent stopped once more before the irreversible step, laid out what deletion would destroy and roughly how long the rebuild would take, and waited for an explicit *yes*.

Before anything was deleted

The reversible work came first.

The k3s version was pinned in **nine** files: seven cluster manifests and templates, a template-variables file, and the golden-image bake script. All nine moved to v1.37.0+k3s1. The ninth held a surprise. `bake-common.sh` documented a `K3S_VERSION` variable as the way to choose which k3s to bake, and set it at the top of the file. But the script that actually runs inside the bake VM is written out through a *quoted* heredoc — `cat << 'CLOUDINIT'` — which copies its contents verbatim and expands no variables at all. Inside it sat a second, hard-coded `K3S_VERSION`. The documented knob had never controlled anything. Anyone who had used it would have baked the old version and been told they had baked the new one.

The client library moved to v0.37.0; the backend built, passed `go vet`, and passed its tests. The Istio version default moved to 1.31.0. The `kind` binary on the workstation moved to 0.33.0.

Then came the inventory that matters most in this chapter: what exists in the running clusters that the repository does not create. It was checked against the scripts, the Makefile, and the README, and the answer was uncomfortable.

- **The Kind clusters themselves.** No configuration file, no `kind create` command anywhere in the repository. They had been created by hand, with defaults.
- **KubeVirt 1.9.0** and its configuration — the operator, the custom resource, and the tuned CPU limits from Chapter 8.
- **CDI 1.66.0**, the component that imports disk images into the cluster.
- **The base Ubuntu image**, a DataVolume called `ubuntu-noble-dv`, which every golden-image bake starts from.

Each one's specification was read out of the live cluster and saved before anything was deleted. Without that step, the rebuild would have begun by reverse-engineering the platform from memory.

One item on that list was wrong, and the error is part of the story. It comes back at the end of the chapter.

Rebuilding from nothing

The destructive part took two commands.

```
kind delete cluster --name cluster1
kind delete cluster --name cluster2
```

The local container registry survived, because it is a separate container rather than part of either cluster, and it was still attached to the Kind network when the new clusters came up. Both clusters were recreated from the `kindest/node:v1.37.0` image and reported Ready at v1.37.0 within minutes. KubeVirt and CDI went back on at their previous versions and reported Deployed. MetalLB went back on `cluster2`. Then Cluster API was initialized, which is where the day stopped being mechanical.

The dead image that reported success

Cluster API's initialization finished and reported success. The first attempt to create the target cluster failed immediately:

```
failed calling webhook "default.kthreescontrolplane.controlplane.cluster.x-k8s.io":
... connect: connection refused
```

The k3s provider's admission webhook was not running. Its pod had been stuck for eighteen minutes with one of its two containers failing to start, on an image that no longer exists:

```
Failed to pull image "gcr.io/kubebuilder/kube-rbac-proxy:v0.16.0": ... not found
```

The provider's release manifest — still version 0.3.0, unchanged — named a proxy image at an address that no longer serves it. And the setup script had not noticed, because every readiness wait in it ends in `|| true : wait five minutes, fail, fall back to a second five-minute wait, fail again, and carry on as though nothing had happened.`

The clue to the fix was in the cluster that had just been deleted. Before deletion, its providers had been running a different image, `quay.io/brancz/kube-rbac-proxy:v0.16.0`, which still pulls. Nothing in the repository, the Cluster API configuration, or any script recorded that change. It had been made once, in the running cluster, and it existed nowhere else. It was re-applied to both provider deployments, the target cluster was created, and the fix now lives in `02-capi-init/init-management-cluster.sh`, which patches the dead image right after initialization.

The setting that went missing

Chapter 8 ended with a warning: the CPU ceiling worth eight seconds on every VM start lives in the running cluster's KubeVirt configuration, not in Git. The script that applies it is called at the very end of the Cluster API setup script, so that a rebuild restores it automatically.

After the rebuild, it was not there. It was applied by hand before any VM was built, which is why the measurements later in this chapter include it.

The reason it was missing is worth telling precisely, because the first explanation written down was wrong. The setup script's output stopped at a single line — `Waiting for k3s bootstrap controller...` — and nothing after it ran. The agent had launched the script under a 590-second time limit. The script's readiness waits for that provider, stuck on the dead proxy image, could take up to 600 seconds on their own. The time limit killed the script partway through its waiting, before it ever reached the step that restores the CPU ceiling.

And the run still reported success. The command's output had been piped through `tail` to shorten it, so the exit code that came back belonged to `tail`, which had succeeded, and not to the script, which had been killed. The same trap caught a failed installation later the same day.

The agent noticed the missing setting and restored it. Then, in the repository's notes, it recorded the cause as a race with KubeVirt's startup. The log does not support that; the evidence points to the time limit instead. The correction is in this chapter's draft notes, and it belongs in `CLAUDE.md` before anyone relies on the wrong one.

The image that still had the old version inside

The golden images were next, and here the surviving registry turned from an asset into a trap. It still held all three of the platform's k3s images — `:warm`, `:preinit`, and `:latest` — and all three had k3s v1.31.4 baked inside them.

The Makefile's `ensure-warm-image` step checks whether a `:warm` image is *present*. It does not check what is inside. Left alone, it would have found the old image, reported it present, and booted every new cluster on the version the upgrade was meant to replace.

The `:warm` image — the default path — was rebaked from scratch: a bake VM ran for 160 seconds and downloaded k3s v1.37.0; the result was packaged as a 1.35 GB container image and pushed over the old tag. The `:preinit` and `:latest` images, used only by older deployment paths, were **not** rebaked, and still carry v1.31.4.

The status that disagreed with the nodes

The target cluster came up, and its control-plane object promptly reported a problem:

`ControlPlaneComponentsUnhealthy`, *"Following machines are reporting control plane errors,"* with its ready flag never set.

The machines themselves reported every condition true. The cluster's own API, reached through its kubeconfig, showed both nodes `Ready` on `v1.37.0+k3s1`, Ubuntu 24.04.5, containerd 2.3.4. The repository's verification script passed. The status field was stale, and a wait loop polling it was stopped rather than left to time out. It is the smaller cousin of Chapter 15's lesson: a status field is a claim too.

Twenty-two green checks

The mesh was the part running outside its tested range, so it was the part that most needed evidence. The repository's mesh demonstration is five acts, and each act ends in assertions rather than screenshots.

Istio 1.31.0 installed cleanly on both clusters, and both control planes adopted the same root certificate — identical fingerprints, the precondition for joining them into one mesh. Then the five acts ran, on Kubernetes 1.37:

- **Gateway API.** HTTPS served with a locally rooted certificate; a 90/10 canary split measured 93/7; header-based routing sent 10 of 10 internal requests to v2; request mirroring and the HTTP-to-HTTPS redirect worked; and a rogue tenant's attempt to claim the hostname was refused with `NotAllowedByListeners`.
- **Waypoints.** Identity-based authorization: the trusted identity got `200`, the untrusted one `403`; a `POST` was denied while a `GET` was allowed; swapping a workload's identity revoked its access with no policy change; the client's identity arrived in a forwarded-certificate header; an in-mesh traffic split measured 24/26.
- **Multicluster.** Each cluster's control plane synced with its peer. Then every local replica of a service was scaled to zero, and calls kept being served from the other cluster, with no configuration change. (The act's own tally, `success=20 failed=0`, proves less than it appears to — it prints the same line when both clusters are down; see Chapter 11. Re-measured by HTTP status on 17 September: 30 of 30 calls returned `200`, every one served by cluster2.)
- **AI gateway.** Model traffic routed through the mesh to the host's Ollama; a completion came back `"ok"`.
- **Observability.** Prometheus held the mesh's request metrics.

The verification script's final table, once the agent platform was reinstalled: **22 passed, 0 failed, 0 skipped**.

It is worth being exact about what that proves. *Supported* means the Istio project ran its test suite on this combination. *Passing* means this repository's twenty-two assertions held. Those are different sets of checks, and the gap between them is not empty.

Act 1's output contained one line that no assertion examines. The act lets a tenant contribute its own listener to the shared gateway through a `ListenerSet`, prints the result, and moves on. The result was:

```
Accepted=False (ListenersNotValid)
Programmed=False (ListenersNotValid)
```

— *"None of the ListenerSet's listeners are valid."* Yet the gateway's Service exposes that tenant listener's port, `8443`, exactly as it did before the rebuild. The status says the listener is invalid; the deployed Service says it exists. Whether traffic actually flows on port 8443 was never tested, and nobody kept

the pre-upgrade status to compare. It may be a regression, a long-standing quirk, or a status bug. It is a question, not a finding — and it would have stayed invisible behind a clean 22 of 22.

Five known traps, one attempt each

Reinstalling the agent platform on the new `cluster2` produced five errors in a row. None needed an investigation, because every one was already written down in the repository's notes, left there by earlier upgrades.

1. **The admission race.** The first install failed with `failed calling webhook "vskillpack.symposium.ai" ... connection refused`: the chart applies its built-in skill packs before its own validating webhook is serving. Documented. Wait for the webhook, retry.
2. **The wait that never ends.** The retry, run with Helm's `--wait`, timed out on an example `postgres` tool server that is deliberately left suspended and never becomes ready. Documented, with the instruction to drop `--wait` and verify by hand. (This was the second failure of the day to report success through a pipe to `tail .`)
3. **The field-ownership conflict.** Without `--wait`, the install stopped on a conflict over `.spec.sidecar.mountWorkspace`, a field the platform's own controller rewrites after install. Documented, with the exact flags: `--server-side=true --force-conflicts`. The next attempt deployed.
4. **Runs that never retry.** Recreating the serving agents left two of their runs in a terminal `Failed` state — *"server Deployment not found"* — because they had been recreated at the same moment their deployments were being replaced. Documented: the controller does not retry a failed run; delete it and let it be regenerated.
5. **The silent crash loop.** One serving pod crash-looped with no log output. Documented, with three distinct past causes and the same first move for all of them: regenerate the run. It came up healthy.

The per-run service-account fix from Chapter 15 had already been written into the manifests that morning, so the reinstall applied the corrected version rather than rediscovering the problem.

Five errors, five single attempts. The difference between a known failure and an unknown one is the difference between a minute and an afternoon, and the only thing that separated them was a text file.

No faster than before

With everything rebuilt, the target cluster's time-to-ready was measured with the same instrument as Chapter 8: three clean runs, host load recorded before each. It reported **23.9, 24.1, and 23.7 seconds** — ten seconds faster than the 34.5 measured a week earlier, and steadier.

That result was wrong, and finding out why is the second half of Chapter 7. In short: the warm golden image carries a ghost Node left over from the machine it was baked on, still marked `Ready`. For the first seconds of every new cluster, the ghost and the control plane together satisfy any

stopwatch that waits for "two nodes Ready" — so the instrument stopped its clock before the worker had joined, and reported the control plane's time as the total.

Measured again the next day, naming each node rather than counting them, three clean runs:

	Run 1	Run 2	Run 3	Median
Control plane Ready	23.4 s	25.1 s	24.7 s	24.7 s
Worker Ready — both nodes	30.4 s	35.1 s	33.7 s	33.7 s

33.7 seconds, against 34.5 before the upgrade. Within the noise of three runs, the upgrade changed nothing. Six minor versions of k3s and two of Kubernetes, and a cluster that builds in the same time — which is a perfectly good result, as long as nobody claims it is a better one.

Tearing the cluster down was measured separately, once: the delete command returned after **10.15 seconds** while Cluster API worked through its finalizers, and every VM and machine object was gone at **18.9 seconds**. That was a single run, and it did not survive being measured again: five runs in Chapter 17 put it at **45.73 seconds**, with the machines themselves stopping at 35.6 s. The figure below is left as it was recorded on the day.

The balance sheet of memory

Here is the day's accounting in the terms the risk committee asked for.

Rebuilt from the repository: both host clusters, the golden image, the target cluster, the service mesh with all five acts, and the agent platform with its three agents. **Recovered from the old cluster, because nothing else held it:** the Kind cluster setup, the KubeVirt and CDI installation, and a working proxy image for the k3s providers. **Found wrong in the repository:** a version variable that controlled nothing, an image check that tested presence rather than contents, and a service-account binding that had quietly stopped applying. **Already known, and therefore cheap:** five errors that cost one attempt each.

And one item from the inventory was simply wrong. The base image was listed as something no script creates, and the repository's notes were updated to say so. While this chapter was being researched, the script turned up: `scripts/import-base-images.sh`, which imports the base image *and* a second, minimal Ubuntu image that the rebuild never recreated. It exists. It is mentioned in exactly one document — a post-mortem about something else. Nothing on the setup path — not the Makefile, not the README, not the operating notes — points to it. A script nobody can find is not institutional memory. It is an attic.

That is the real lesson of the day, and it is the one this book's opening chapters were building toward. Sovereignty is usually argued as a question of *where*: whose data center, whose jurisdiction, whose model. But owning a platform also means being able to rebuild it without asking anyone. By that measure, everything that lives only in a running cluster — a hand-applied image patch, a tuned CPU limit, a cluster created with defaults — is not owned. It is on loan from the cluster's continued existence.

The notes paid for themselves in minutes, on a single afternoon. Written knowledge compounds, and unwritten knowledge is deleted along with the cluster that held it.

* * *

MERIDIAN · THE FOLLOWING MONDAY

Vikram brought one page to answer the committee's memo.

Rebuilt from the repository, on current versions, in one afternoon. Four things were not in the repository. One is now automated. The other three are written down but not yet scripted — including a setup script that existed all along and that nobody could find. One component runs outside its vendor's tested range, by decision, recorded next to the version. Cluster build time: 33.7 seconds for both nodes, unchanged by the upgrade — after our own stopwatch briefly said otherwise.

Anita read it and went back to one line. "One afternoon. Because we'd never done it."

"The first time is the expensive one."

"Then it doesn't get to be the only one." She wrote on the memo before handing it back: *Rebuild from nothing — quarterly.*

"A platform we have never rebuilt," she said, "is a platform we only think we own."

The ledger

- **Upgraded:** Kubernetes 1.35.0 → 1.37.0 on both host clusters; k3s v1.31.4 → v1.37.0 in the target cluster; Istio 1.30.3 → 1.31.0; Symposium 0.10.57 → 0.10.75; kind 0.31.0 → 0.33.0; the client library 0.36.2 → 0.37.0.
- **Upgraded without anyone choosing it:** Cluster API 1.14.0 → 1.14.2, installed at whatever version `clusterctl` fetched that day; and cert-manager on `cluster2` at 1.21.1, brought in by `clusterctl`, while the agent platform's install script pins 1.16.2 — it found cert-manager already present and skipped its own.
- **Rebuilt and verified:** both host clusters, the golden image, the target cluster, all five mesh acts (22 of 22 checks), and the agent platform.
- **Measured:** bring-up 33.7 s median for both nodes (control plane alone, 24.7 s) — no measurable change from 34.5 s; teardown 18.9 s — superseded by Chapter 17's 45.73 s across five runs. An earlier reading of 23.9 s was the control plane only (Chapter 7).
- **Written back into the repository:** the dead-image patch, the bake-script warning, the unsupported-mesh decision, and an upgrade entry in the operating notes.
- **Still owed:** the Kind clusters, KubeVirt, and CDI documented as prerequisites but still not scripted; a mesh outside its tested range; a ListenerSet whose status and Service disagree; `:latest` and `:preinit` images still carrying k3s v1.31.4; a minimal base image never re-imported; Kind node images pinned by tag rather than digest; timing scripts that count a ghost node as a worker; and a recorded cause for the missing CPU setting that is wrong.

Ask your team

1. **If we lost the platform tonight, could we rebuild it from what is written down — and when did we last prove it?** Not believe it. Prove it.
2. **What in our running systems exists nowhere else?** Compare what is deployed with what the repository would deploy, and treat every difference as a liability.
3. **Where are we running outside a vendor's supported range, who decided that, and where is the decision written?** An unrecorded exception looks exactly like an accident.

Open the repo

- `CLAUDE.md` — the entry "*Kubernetes/Istio/k3s upgraded 2026-09-16*", with the rebuild's findings.
- `02-capi-init/init-management-cluster.sh` — the dead-image patch, and the comment explaining it.
- `bake-common.sh` — the header comment on the `K3S_VERSION` variable that controlled nothing.
- `07-istio-advanced/README.md` — the note recording the unsupported combination; `make istio-adv-verify` reruns the 22 checks.
- `./scripts/phase-timings.sh --runs 3` — the bring-up measurement protocol.
- `scripts/import-base-images.sh` — the script the rebuild could not find.
- **Without the repo:** Appendix G.7 quotes both `K3S_VERSION` assignments and the quoted heredoc between them; G.8 quotes the dead-image patch.

Chapter 20 — Reclaiming the Margin

The question is no longer whether you can afford to become your own cloud provider, but whether you can afford not to.

— *Reclaiming Margins with Sovereign Cloud, strategic briefing, February 2026*

MERIDIAN · A FRIDAY IN SEPTEMBER 2026 · THE CFO'S OFFICE

The chief financial officer had a single slide on his screen, and he had clearly been waiting all week to show it to someone.

Two columns. *Managed Cloud (5-Year): \$12.5M — escalating costs, no equity. Sovereign Cloud (5-Year): \$8.2M — predictable costs, asset value.* And across the bottom, in green: *Net Savings: \$4.3M (34% reduction) + owned infrastructure assets.*

"A consultant sent it," he said. "It's the number I asked you for a year ago. Which parts of the bill are rent. I want it in the board pack."

Anita read it twice. "Whose twelve and a half million?"

"An enterprise like ours, apparently."

"Then it isn't ours." She sat down. "I'm not saying it's wrong. I'm saying I can't defend it, because I can't find any of it on our invoice. Can we take it apart before the board does?"

The CFO looked at the green line for a moment, then closed the slide and opened a blank spreadsheet. "Line by line," he said. "Go on, then."

* * *

The slide

BRIEFING · FEBRUARY 2026

The slide is the fifth and last chapter of the strategy deck behind this book, and it is the most quoted. It compares five-year total cost of ownership — the full cost of an option over its life, not just its first invoice — for a managed cloud and a self-managed sovereign cloud:

	Five-year cost	The deck's description
Managed cloud	\$12.5M	"Escalating costs, no equity"
Sovereign cloud	\$8.2M	"Predictable costs, asset value"
Net saving	\$4.3M (34%)	"+ owned infrastructure assets"

It explains the saving with four "*Cost Elimination Factors*":

Factor	Saving	The deck's description
Success Tax elimination	-\$1.8M	"No more exponential scaling fees"
Innovation Tax removal	-\$1.2M	"Stop paying premium for stagnation"
Vendor licensing savings	-\$900K	"Open-source alternatives"
Infrastructure optimization	-\$400K	"Right-sized for actual workloads"

Elsewhere, the deck adds the numbers a board will ask about next: "20–35% Cost Reduction", a "5–7yr Payback Period", "12–18 months to full production readiness", an "Initial Investment: \$2–5M for mid-size enterprise deployment", and a "3–5 Year Timeline" for the whole transformation.

None of these figures has a source. They describe an illustrative enterprise, not a measured one. That doesn't make the slide useless. It means the slide has to be read the way a CFO reads any model: by asking whether it holds together.

Three questions for any TCO slide

Do the parts add up?

They do, exactly. $\$1.8\text{M} + \$1.2\text{M} + \$0.9\text{M} + \0.4M is $\$4.3\text{M}$, which is $\$12.5\text{M}$ minus $\$8.2\text{M}$, which is a 34.4% reduction.

That exactness is the problem. All four lines are *savings*. There is no line for anything owning *adds*. Yet the deck's own previous chapter lists what a self-managed platform requires: hardware; data-centre facilities; a platform engineering team of "Kubernetes experts, DevOps, SRE specialists"; automation and tooling; operational processes. The migration itself costs money, and so does the period of paying for both the old estate and the new one. If the four savings add up exactly to the net figure, every one of those costs is either hidden inside the four lines or missing from the model.

Can you find the lines on an invoice?

Two of the four, yes. Licences are an invoice line, and so is infrastructure you can right-size.

The other two are not. "Success Tax elimination" and "Innovation Tax removal" are the deck's names for two ideas from Chapter 1: that costs grow with success, and that money spent standing still is money not spent building. Both are real ideas. Neither is an account. Together they are **\$3.0M of the \$4.3M — about 70% of the saving** — sitting in lines no finance system records.

Do the slides agree with each other?

A 34% saving over five years, and a payback of five to seven years, are hard to hold at the same time. Payback is how long it takes the savings to repay the initial investment. If the model saves money within five years, it has paid back within five years.

The two agree only under assumptions the deck doesn't state. Take the deck's own initial investment of \$2–5 million, and suppose the \$8.2 million *leaves it out*:

Initial investment	Five-year cost of owning	Five-year saving
\$2.0M	\$10.2M	+\$2.3M (18%)
\$3.5M	\$11.7M	+\$0.8M (6%)
\$5.0M	\$13.2M	−\$0.7M (−6%)

Or suppose the \$4.3 million is a steady saving of \$0.86 million a year: then a \$2 million investment pays back in **2.3 years**, and a \$5 million one in **5.8 years**. The deck's five-to-seven-year payback fits only at the expensive end of its own range — where the five-year saving shrinks toward nothing, or turns negative.

None of this means the slide is wrong about the direction. It means the slide is a *shape*, not a number. Its real value is the list of lines. The amounts have to come from somewhere else.

An asset that wears out

PUBLIC RECORD · 2022–2025

The slide's right-hand column claims "*asset value*", and an earlier slide lists "*Asset value appreciation*" and gives the asset value as " ∞ ". Servers do not appreciate. They wear out, and they become obsolete, and accounting says how fast.

The clearest evidence comes from the landlords themselves. Microsoft, in its 2022 fiscal year, and Alphabet, in 2023, lengthened the accounting life of their servers and network equipment from four years to six, which reduced their reported depreciation by billions. In February 2025, Amazon went the other way for part of its fleet, shortening it from six years to five from the start of that year, and citing "*the increased pace of technology development, particularly in the area of artificial intelligence and machine learning.*"

The three largest owners of servers in the world disagree about how long a server lasts. A TCO for a sovereign platform should not assume they last forever. Plan a hardware refresh inside, or just after, a five-year model — and expect AI hardware to age faster than the rest.

What the platform can put on the worksheet

BUILD LOG · 17 SEPTEMBER 2026

The platform in this book cannot supply Meridian's numbers. It can supply something the slide lacks: inputs that were measured. Four of them matter for a TCO.

Speed. A two-machine Kubernetes cluster is built from its declaration to both machines ready in a median of 33.7 seconds, and torn down in 45.7 (Chapters 7 and 17). For steady workloads, "we rent because renting is faster" is no longer a line that has to be paid for.

Reservation versus use. Waste is not only a cloud problem. On the management cluster that afternoon:

- Its 59 running pods had **reserved 22.42 GiB** of memory. Of that, 14.57 GiB was held by the two virtual machines of the target cluster.
- The whole management cluster was **actually using 7.45 GiB** — about a third of what it had reserved.
- The target cluster's control-plane machine had been given 4 cores and 8 GiB. It was using **59 millicores — 1% of its processor allowance — and 1,114 MiB of memory**.
- The node had reserved 4.59 processor cores. At the moment it was sampled, it was using well under half of one.

On owned hardware, capacity is set by what is **reserved**, not what is used. On this host, a second target cluster would not fit — its machines could not reserve the memory — even though two-thirds of the memory already reserved sat idle. That is the deck's "*overprovisioning*", on hardware you own. It is the owner's margin to reclaim or waste, and right-sizing reservations is where it lives.

The energy cost of AI. The model most of the platform's agents use, `qwen2.5:7b`, was measured on the laptop's GPU. With the model loaded and idle, the GPU drew **6.4 W**. While generating text it drew **38–42 W**, at **34–37 tokens per second**. That comes to **1.19 joules per generated token — 0.33 kWh per million tokens**, for the GPU alone.

Multiply that by any electricity price between 10 and 50 cents a kilowatt-hour, and the power for a million generated tokens costs roughly **3 to 17 cents**. The lesson is not that local AI is nearly free. It is that for local AI, the electricity is a rounding error. The cost is the hardware — which, on this GPU, holds one model at a time (Chapter 9) — and the people who run and check it.

The people. The repository behind this book holds **88 commits, on 18 distinct days**, from 11 February to 18 September 2026, from a **single committer, Mahipal**. **70 of those 88 commits** carry an AI coding assistant as co-author. (The count moved while this chapter was being written, because writing the book added commits to the platform — which is its own small lesson about measuring a thing you are standing on.) The platform's operating notes run to 612 lines, including 23 entries in a list of pitfalls — most of them a failure that looked like success.

Read that carefully, because it is the easiest number in this book to misuse. It describes a lab: no on-call rota, no service-level agreement, no production users, no auditors. It says nothing about the cost of *operating* a sovereign platform at Meridian's scale. What it does suggest is narrower, and still worth a line in the model: in 2026, the effort to *build* such a platform is not the effort a 2021 estimate would assume. And the cost that never appears on a TCO slide is the one those 23 pitfalls describe: the continuous work of knowing what is true.

One more fact belongs on the worksheet. The platform's cost-analysis agent — the persona written in April to find idle virtual machines and propose stopping them — cannot currently be installed; it was written for a resource type the agent platform no longer ships (Chapter 16). The tool meant to measure the owner's waste is itself waiting to be fixed.

The worksheet

Put the slide's shape and the platform's measurements together, and the chapter's real deliverable is a worksheet: the lines a sovereign-cloud TCO must contain, and what this book can say about each. The amounts come from your own invoice.

Line	Renting	Owning	What this book measured
Capacity	On the invoice	Hardware, amortized over its refresh cycle (4–6 years, per the public record)	22.42 GiB reserved against 7.45 GiB used
Speed to capacity	Minutes	The factory's build time	33.7 s to a two-machine cluster; 45.7 s to remove one
Software	Managed-service fees and licences	Open source, plus optional support	The whole stack is open source (Chapters 3–4)
Leaving	Switching fees (banned in the EU from 2027) and rebuilding proprietary services	Migration project, and paying for both during it	No migration tooling in the repository (Chapter 4)
Power and facilities	Inside the rent	Power, cooling, space	0.33 kWh per million generated tokens, GPU only
AI inference	Per token, hosted	Hardware sized for the models	34–37 tokens/s; one model at a time on a 6 GB GPU
People	The supplier's staff	Platform team, on-call, training	A lab: one committer (Mahipal), 70 of 88 commits AI-co-authored, over 18 days
Upgrades	A button and a maintenance window	Yours	A week, for one major upgrade (Chapter 19)
Knowing what's true	The supplier's status page	Your own checks	Gauges found wrong in Chapters 4, 5, 11 and 12

One workload, metered both ways

BUILD LOG · 18 SEPTEMBER 2026

A worksheet with no amounts in it is still a slide. So the test this chapter proposes — meter one workload both ways, in parallel — was run at the only scale available: one workload, on the laptop.

The workload. An `nginx` deployment on the target cluster, served continuously and probed every twenty seconds for ten minutes. It answered 30 of 30 probes with **200**. It is a small, steady, unglamorous workload, which is exactly the kind the a16z paper says is cheapest to own.

What it consumed, measured. The platform as a whole ran at a **median 0.549 of the host's 24 threads — 2.29% of its processor — and 8.71 GiB of 29.96 GiB**. The target cluster's control-plane node used 67 millicores and 1,142 MiB.

But use is not what capacity costs, and this is where the lab produced its sharpest number. The two virtual machines **reserve 14,920 MiB — 48.6% of the host's memory** — while reserving only 400 millicores each, 3.3% of its processor. The `cores: 4` in each machine's declaration is a guest topology setting, not a claim on the host's CPU at all. Memory is the binding constraint, and one cluster of this shape holds down **roughly half the machine**, whatever it is doing.

What could not be measured, and therefore is not claimed. Whole-host power. The CPU energy counter at `/sys/class/powercap/intel-rapl:0/energy_uj` is present but not readable without root on this kernel, and the battery reads zero on AC. So host power below is an assumption with a range, not a measurement. The GPU was measured — a steady 5.22 W — and is idle: this workload never touches it.

The rented side, priced from published list prices, accessed 18 September 2026, with no committed-use discount:

Item	Price
DigitalOcean CPU-Optimized Droplet, 4 vCPU / 8 GB × 2	\$0.125 / hour each
DigitalOcean Kubernetes standard control plane	free
<i>(alternative)</i> high-availability control plane	\$40 / month ≈ \$0.055 / hour
<i>(for comparison)</i> Amazon EKS control plane, standard support	\$0.10 / cluster-hour

Rented: \$0.250 per workload-hour for the same two-machine shape with a free control plane — \$0.305 with a highly available one. A managed control plane is not always free, which is what the EKS line is doing there.

The owned side needs four inputs. One is measured, three are assumptions, and the assumptions are stated so a reader can substitute their own:

Input	Value	Measured?
Share of the host this workload holds	48.6%	measured (memory reservation)
Host capital cost	\$1,500–\$3,000	assumption
Amortization	4 years — the conservative end of the public record above	assumption
Host power under this load	25–65 W, at \$0.10–\$0.50/kWh	assumption; RAPL unreadable

That gives hardware at **\$0.021–\$0.042** per workload-hour and power at **\$0.001–\$0.016** — call it **\$0.02–\$0.06 per workload-hour** for the metal.

Against \$0.250 rented, that looks like a rout: four to eleven times cheaper. It is also wrong, because it leaves out the only line that actually decided this.

The line that decides it

The repository carries commits on **18 distinct days** between 11 February and 18 September 2026. That is a measured lower bound on the human effort — real work happens on days without commits, and a commit day is not a full day — but it is the only effort figure this book has that is not a guess.

Put that on the same worksheet. At an assumed \$400–\$1,200 per engineer-day, spread over four years:

Steady workloads sharing the build effort	Build effort, per workload-hour	Total owned, per workload-hour	Against \$0.250 rented
1	\$0.21–\$0.62	\$0.23–\$0.68	loses, or breaks even at best
2	\$0.10–\$0.31	\$0.12–\$0.37	roughly a wash
3	\$0.07–\$0.21	\$0.09–\$0.27	wins, except at the pessimistic end
5	\$0.04–\$0.12	\$0.06–\$0.18	wins
10	\$0.02–\$0.06	\$0.04–\$0.12	wins comfortably

The crossover is at roughly two to three steady workloads. Below it, the build effort swamps everything and the cloud is cheaper. Above it, the hardware arithmetic starts to look like the a16z paper's "one-third to one-half," and keeps improving.

That is the first line in this book that is true about money, and it is worth being precise about how small it is. It is one workload, on one laptop, for ten minutes, against one provider's list price, with three of its four cost inputs assumed rather than measured. It does not capture high availability, replicated storage, backup, support contracts, on-call, an SLA, facilities, networking hardware, or any of the right-hand column of Chapter 18 — every one of which the rented price *includes* and the owned price does not.

So the honest statement of the result is narrow, and it is still the most useful sentence the platform has produced about cost:

On steady workloads, owning beats renting on the metal by roughly an order of magnitude — and the build effort is so much larger than the metal that it, not the hardware, decides whether owning pays. At this lab's effort level, the answer turns positive somewhere around the second or third workload.

Which is the same conclusion the deck reached, by a different route, and with one difference that matters: the deck put 70% of its saving in lines no invoice records. This puts the deciding line — people — in the open, where a CFO can argue with it.

The journey

BRIEFING · FEBRUARY 2026

The deck's last recommendations slide is a seven-step *"Sovereign Cloud Journey"*. It holds up better than its TCO slide, because each step is an action, not a number. Here are the steps, with where each one shows up in this book:

1. **Assess current state** — *"Audit cloud spend, identify lock-in points, catalog proprietary service dependencies."* Meridian's workload spreadsheet, and its dependency table (Chapters 1 and 2).
2. **Evaluate workload portability** — which workloads move easily, and which need rework. The steady ones first (Chapter 1).
3. **Start with non-critical workloads** — *"Build expertise on development, testing, or low-risk production systems."* The pilot, and Parts II to IV.
4. **Invest in CAPI & Kubernetes** — the factory (Chapter 4 and Part II).
5. **Consider hybrid approaches** — *"Maintain some managed services during transition to reduce risk."* Meridian is still renting, on purpose.
6. **Build platform engineering** — the blank fourth column of Vikram's table in Chapter 2: who could fix the boiler.
7. **Plan a 3–5 year timeline** — *"Full transformation requires patience but delivers lasting value."*

The deck lists four *"Key Success Factors"*: executive sponsorship, adequate budget, skills development and hiring, and incremental delivery. This book would add a fifth, ahead of the others: **measurement before savings**. Nothing in steps 1 to 7 produces a saving until someone has metered the same workload both ways.

A cloud owning culture

The campaign flyer promises *"Save \$Millions"* and names the means: *"Creating a Cloud Owning Culture."* The book agrees with the second half, and has spent eighteen chapters finding out what it means.

Owning a cloud is not owning servers. It is owning the work a landlord used to do out of sight: knowing what is true about the platform, and being able to show someone else how you know. At Meridian that took the form of rules, each written after a measurement proved an assumption wrong:

- Don't say *ready*; say *working*, and only once something has worked (Chapter 5).
- Anything we keep warm, we exercise (Chapter 9).
- Keep the before picture, and label it (Chapter 10).
- No check goes on a dashboard until someone has watched it fail on purpose (Chapter 11).
- A model gets what every service gets: an identity, a lock, and a meter (Chapter 12).
- An agent may advise; it may not act — until it has earned it, in writing (Chapters 15 and 16).

None of those rules appears on a TCO slide. All of them are part of the cost of owning, and all of them are where the margin is actually kept.

What to ask your vendors — including the one you become

When a sovereign-cloud programme starts, the organization becomes a supplier to itself. The same questions apply to both.

1. **Price.** How much notice do we get before a price change, and is there a cap?
2. **Leaving.** What does it cost to leave, in fees and in rebuilding, and in what format do we get our data back?
3. **Control.** Which settings we need can we *not* change, and who decides when that list changes?
4. **Jurisdiction.** Where are our data, our logs and our AI prompts processed — and under whose law?
5. **Metering.** Can we see tokens used, by model and by team, every month?
6. **Truth.** When your status page says *healthy*, what did you check, and how would we check it ourselves?

A platform team that cannot answer those six for its own platform has rebuilt the landlord, and moved it indoors.

* * *

MERIDIAN · THE SAME FRIDAY, TWO HOURS LATER

The spreadsheet had nine rows, the same nine as the worksheet, and three columns: *renting today*, *owning, estimated*, and *owning, measured*. The first column was nearly full; the finance team had the invoices. The second was full of ranges. The third was empty except for three cells Vikram had sent over: a build time, a reservation figure, and an energy number.

"So what goes in the board pack?" the CFO said. "They asked for a saving."

"They'll get one," Anita said. "Not from a consultant. From us." She pointed at the first row of the workload spreadsheet from January — the steadiest system Meridian had. "We run that on the pilot platform for a quarter, in parallel. We meter both sides — the invoice on one, and hardware, power and hours on the other. People's time goes in the model, not in a footnote. At the end of the quarter we have one number we can defend, for one workload."

"One workload isn't four point three million."

"No. It's the first line of it that's true." She closed the consultant's slide. "The board has been shown a lot of green this year. I'd like the first number we give them about the money to be one we measured."

The CFO typed the board-pack line himself. *Sovereign cloud: no savings claimed. One workload metered in parallel, Q4; first measured result in January.*

He read it back. "It's a very boring sentence."

"Good," Anita said. "Boring sentences are the ones that turn out to be true."

The ledger

- **The deck's model (briefing):** \$12.5M rented against \$8.2M owned over five years; a \$4.3M (34.4%) saving that is exactly the sum of four savings lines, with no line for anything owning adds. About 70% of the saving (\$3.0M) is in two lines no invoice records. The 5–7 year payback agrees with the five-year saving only at the expensive end of the deck's own \$2–5M investment range.
- **Depreciation (public record):** Microsoft (FY2022) and Alphabet (2023) moved servers from four- to six-year lives; Amazon moved part of its fleet from six to five years from January 2025, citing AI.
- **Owned waste (build log, 17 September 2026):** 22.42 GiB reserved against 7.45 GiB used on the management cluster; a control-plane VM using 1% of its CPU and 14% of its memory.
- **AI energy (build log):** qwen2.5:7b on an RTX 4050 laptop GPU at 6.4 W idle and 38–42 W generating; 34–37 tokens/s; 1.19 J per token; 0.33 kWh per million generated tokens, GPU only.
- **Effort (build log, 18 September 2026):** 88 commits from one committer, Mahipal, over seven months, 70 of them co-authored with an AI coding assistant, on 18 distinct days — a lab, not an operation. (The figure moved while this book was being written, because writing it added commits.)
- **Metered, one workload both ways (build log, 18 September 2026):** an nginx deployment on the target cluster, 30 of 30 probes 200 over ten minutes. It holds 48.6% of the host's memory by reservation while using 2.29% of its processor. Rented equivalent, at published list prices: \$0.250 per workload-hour. Owned metal, on stated assumptions: \$0.02–\$0.06. Owned including the build effort: \$0.23–\$0.68 at one workload, falling below the rented price at roughly two to three steady workloads.
- **Not claimed:** any saving for Meridian; any figure for high availability, backup, support, on-call or an SLA, none of which the owned side includes and all of which the rented price does.

Ask your team

1. For our biggest workload, what would it cost to meter it both ways, in parallel, for one quarter?
2. In any savings model we have been shown, which lines appear on an invoice — and which lines for new costs are missing?
3. What do we reserve against what we actually use — in the cloud, and on the hardware we already own?

Open the repo

- `kubect1 --context kind-cluster2 describe node cluster2-control-plane` — reserved against allocatable, under *Allocated resources*.
- `docker stats --no-stream` — what the clusters actually use.
- `kubect1 --kubeconfig target-cluster-kubeconfig top nodes` — what a target-cluster machine uses of what it was given.
- `nvidia-smi --query-gpu=power.draw --format=csv -lms 200`, alongside ollama's `eval_count` and `eval_duration` — energy per token.
- `git log --format='%B' | grep -c 'Co-Authored-By'`, `git shortlog -sn`, and `git log --format=%ad --date=short | sort -u | wc -l` — how the platform was built, and on how many distinct days.
- `measurements/ch19-parallel-metering-20260918.tsv` and `measurements/ch19-meter.sh` — the metering run's raw samples and the script that produced them; `measurements/README.md` holds the full assumptions table and the dated price sources.
- `06-symposium/cost-analyzer.yaml` — the FinOps agent, and why it does not install (Chapter 16).
- The strategy deck, slides 16, 18 and 19; the unbranded sovereign-cloud flyer, version 2. (Neither is a repository file; see Appendix H — Sources.)

Epilogue — Own Your Cloud

Open standards. Your data, your platform, your rules — always.

— Sovereign Cloud campaign flyer, March 2026

MERIDIAN · A MONDAY IN SEPTEMBER 2026 · THE CTO'S OFFICE

The invoice came on the first working day of the month, as it always did.

Anita read the slope before the total, as she always did. It had not changed much. Nothing in production had moved. The pilot was a pilot, and she had refused, more than once, to let anyone call it anything else. Twelve months ago she had written a question above the total. The answer, so far, was a platform on a single rack in the basement and a folder of measurements.

Under the invoice was a proof from the marketing team: a one-page flyer for the internal launch, due out on Friday. Across the top it said *Sovereign Cloud*. Down the page were six promises, each with a short line of copy, and a band at the bottom: *Save \$Millions — Creating a Cloud Owning Culture*. Across the bottom right corner, in the largest type on the page: *Own Your Cloud*.

She uncapped a pen and read it the way she had learned to read everything this year: one claim at a time, against something measured.

AI-Driven Intelligence — *"Built-in AI models and automated workflows that learn, adapt, and act."*

She underlined *act*. In September she had written the rule for the agent pilot herself: the agent may triage alerts; it may not act on them. The models ran on hardware in the building. The smallest had once invented a cluster's output when its tool failed quietly. The better one had got a planted fault right three times out of three. She crossed out *and act* and wrote in the margin: *and advise — every answer checked*.

Service Mesh Visualization — *"Live visual map of your service mesh — traffic, health and topology."*

She remembered the evening before the board meeting: a map labelled *live* on which fourteen of twenty-one things did not exist, thirteen of them drawn green. The mesh itself was real; failover between the clusters had been measured by status code, thirty calls out of thirty. She wrote: *live where checked. Every box traceable to a lookup that could come back empty*.

Zero Vendor Lock-In — *"Open standards. Your data, your platform, your rules — always."*

That one was nearly true, and she let it stand. Every component was open source. But she added a note underneath, because the year had earned it: *no vendor lock-in is not the same as no dependency*. A young provider had reported a working cluster as unavailable. A moving version tag had crashed a service thousands of times. An upgrade had taken a week. *Our rules — and our maintenance*.

Sovereign & Secure — *"Data center-grade security. Meets residency, privacy and compliance mandates."*

The models were local. A new cluster booted without touching the internet. The mesh's root of trust had been generated in-house. And the model gateway had shipped without a lock, the observability stack was still downloaded from a public repository at install time, and some demonstration keys

were sitting in version control. She wrote: *sovereign: yes, audited quarterly. Secure: in progress — list attached.*

Autonomous Operations — *"Self-healing infra with auto-scaling and hands-free management."*

She struck out *hands-free* with one line, and did not replace it.

Full Observability — *"Deep telemetry across compute, network and storage layers."*

Deep, yes. But the model traffic — the traffic the board asked about most — had been counted in a metric nobody displayed, with no model name and no token count. She wrote: *add: tokens, by model. That's the bill.*

Then she came to the band at the bottom, and stopped.

Save \$Millions.

Nobody had measured that — not here. The platform could build a cluster from a written order with both machines ready in a median of 33.7 seconds, and tear it down in 45.7. It could keep a spare cluster warm and hand it over in under half a second. Those were measurements.

And in September the pilot had finally metered one workload both ways: a single web service, running steadily, priced against a public list price for the same two machines. Owning the metal came out roughly an order of magnitude cheaper per hour. Then the engineering time went on the same sheet, and the picture turned over: at one workload, owning lost. The line crossed somewhere around the second or third steady workload, and after that it ran away.

Anita had read that result twice, and what she took from it was not the ratio. It was which line decided. The hardware was a rounding error next to the people, and the people were the line the consultant's slide had left out.

One workload, on one laptop, against one price list, with three of its four inputs assumed. It was the smallest true thing in the folder, and it was the first line about money that was hers rather than someone else's.

She drew a box around *Save \$Millions* and wrote inside it: *one workload says yes, past the third. Ours, on our invoice: not yet measured.*

* * *

Vikram knocked on the open door, saw the flyer, and winced at the amount of ink on it.

"Marketing wants it back by Wednesday," he said.

"They'll have it." She turned the page round so he could read the margins. "Most of it is right. It's just written as if we'd finished."

He read it from the top. "You kept *Open standards*."

"It's true."

"And the band at the bottom?"

"*Creating a Cloud Owning Culture* stays." She tapped it. "But I want to know what we mean by it before we print it. I don't think we mean the hardware."

"What, then?"

Anita thought about the year. A dashboard that said *Ready* about a machine that couldn't work, and the rule that followed: say *working*, and only when something has worked. A demo whose success didn't depend on the step before it, kept on purpose and labelled as the before picture. A failover check that could not fail, and the rule that no check reaches a dashboard until someone has watched it fail. A warm standby nobody exercised. An AI control point the AI didn't use.

"Owning isn't the servers," she said. "Anybody can buy servers. Owning is knowing what's true about them, and being able to show someone else how you know. Renting lets you skip that. It's most of what the rent pays for." She capped the pen. "The culture is the checking."

Vikram looked at the flyer for a while, at the crossings-out and the boxes and the margin notes that now took up more of the page than the copy.

"So what's left, once you've taken out everything we haven't proved?"

Anita turned the page back round, and put her finger on the line in the corner that she had not marked at all.

Own Your Cloud.

The ledger

The year, measured — collected from the chapters of this book.

- **The factory:** a two-machine Kubernetes cluster built from a written declaration, both machines ready in a median of 33.7 s and torn down in 45.7 s, on open-source Cluster API and KubeVirt, on owned hardware (Chapters 4, 7 and 8).
- **The fabric:** failover between two clusters measured by status code, 30 calls out of 30, with one label replacing hand-written cross-cluster plumbing (Chapters 10 and 11).
- **The brain:** AI agents on models hosted in the building; a planted fault triaged correctly in 3 runs out of 3; a smaller model that invented output when its tool failed quietly (Chapters 13–16).
- **The gauges:** a worker reported `Ready` that could not run work; a cluster reported unavailable while it served workloads; a map labelled `live` drawing fourteen things that did not exist; a failover check that printed success during a total outage (Chapters 4, 5, 11 and 12).
- **The money:** one workload metered both ways on the lab platform (Chapter 20) — rented \$0.250 per workload-hour at list price, owned metal \$0.02–\$0.06, and owning only paying for its own build effort at roughly the second or third steady workload. Against **Meridian's** own invoice: still not measured. That is the difference the chapter is careful about.

Ask your team

1. Which of the claims we make about our own platform have we measured — and which are we still printing?
2. When did each of our important checks last fail, on purpose, in front of someone?
3. If we owned it all tomorrow, who would know what is true about it — and how would they show us?

Appendices — The Owner's Reference

Every check asserts on a value read back from the cluster or an HTTP response code. No check passes on the absence of an error, and no screenshot counts as evidence.

— `07-istio-advanced/verify.sh`

Eight appendices, for readers who want to check the book rather than take its word. Each was checked against the repository and the running platform on 17 September 2026 (Appendix G, on 18 September). Where the repository's own documentation disagrees with what was running, both are shown.

- **A. The repo, stage by stage** — what each directory does, and the `make` targets that matter.
- **B. Addresses that must not move** — the network address plan, as reserved and as found.
- **C. The failure-mode catalog** — every documented failure, grouped by what you would see.
- **D. The measurement protocol** — the figures as last measured, and the rules for measuring them again.
- **E. The ten labs** — the agent-platform labs, with their original honesty markers and what this book re-checked.
- **F. A glossary for the boardroom** — one plain-English sentence per term.
- **G. Key files, quoted** — the load-bearing code and manifests, verbatim, for readers without the companion repository.
- **H. Sources** — what the book draws on, and which of it a reader can check.

Appendix A — The repo, stage by stage

BUILD LOG · 17 SEPTEMBER 2026

The numbered directories run in order. Four prerequisites are **not** created by any script, and are easy to forget when reproducing the platform: the two Kind clusters themselves (`cluster1` and `cluster2`); KubeVirt (v1.9.0) and CDI (v1.66.0) on `cluster2` ; and the `ubuntu-noble-dv` base disk image the golden-image bake starts from.

00-prereqs

`install-clusterctl.sh` installs `clusterctl` , the Cluster API command-line tool. *Target:* `make prereqs` .

01-metallb

Installs MetalLB, which gives Services on the laptop's Docker network real IP addresses, and configures the address pool for `cluster2` (now `.211-.225` , widened by `07-istio-advanced`). **Do not run it against `cluster1`** : it tries to substitute a placeholder that does not exist in its configuration

file, so it would advertise `cluster2`'s range on the shared network (Appendix B). *Target:* `make metallb`.

02-capi-init

`init-management-cluster.sh` turns `cluster2` into a cluster factory. It installs Cluster API v1.14.2, the KubeVirt infrastructure provider (CAPK) v0.11.2, and the k3s bootstrap and control-plane providers v0.3.0. It then replaces a dead proxy image those providers ship (Chapter 19), and runs `scripts/configure-kubevirt-perf.sh`, which applies the CPU setting that cut eight seconds from every VM start (Chapter 8). That auto-run can silently do nothing — not, as first diagnosed, because it races KubeVirt's start-up, but because the script's own readiness waits can outlast whatever time limit its caller imposes, and it is killed before reaching that last step (Chapter 19). Check afterwards that the setting landed. *Targets:* `make capi-init`, `make kubevirt-perf`.

03-target-cluster

The cluster declarations. `target-cluster-warm.yaml` is the default: seven objects in 258 lines, parallel worker boot, fixed demonstration certificate authorities in `warm-ca/` (Chapter 4). `target-cluster-lite.yaml`, `target-cluster.yaml` and `target-cluster-parallel.yaml` are older, slower paths on the `:latest` image. `generate-cluster.sh` and the `.tmpl.yaml` files render variants. The golden-image tooling sits at the top level (`bake-common.sh`, `bake-golden-image*.sh`, `build-containerdisk.sh`) and in `scripts/` (`seed-cluster-secrets.sh`, `gen-warm-ca.sh`, `fix-registry-hosts.sh`). *Targets:* `make target-cluster` (which runs `registry-fix`, then `ensure-warm-image`, then seeds the secrets and applies the warm manifest); `make clean`.

Caution: `ensure-warm-image` checks that the image *tag* exists, not what is inside it. After the September upgrade, the `:latest` and `:preinit` images still carried the old k3s binary (Chapter 19).

04-verify

`verify-cluster.sh` waits for the target cluster's virtual machines and checks its health. It also refreshes the agents' copy of the target cluster's credentials. *Target:* `make verify`.

05-istio

The "before" picture, kept on purpose (Chapter 10). `install-istio-ambient.sh` installs Istio 1.28 on the target cluster — outside Istio's tested range on Kubernetes 1.37. `cross-cluster-demo.sh` is the eight-step hand-plumbed demo, and `cross-cluster-mindmap.html` is its diagram. *Target:* `make istio`.

06-symposium

The agent platform. `install-symposium.sh` pins chart 0.10.75. The three serving agents are `cluster2-agent.yaml`, `target-cluster-agent.yaml` and `mesh-sre-agent.yaml`. Their extra permissions are `agent-kubevirt-rbac.yaml` and `agent-istio-rbac.yaml`, now bound to a group rather than one account (Chapter 15). Their network rule is `agent-egress-ollama.yaml`. The `fix-`

*.sh scripts each answer a pitfall in Appendix C; labs/ holds the ten labs (Appendix E); demo-mesh-sre.sh is the recorded agent demo. cost-analyzer.yaml and incident-responder.yaml cannot be applied on the current chart (Chapter 16). *Targets:* make symposium-install, make symposium-lb.

07-istio-advanced

The "after" picture: five acts on Istio 1.31 across cluster1 and cluster2 — Gateway API, waypoints, multicluster, the AI gateway, and observability — plus verify.sh (22 checks) and clean.sh (Chapters 11 and 12). *Targets:* make istio-adv-prereqs, istio-adv-install, istio-adv-act1 to act5, istio-adv-demo, istio-adv-verify, istio-adv-clean.

scripts/

The measurement instruments phase-timings.sh and time-to-ready.sh (Appendix D), along with configure-kubevirt-perf.sh, fix-registry-hosts.sh, seed-cluster-secrets.sh, gen-warm-ca.sh, render-cluster.sh, render-demo-video.sh and import-base-images.sh.

ui/

The web console: a React frontend and a Go backend (port 8080), with SRE, AI and Visual modes (Chapters 12 and 13). run-ui.sh starts both for development. The production manifest in ui/k8s/ was not deployed on 17 September: the kubeui namespace did not exist. *Targets:* make ui, make ui-build.

docs/, and the top level

docs/sub-60s-cluster-strategy.md records the cluster-speed work and its dead ends (Chapters 7 and 8). docs/warm-pool-strategy.md covers the standby (Chapter 9), and docs/demo/ holds the recorded agent demo. At the top level: CLAUDE.md, the platform's operating memory and pitfall list; Sovereign_Cloud_Agentic_Strategy.md, the agent strategy (Chapters 3, 14 and 16); demo.sh and show-cluster.sh, the original demo, whose narration has drifted (Chapter 5); and target-cluster-kubeconfig, the target cluster's credentials file.

The make targets that matter

The Makefile defines 65 targets. These are the ones this book used.

Target	What it does	Note, 17 September 2026
<code>make all</code>	<code>prereqs</code> → <code>metallb</code> → <code>capi-init</code> → <code>target-cluster</code>	Assumes the four unscripted prerequisites above
<code>make target-cluster</code>	Builds the default warm cluster	33.7 s to both nodes ready (Appendix D)
<code>make clean</code>	Deletes the target cluster	45.7 s until every object is gone; the machines stop at 35.6 s (Chapter 17)
<code>make verify</code>	Checks the target cluster and refreshes agent credentials	
<code>make kubevirt-perf</code>	Applies the KubeVirt CPU setting	Cluster state, not manifest state: re-run after rebuilding <code>cluster2</code>
<code>make phase-timings</code>	Per-phase breakdown of a cluster build	Counts nodes, not names — see Appendix D
<code>make registry-fix</code>	Maps the registry alias for <code>containerd</code>	See Appendix C.5
<code>make istio</code>	Istio 1.28 on the target cluster	The legacy path; unsupported on Kubernetes 1.37
<code>make symposium-install</code> <code>/ symposium-lb</code>	Installs the agent platform; gives its Services addresses	Re-run <code>symposium-lb</code> after any serving Deployment is regenerated
<code>make symposium-warm</code>	One-shot model warm-up	Its help text still names llama3.2
<code>make symposium-demo-agents</code>	Applies the April cost and incident personas	Fails on the current chart (Chapter 16)
<code>make ui</code>	Starts the web console	
<code>make istio-adv-demo</code> / <code>istio-adv-verify</code>	Runs the five acts / the 22 checks	22 passed on 17 September; the Act 3 checks test configuration, not failover (Chapter 11)

`make help` lists most targets, not all. It omits `kubevirt-perf`, `phase-timings` and every `istio-adv-*` target, and it still describes the default cluster as a "*~40s target*."

Appendix B — Addresses that must not move

BUILD LOG · 17 SEPTEMBER 2026

Both Kind clusters live on one Docker network, `172.18.0.0/16`, and both run MetalLB announcing addresses on that same network. If their pools overlapped, two clusters would answer for one address, and traffic would go to whichever answered first. So each cluster has its own pool, and every Service that needs a stable address pins one inside its own pool.

Pools: `cluster1` has `172.18.255.200-.210`, installed by `07-istio-advanced/00-prereqs/`. `cluster2` has `172.18.255.211-.225`, originally `.211-.220`, widened by `07-istio-advanced`.

Address	Reserved for	Cluster	Allocated, 17 Sept	Files naming it
.200	httpbin-lb — legacy 05-istio demo	cluster1	No	9
.201	Act 1 edge gateway	cluster1	Yes	4
.202	Act 3 east-west gateway	cluster1	Yes	2
.203	Act 4 AI gateway	cluster1	Yes	3
.204	Kiali	cluster1	Yes	3
.211	Web console frontend	cluster2	No — not deployed	10
.212	Symposium API and dashboard	cluster2	Yes	4
.213	cluster2-agent	cluster2	Yes	4
.214	target-cluster-agent	cluster2	Yes	4
.215	Target cluster API server	cluster2	Yes	18
.216	Target-cluster nginx proxy — legacy 05-istio , only while it runs	cluster2	No	5
.217	Security agent	cluster2	No	4
.218	cost-analyzer	cluster2	No	3
.219	incident-responder	cluster2	No	3
.220	Host Ollama load balancer (optional)	cluster2	No	4
.221	Act 3 east-west gateway	cluster2	Yes	3
.222	mesh-sre-agent	cluster2	Yes	3

"Files naming it" counts shell, YAML, Go, TypeScript, env and Makefile files outside `book/` . Seven of the seventeen reserved addresses were not in use on 17 September, and every one of them is still written into between three and ten files.

The node addresses on the same network matter too. `172.18.0.1` is the host, where the models run. `.2` , `.3` and `.4` were `cluster1-control-plane` , `cluster2-control-plane` and the image registry — in that order *on that day*. Docker hands those addresses out in the order containers join the network, which is why `172.18.0.2:5000` , the registry address baked into image references, is an alias that only containerd resolves (Appendix C.5).

Pool discipline

- **Never overlap pools.** One cluster's installer must never advertise another's range.
- **Pin, don't request.** Every Service that anything else depends on names its address in an annotation. Both MetalLB annotation spellings appear in the repository,

`metallb.universe.tf/loadBalancerIPs` and `metallb.io/loadBalancerIPs` ; both worked on 17 September.

- **Change all three places.** The project's rule: never reassign an address without updating its source file, `sympozium-lb-setup.sh` and `run-ui.sh`.
- **A regenerated Service forgets its address.** When the agent platform regenerates its serving Deployments, their Services come back without MetalLB addresses; re-run `make sympozium-lb`.
- **An idle reservation is still a reservation.** `.216` stays free for the legacy demo; when `cluster2` needed another address, the pool was widened instead (Chapter 10).
- **Retire addresses on purpose.** A reservation nobody uses, named in ten files, is a future collision waiting for someone who reads only the table.

Appendix C — The failure-mode catalog

BUILD LOG · FEBRUARY–SEPTEMBER 2026

The platform's operating notes list 23 pitfalls, and the chapters of this book found more. They are grouped here by **what you would see**, because that is what an engineer has at 2 a.m. The most expensive group is not the crashes. It is the failures that look like success.

C.1 A crash loop with exit code 2 and no log output

One sidecar — the agent platform's `web-proxy` — produced this identical symptom from four different causes.

Cause	When	Fix or status
The image could not run with a read-only root filesystem	Chart 0.10.38	Fixed upstream in 0.10.47; <code>fix-web-proxy-rootfs.sh</code> retired
The sidecar image was pinned to a moving <code>:latest</code> tag and drifted to an incompatible version (observed at about 5,680 restarts)	0.10.47 era	<code>fix-web-proxy-image.sh</code> pins it to the chart version and purges the cached image
The chart turned on NATS authentication, but the serving sidecar was never given credentials	0.10.57	<code>nats.auth.enabled: false</code> in <code>06-sympozium/values.yaml</code>
Two long-running crash loops (600+ restarts), cause not diagnosed	Before the 0.10.75 upgrade	Cleared when the post-upgrade re-pin regenerated the Deployments

C.2 The script reported success; nothing happened

Cause	Where	Fix or status
A provider shipped a proxy image that no longer exists; readiness waits timed out into an <code> true</code> fallback, so setup reported success	<code>make capi-init</code> (Chapter 19)	Image patched automatically in <code>init-management-cluster.sh</code>
The KubeVirt CPU setting's applying step never ran because the setup script itself was killed by an external time limit mid-readiness-wait (corrected 18 September; first misdiagnosed as a race with KubeVirt's own start-up)	<code>configure-kubevirt-perf.sh</code> auto-run (Chapter 19)	Check <code>supportContainerResources</code> after any <code>capi-init</code> ; re-run <code>make kubevirt-perf</code>
A documented <code>K3S_VERSION</code> override never reached the image being baked	<code>bake-common.sh</code> (Chapter 19)	Both version literals updated and commented
Image checks tested whether a tag existed, not what it contained	<code>ensure-warm-image</code> , <code>pre-pull</code>	<code>:warm</code> rebaked; <code>:latest</code> and <code>:preinit</code> still stale
Built-in SkillPacks "silently dropped" at install — really the chart applying them before its own admission webhook was serving	<code>helm install</code>	<code>fix-missing-builtin-skillpacks.sh</code>
Agent runs created in quick succession raced into a terminal <code>Failed</code> , which the controller never retries	Serving re-pin	Delete the failed run and let the controller recreate it
A failover check counted any non-empty body as success, so a total outage printed <code>success=20 failed=0</code>	Act 3 <code>run.sh</code> (Chapter 11)	Not fixed; the console's failover action counts HTTP 200

C.3 The status says healthy; it is not

Cause	Where	Fix or status
The worker node was <code>Ready</code> but tainted as uninitialized, so it accepted no workloads	Target cluster (Chapter 5)	Not fixed
A ghost node baked into the golden image satisfied "2 nodes Ready" before the worker joined	Both timing scripts, and (found 18 September) the UI backend's own <code>targetNodesReady</code> — used by both the on-demand deploy path and the warm pool (Chapter 7, Chapter 9)	Measure/count by name (Appendix D); fixed in <code>pool.go</code> 18 September, standalone scripts still not fixed
A missing worker timestamp was treated as zero, so "both nodes" meant the control plane only	<code>phase-timings.sh</code> (Chapter 7)	Not fixed
The deterministic permission check tested an identity running pods no longer used	Agent RBAC after 0.10.75 (Chapter 15)	Bindings moved to a group; check the live pod's account
A topology view labelled <code>live</code> drew 14 nodes that did not exist, 13 of them healthy	Visual mode (Chapter 12)	Not fixed
A tenant's listener was refused, but the act printed its status instead of asserting it	Act 1 ListenerSet (Chapter 11)	Not fixed
"Locality-aware routing prefers in-cluster endpoints" was printed without comparing counts; measured traffic split 30/30	Act 3 (Chapter 11)	Not fixed

C.4 The status says broken; it is not

Cause	Where	Fix or status
Cluster API reported <code>Available=False</code> (<code>WaitingForKthreesServer</code>) for a working cluster	k3s control-plane provider v0.3.0 (Chapter 4)	Not investigated

C.5 Unreachable from here, reachable from there

Cause	Where	Fix or status
<code>172.18.0.2:5000</code> is a registry alias only containerd can resolve; plain HTTP clients get connection refused	Web console Registry tab	Backend dials a real address, displays the alias
A host-network probe looked for Ollama on <code>127.0.0.1</code> ; the redirect that fixed it was lost on reboot	Symposium node probe	Redirect re-asserted every 60 s by a <code>DaemonSet</code>
A VM's pod address is reachable from its node but not from other pods, costing a 10 s timeout	Worker joining the control plane (Chapter 8)	<code>advertise-address</code> set on the control plane
Remote-cluster credentials embedded <code>127.0.0.1:<port></code> , so the peer sat at timeout while all else looked healthy	<code>istioctl create-remote-secret</code> on Kind	Pass <code>--server https://<node-ip>:6443</code> ; Act 3 asserts <code>syncd</code>
<code>host.docker.internal</code> does not resolve inside Kind pods	Model endpoint	Use <code>172.18.0.1</code> through an in-cluster shim Service
Both Kind clusters number pods from <code>10.244.0.0/24</code> , so a pod address means different machines in each	Cross-cluster calls (Chapter 10)	Hand-plumbed proxy (before); east-west gateways (after)

C.6 The model gave a confident, wrong answer

Cause	Where	Fix or status
Stale target-cluster credentials made the tool fail TLS quietly; llama3.2 invented plausible output	<code>target-cluster-agent</code> (Chapter 15)	<code>refresh-target-kubeconfig.sh</code> , auto-run by verify and deploy paths
Chat runs got no tool sidecar, so the model gave up or improvised	Symposium 0.10.38	Skills declared on the <code>Agent</code> ; the webhook workaround is now redundant
A refusal rule in the system prompt was not reliably obeyed	<code>mesh-sre-agent</code> , qwen2.5:7b (Chapter 14)	Enforce below the model, not in the prompt
A tool denied by policy was used anyway	Lab 04 (Chapter 14)	Not fixed on 0.10.75
After the per-run identity change, the agent reported a plausible "permission issue"	<code>mesh-sre-agent</code> after 0.10.75	Group binding (C.3)

C.7 503 on every request, with every component healthy

Cause	Where	Fix or status
An in-mesh route pinned traffic to non-global backends, overriding cross-cluster routing	Act 2's <code>echo-internal-split</code> during Act 3	Act 3 deletes the route; still reproduces on Istio 1.31 (Chapter 11)
A waypoint resolved only local endpoints	Istio 1.30.3	Did not reproduce on 1.31; the console's opt-out label is now harmless (Chapter 11)

C.8 It waits forever, though the thing is done

Cause	Where	Fix or status
<code>kubectl wait --for=condition=Accepted</code> reads top-level conditions; a route's live under <code>status.parents</code>	HTTPRoute	<code>wait_route_accepted</code> in <code>07-istio-advanced/lib/common.sh</code>
<code>helm upgrade --wait</code> polls an intentionally suspended example MCPServer	Symposium upgrades	Drop <code>--wait</code> ; verify by hand
Gateway API CRDs exceed the annotation size limit of client-side apply	Gateway API install	Apply with <code>--server-side</code>

C.9 The screen is empty, or shows the wrong thing

Cause	Where	Fix or status
The dashboard's namespace picker defaults to <code>default</code> , where nothing lives	Symposium dashboard	Switch to <code>symposium-system</code> (stored per browser)
The GPU detector shells out to <code>nvidia-smi</code> , which is not in its container, so it reports the integrated GPU	Model-fit hardware view	<code>fix-llmfit-nvidia-smi.sh</code> replay shim
The agent drop-down queries a resource type the chart no longer ships	Web console AI mode	Falls back to the default agent only; not fixed
Model calls through the AI gateway are absent from the mesh request metric and from Kiali	Observability (Chapter 12)	Not fixed

C.10 It changed its name or its shape

What changed	Where	Consequence
<code>SymposiumInstance</code> stopped being reconciled, then stopped shipping	Symposium 0.10.38 → 0.10.47	Agents must be <code>Agent</code> resources
Skill sidecars moved from one shared account to one account per run	Symposium 0.10.75	Permissions bound to a single account silently stopped applying
<code>XListenerSet</code> graduated to <code>ListenerSet</code>	Gateway API v1.5	Most write-ups use the old name, which fails
<code>istioctl verify-install</code> was removed	Istio 1.30	Use <code>istioctl analyze</code>
<code>--server-side</code> started requiring a value	Helm 4	A bare <code>--server-side --force-conflicts</code> fails
The three GatewayClasses name their generated Services differently; the Prometheus container is <code>prometheus-server</code>	Istio add-ons	Look Services up by label; name the container explicitly
VM networking moved from masquerade to bridge, in the same commit as the diagram describing masquerade	March 2026 (Chapter 10)	The mind map and README diagram still describe the old network

Appendix D — The measurement protocol

BUILD LOG · JUNE–SEPTEMBER 2026

D.1 The figures, as last measured

Measurement	Value	When, and how
Cluster bring-up, both nodes ready	33.7 s median (30.4 / 35.1 / 33.7)	17 September, by node name, warm manifest, Kubernetes 1.37
Control plane alone	24.7 s median	Same runs
Teardown	Delete command returns at 10.13 s; machines stop at 35.6 s; last object gone at 45.73 s	19 September, N=5 (spread 1.24 s). Supersedes an 18.9 s figure from a single run on 17 September
Standby claim	467 ms median (1.059 / 0.467 / 0.465)	18 September, by name-based readiness check, warm-image standby; was 194 ms in June
Model load, cold vs warm	2.3–2.6 s cold; 0.14–0.17 s warm	17 September, GPU, qwen2.5:7b and llama3.2 (Chapter 9)
Local model energy	1.19 J per generated token; 0.33 kWh per million tokens (GPU only)	17 September, qwen2.5:7b (Chapter 20)
Cross-cluster failover	30 of 30 HTTP 200, all from cluster2	17 September, Istio 1.31 (Chapter 11)

And the history behind the headline number:

Stage	Median	Runs	Spread
Baseline (June)	50.0 s	52.7 / 50.0 / 46.6	6.1 s
+ advertise-address	42.7 s	42.7 / 42.8 / 41.2	1.6 s
+ supportContainerResources	34.5 s	34.5 / 34.2 / 34.7	0.5 s
After the Kubernetes 1.37 upgrade	33.7 s	30.4 / 35.1 / 33.7	4.7 s

Notice the last column. The spread that had collapsed to half a second is back to 4.7 seconds after the rebuild. With three runs, a change smaller than about five seconds cannot be seen.

D.2 The instrument

scripts/phase-timings.sh attributes a build's wall clock to its phases, read-only and without logging in to the guests. Its sources, in its own words:

- Cluster API and CAPK object creation timestamps — the controller chain.
- VM instance phase transitions — scheduling to running.
- The launcher container's start time — disk unpack and hypervisor start.

- The guest serial console (`kubectl logs -c guest-console-log`) — kernel, initramfs, systemd, cloud-init, and k3s start.
- Node `Ready` transitions — the metric itself.
- A host-side probe loop — the API server's first TCP accept, and its first healthy answer.

Guest console timestamps count seconds since boot, not time of day. The script anchors them to the wall clock using cloud-init's own log lines, which carry both: *"running 'init' at Wed, 09 Sep 2026 06:44:19 +0000. Up 6.81 seconds."* Subtract the uptime from the date, and every console line lands on the same clock as the Kubernetes objects.

`./scripts/phase-timings.sh --collect-only` rebuilds the table for whatever is running now, and changes nothing. `--runs N` deploys and measures `N` clean runs.

D.3 The rules

1. **Instrument before you optimise.** Two dead ends each cost a full image bake to disprove before the phases were measured (Chapter 8).
2. **Name what you count.** Wait for the control-plane node and the worker node *by name*, never for "two nodes". A stale node object inside the golden image satisfied the count (Chapter 7).
3. **Missing is not zero.** An absent timestamp must fail the measurement, not add nothing to it (Chapter 7).
4. **Publish every run and the spread, not just the median.** With a six-second spread, three runs cannot resolve a lever smaller than about five seconds. Use seven or more runs, a quiet host, or both.
5. **When a number improves, check the instrument first.** The 23.9-second figure was the control plane alone (Chapter 7).
6. **Keep the raw data.** The raw files behind the 9 September phase table no longer exist to re-check.
7. **Judge a request by its status code, never its body.** An error page is a body (Chapter 11).
8. **Watch every check fail once, on purpose,** before it goes on a dashboard (Chapter 11).
9. **Measure use against reservation.** On owned hardware, reservations set capacity (Chapter 20).
10. **For models, separate cold from warm, and meter energy against tokens.** Read the model server's own `load_duration`, and sample GPU power every 200 ms alongside its `eval_count` (Chapters 9 and 20).
11. **Record what you changed, and put it back.** Every live experiment in this book ran with a restore step, and was checked afterwards.

D.4 Measuring bring-up by name

The repository's two original timing scripts, `time-to-ready.sh` and `phase-timings.sh`, still wait for a node *count*, which a ghost node can satisfy before the worker joins. The procedure below is now committed as `scripts/time-to-ready-by-name.sh` (`make time-to-ready-by-name`), which waits by

name, ignores and *names* any node matching neither pattern, and fails loudly on a missing timestamp rather than reporting it as zero. The 33.7-second figure was measured this way, three times:

1. Delete the cluster, and wait until no virtual machine or VM instance named `target-cluster` remains.
2. Seed the fixed certificate and token secrets (`scripts/seed-cluster-secrets.sh`).
3. Start the clock, and apply `03-target-cluster/target-cluster-warm.yaml` .
4. Every half-second, fetch the target kubeconfig if you don't have it yet, and list nodes with their `Ready` condition.
5. Record, separately, the first moment a node whose name contains `-cp-` is `Ready` , and the first moment a node whose name contains `-workers-` is `Ready` .
6. Cross-check both against each node's `Ready` `lastTransitionTime` from the API. (*The committed script does not do this step — it was performed by hand for the 17 September figures. Either do it by hand, or treat the polled values as the measurement, which is what Chapter 17's drill does.*)
7. Afterwards, refresh the agents' copy of the target credentials (`06-symposium/refresh-target-kubeconfig.sh`).

Run it directly:

```
make time-to-ready-by-name RUNS=3
```

(The script previously required the manifest positionally before `--runs` , so that invocation failed with `manifest not found: --runs` . Fixed 19 September; either argument order now works.)

D.5 The recovery drill

Chapter 17's cycle — destroy, rebuild, deploy, serve — measured on 19 September 2026, N=5 after one discarded rehearsal, with the boundary table committed to `measurements/README.md` **before the first measured run**. Raw data: `measurements/ch17-recovery-drill-20260919-163611.tsv` ; script: `measurements/ch-recovery-drill.sh` .

Figure	Median	Spread
Destroy to first served request	81.96 s	3.63 s
Delete command returns	10.13 s	0.98 s
Teardown, last object gone	45.73 s	1.24 s
Build, both nodes Ready	34.39 s	4.24 s
Build, control plane alone	25.60 s	0.78 s
Workload to first HTTP 200	1.34 s	0.58 s

Pre-flight gates abort rather than adjust: the `:warm` image must already be cached (recording its image ID, not just its tag), `supportContainerResources` must be present at `cpu: "1"`, and no warm-pool controller may be running. The workload is `rancher/mirrored-library-busybox`, baked into the golden image — every run recorded zero `Pulling` events, so nothing was fetched from the internet inside the clock.

D.6 The parallel-metering run

Chapter 20's cost comparison was measured on 18 September 2026 with `measurements/ch19-meter.sh`: an `nginx` deployment on the target cluster, sampled every 20 seconds for ten minutes, recording the host container's CPU and memory, the GPU's power draw, the target control plane's CPU and memory, and one HTTP probe of the workload per sample. All 30 samples returned `200`.

The raw samples are kept at `measurements/ch19-parallel-metering-20260918.tsv`, and `measurements/README.md` holds the full assumptions table and the dated price sources. Two things could not be measured on this host and are therefore assumptions, not figures: whole-host power (`/sys/class/powercap/intel-rapl:0/energy_uj` is root-only on this kernel, and the battery reads zero on AC) and the hardware's purchase price.

Appendix E — The ten labs

BUILD LOG · JULY–SEPTEMBER 2026

The labs in `06-symposium/labs/` were written in July against Symposium 0.10.38. Each one is standalone, cleans up after itself, and shows real captured output. Their index marks every lab  or , and explains every  in the lab's own README. Those markers are preserved below. The last column is what this book re-checked on 0.10.75 and Kubernetes 1.37.

#	Lab	What it shows	Marker (July, 0.10.38)	Re-checked in this book
01	chat-serving	An agent as an OpenAI-compatible chat endpoint	✓ works	Serving endpoints answered on 17 September (<code>verify.sh</code>); live completions through all three agents on 16 September (Chapter 19)
02	agentrun	A one-shot task as a Kubernetes object	✓ works — needs explicit <code>agentId</code> , <code>model</code> , <code>sessionKey</code> , <code>skills</code>	Re-run; required fields confirmed by a server dry run (Chapter 13)
03	schedule	Timed agent work	⚠ prompt-only tasks work; skill-using tasks don't	The platform's warm-up schedule was found stopped (Chapters 9 and 14)
04	policy	Tool gating, egress, sandbox	⚠ explicit override blocked; default deny not enforced; network policy inert	Override still refused; denied tool still used; network policy <i>is</i> now enforced, but the sandbox policy selects no pods (Chapter 14)
05	mcpserver	External tool servers	✓ discovery works; ⚠ tool calls flaky with the model	Not re-checked
06	ensemble	Agents handing work to each other	✓ works — <code>spec.enabled: true</code> required	Re-run end to end and cleaned up (Chapter 16)
07	model-fit	Which models fit the hardware	✓ works	Re-run: 1,365 candidate models scored against the host (Chapter 16)
08	capstone	A real task under policy	✓ works, partially — an RBAC gap surfaces	Not re-checked
09	metrics	Tokens, tool calls and latency compared across models	✓ works	Not re-run; figures predate three chart upgrades (Chapter 15)
10	sre-triage	A scored diagnosis of a broken workload	✓ works — 3/3 root-caused an <code>ImagePullBackOff</code>	Not re-run (Chapter 15)

Two statements in the labs index are now out of date. It describes `SymposiumInstance` objects in the agent manifests, and that resource type no longer ships with the chart. And its conventions describe *"observed 0.10.38 behavior"*, three upgrades ago.

Appendix F — A glossary for the boardroom

- **Agent** (Symposium) — an AI assistant declared as a Kubernetes object: its model, its tools, its policy, and its briefing, all in one reviewable file.
- **AgentRun** — one piece of work given to an agent, recorded as an object with its result, its token usage and its duration.
- **Agentic control plane** — software that lets AI agents observe and operate infrastructure through the same permissions, policies and audit trail as people.
- **Air-gapped** — able to run with no connection to the internet.

- **Ambient mesh** — Istio's service-mesh design with no proxy inside each application; a shared component on each machine handles encryption and identity.
- **CAPK** — the Cluster API provider for KubeVirt: the plug-in that lets the cluster factory build machines as virtual machines.
- **CDI** (Containerized Data Importer) — the KubeVirt companion that imports and prepares virtual-machine disk images.
- **Cluster API (CAPI)** — open-source software that uses one Kubernetes cluster to build and maintain others from written declarations.
- **Container** — a packaged application with everything it needs to run, lighter than a virtual machine because it shares the host's operating system.
- **Control plane** — the part of a Kubernetes cluster that decides what should run where; the cluster's brain.
- **Declaration and reconciliation** — writing down what should exist, and software continuously making reality match it.
- **Depreciation** — spreading the cost of hardware over its useful life; the landlords themselves disagree on whether a server lasts five years or six.
- **East-west gateway** — the door one cluster's mesh uses to reach another's, carrying encrypted, identity-bearing traffic between them.
- **Egress** — traffic leaving a system; also the fee some providers charge for data leaving them.
- **Gateway API** — the Kubernetes standard for declaring how traffic enters and moves between services, with ownership split between platform and application teams.
- **Golden image** — a pre-built machine image with everything a cluster node needs baked in, so nothing is downloaded at boot.
- **Istio** — the open-source service mesh this platform uses.
- **k3s** — a small, single-binary distribution of Kubernetes.
- **Kind** — "Kubernetes in Docker": complete Kubernetes clusters running as containers on one machine; used here as the laptop's data centre.
- **KubeVirt** — open-source software that runs virtual machines as Kubernetes objects, beside containers.
- **Kubernetes** — the open-source system that runs and manages containers across many machines.
- **Local-LLM-first** — the rule that AI models run on hardware the organization controls, with hosted models used only where explicitly allowed.
- **Management cluster / workload cluster** — the factory, and the clusters it builds.
- **MetalLB** — software that gives Kubernetes Services real network addresses on ordinary networks.
- **mTLS** (mutual TLS) — encryption in which *both* sides prove their identity, not just the server.
- **NetworkPolicy** — a Kubernetes rule stating which workloads may talk to which, enforced by the network layer.

- **Ollama** — the open-source server that runs AI models on the platform's own GPU.
- **Repatriation** — moving workloads from a rented cloud back onto infrastructure the organization owns.
- **ServiceEntry** — an Istio record that tells the mesh about a service outside it; the "before" picture of cross-cluster discovery.
- **SkillPack** — a bundle of tools an agent may use, with the permissions those tools need.
- **Sovereign cloud** — infrastructure whose data, technology and operations stay under the control of the organization and the jurisdiction it chooses.
- **SPIFFE identity** — a standard, cryptographically verifiable name for a workload, such as `spiffe://cluster.local/ns/demo-apps/sa/trusted`.
- **Switching charges** — fees for moving to another provider; banned in the EU from 12 January 2027 under the Data Act.
- **Token** — the unit AI models read and write — roughly a short word — and the unit most AI is metered and billed in.
- **Total cost of ownership (TCO)** — the full cost of an option over its life: purchase, operation, people, and exit.
- **Waypoint** — an optional proxy that adds request-level rules (paths, methods, identities) to an ambient mesh, one per namespace.
- **Warm standby** — a pre-built spare kept running so it can be handed over instantly; fast, but it holds capacity idle.
- **ztunnel** — the ambient mesh's per-machine component: it encrypts traffic and carries identities, but sees connections, not individual requests.

Appendix G — Key files, quoted

Read from the companion repository on 18 September 2026, branch `upgrade/k8s-1.37-istio-1.31-symposium-0.10.75`.

The chapters' **Open the repo** blocks point at files in the companion repository. Readers who do not have it should not have to take the book's word for what those files say, so the load-bearing ones are reproduced here, verbatim, each captioned with its path.

Two conventions. Where a file is long, an excerpt is shown and the elision is marked `# ...`; nothing inside a quoted block is paraphrased or tidied. And where an SSH public key appears in a committed manifest, it is elided as `[public key elided]` — the key is real, and it is in the repository, but it is the author's and not the reader's business.

G.1 The warm cluster manifest — seven objects

Chapter 5 describes a cluster as a written declaration; this is that declaration. Seven objects, applied together, produce two virtual machines running a Kubernetes cluster. Chapters 7, 8 and 9 all measure this file.

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – header comment)
# WARM-boot variant of target-cluster.yaml – target time-to-ready <40s.
#
# Combines two ideas:
# 1. WARM IMAGE (BAKE_MODE=warm, image tag :warm). The golden image was baked
#    with a FIXED custom CA set (03-target-cluster/warm-ca/) and a fixed token,
#    with the API VIP in the serving-cert SAN. Its etcd DB, certs, and system
#    charts are fully initialized. First boot does NO `k3s server
#    --cluster-reset` and NO cert purge – it just starts k3s from warm state.
# 2. PARALLEL WORKER (from target-cluster-parallel.yaml). The worker boots
#    alongside the control plane via a static bootstrap secret + skip-preflight,
#    so it is not serialized behind the CP.
#
# — REQUIRED ORDER —————
# ./scripts/seed-cluster-secrets.sh          # seed ca/ccs/etcd/token FIRST
# kubectl apply -f 03-target-cluster/target-cluster-warm.yaml
#
# DEMO USE ONLY: fixed CA + token are committed (see warm-ca/).
```

The seven objects, in the order they appear in the file:

#	Kind	Name	What it is
1	Cluster	target-cluster	The cluster itself: pod and service address ranges, and pointers to the two objects below
2	KubevirtCluster	target-cluster	The infrastructure side: pins the API server's address to 172.18.255.215
3	KThreesControlPlane	target-cluster-control-plane	The control plane: k3s version, one replica, and the boot commands
4	KubevirtMachineTemplate	target-cluster-cp	The control-plane machine: 4 cores, 8 GiB, booting the :warm image
5	Secret	target-cluster-workers-bootstrap	The worker's static cloud-init, so it need not wait for the control plane
6	MachineDeployment	target-cluster-workers	One worker, with CAPI's stability preflight skipped
7	KubevirtMachineTemplate	target-cluster-workers	The worker machine: 4 cores, 6 GiB, same image

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – object 1 of 7)
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: target-cluster
  namespace: default
spec:
  clusterNetwork:
    pods:
      cidrBlocks:
        - 10.42.0.0/16
    services:
      cidrBlocks:
        - 10.43.0.0/16
  controlPlaneRef:
    apiVersion: controlplane.cluster.x-k8s.io/v1beta2
    kind: KThreesControlPlane
    name: target-cluster-control-plane
  infrastructureRef:
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
    kind: KubevirtCluster
    name: target-cluster
```

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – object 4 of 7)
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
kind: KubevirtMachineTemplate
metadata:
  name: target-cluster-cp
  namespace: default
spec:
  template:
    spec:
      virtualMachineBootstrapCheck:
        checkStrategy: none
      virtualMachineTemplate:
        metadata:
          namespace: default
        spec:
          runStrategy: Always
          template:
            spec:
              domain:
                cpu:
                  cores: 4
                memory:
                  guest: 8Gi
                ioThreadsPolicy: shared
              devices:
                disks:
                  - disk:
                      bus: virtio
                      name: systemdisk
                interfaces:
                  - bridge: {}
                    name: default
              networks:
                - name: default
                  pod: {}
              volumes:
                - containerDisk:
                    image: 172.18.0.2:5000/ubuntu-noble-k3s:warm
                    name: systemdisk
```

And the worker's static bootstrap — the object that lets the worker boot beside the control plane rather than behind it (Chapter 8):

```

# 03-target-cluster/target-cluster-warm.yaml (excerpt – object 5 of 7)
# Static worker bootstrap (identical to target-cluster-parallel.yaml). Embeds the
# fixed WARM token so the agent can join the moment the CP API answers. The k3s
# agent retries the server VIP until it is reachable.
apiVersion: v1
kind: Secret
metadata:
  name: target-cluster-workers-bootstrap
  namespace: default
type: cluster.x-k8s.io/secret
stringData:
  value: |
    #cloud-config

    write_files:
      - path: /etc/rancher/k3s/config.yaml
        owner: root:root
        permissions: '0640'
        content: |
          kubelet-arg:
            - cloud-provider=external
          server: https://172.18.255.215:6443
          token: f00dcafef00dcafef00dcafef00dcafe

    runcmd:
      - "mkdir -p /home/ubuntu/.ssh && echo '[public key elided]' >
/home/ubuntu/.ssh/authorized_keys && chown -R ubuntu:ubuntu /home/ubuntu/.ssh && chmod 700
/home/ubuntu/.ssh && chmod 600 /home/ubuntu/.ssh/authorized_keys"
      # `agent` MUST be positional: the baked /opt/install.sh parses only positional
      # args and ignores INSTALL_K3S_EXEC, so without it the worker starts the k3s
      # SERVER unit, which fails ~0.8s later and – because install.sh runs under
      # `set -e` – aborts the script before writing the bootstrap-success sentinel.
      - INSTALL_K3S_SKIP_DOWNLOAD=true INSTALL_K3S_EXEC='agent' sh /opt/install.sh agent && mkdir -
p /run/cluster-api && echo success > /run/cluster-api/bootstrap-success.complete
      - "systemctl daemon-reload && systemctl enable k3s-agent && systemctl start k3s-agent"

```

Note the `kubelet-arg: cloud-provider=external` in that worker configuration. It is the line behind Chapter 5's finding: the worker comes up `Ready` and carries a taint no controller ever clears, so it accepts no workloads.

The annotation that lets the worker start at all, rather than waiting:

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – object 6 of 7)
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineDeployment
metadata:
  name: target-cluster-workers
  namespace: default
  annotations:
    # Skip CAPI's "ControlPlaneIsStable" preflight so the worker MachineSet scales
    # up immediately and the worker VM boots in parallel with the control plane.
    machineset.cluster.x-k8s.io/skip-preflight-checks: All
```

G.2 advertise-address — the first of the two one-line fixes

Chapter 8's headline: two configuration lines cut the build from 50.0 to 34.5 seconds. This is the first, with the comment that explains it, quoted whole because the comment is the evidence.

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – inside KThreesControlPlane.preK3sCommands)
# Advertise the API VIP, not this VM's pod IP. KubeVirt bridge binding gives
# the VM the launcher pod's address, and that address is reachable from the
# NODE but NOT from other pods (measured: node->VM:6443 = 200 in 0.3ms,
# pod->VM:6443 = timeout). The worker's k3s agent is itself a pod, so when
# the supervisor hands it the CP's pod IP it dials an address it can never
# reach and eats the full 10.000s dial timeout – and the kubelet does not
# start until that expires. Advertising the VIP, which every VM already
# reaches, makes the first dial succeed: median time-to-ready 50.0s -> 42.7s,
# run-to-run spread 6.1s -> 1.6s. Must run after KThrees writes config.yaml.
- "echo 'advertise-address: 172.18.255.215' >> /etc/rancher/k3s/config.yaml"
```

The same file also carries the ghost node of Chapter 7 — the leftover `Node` object baked into the image, and the asynchronous cleanup whose lateness let a stopwatch stop early:

```
# 03-target-cluster/target-cluster-warm.yaml (excerpt – KThreesControlPlane.postK3sCommands)
postK3sCommands:
  # Reap the bake VM's leftover Node object (carried in the warm SQLite
  # state) once the API is up. Transient unit so it survives cloud-init.
  - systemd-run --unit=warm-ghost-node-cleanup --no-block /bin/bash -c 'for i in $(seq 1 60);
do /usr/local/bin/k3s kubectl delete node ubuntu-bake-vm-warm --ignore-not-found && break; sleep 2;
done'
```

G.3 supportContainerResources — the second one-line fix

Chapter 8's second fix, worth 8.2 seconds on every cluster build, and the one that is cluster state rather than repository state — which is why Chapter 19 found it missing after the rebuild.

```
# scripts/configure-kubevirt-perf.sh (excerpt)
{"spec":{"configuration":{"supportContainerResources":[
  {"type":"container-disk",
    "resources":{"requests":{"cpu":"100m","memory":"40M"},"limits":
{"cpu":"${CPU_LIMIT}","memory":"100M"}}},
  {"type":"guest-console-log",
    "resources":{"requests":{"cpu":"100m","memory":"60M"},"limits":
{"cpu":"${CPU_LIMIT}","memory":"100M"}}}
]}}}
```

G.4 Two ways to count a failover

Chapter 11's sharpest finding is a comparison, so both sides belong here. First, the demonstration script — which counts any non-empty response body as a success, and therefore prints `success=20 failed=0` when both clusters are down:

```
# 07-istio-advanced/act3-multicluster/run.sh (excerpt)
call() { $K1 exec -n "$NS_APPS" deploy/client-trusted -c curl -- \
    curl -s --max-time 10 http://echo/ 2>/dev/null || true; }

# ...

b=$(call)
if [[ -n "$b" ]]; then okc=$((okc+1)); else failc=$((failc+1)); fi
done
detail "success=${okc} failed=${failc}"
if (( okc >= 18 )); then
    ok "Traffic failed over to cluster2 with zero config change and no app restart"
else
    fail "Failover incomplete: ${failc}/20 requests failed"
fi
```

And the web console, checking the same thing correctly — on the status code:

```
// ui/backend/handlers/cross_cluster.go (excerpt)
ok := func(p ProbeResult) bool { return p.StatusCode >= 200 && p.StatusCode < 300 }

httpbinAlive := ok(probes[0])
nginxAlive := ok(probes[2])
// Failover: a local service is down but the cross-cluster path is still up
failoverActive := (!httpbinAlive && ok(probes[1])) || (!nginxAlive && ok(probes[3]))
```

Two checks of one behaviour, in one repository, written weeks apart. One cannot fail.

G.5 The egress allowlist that allows more than it names

Chapter 14 grades the eight architecture constraints. This is the file behind the verdict on the egress rule: the last two rules carry ports but no destination, so they permit those ports to *anywhere*.

```

# 06-symposium/agent-egress-ollama.yaml (quoted whole)
---
# AgentRun pods are labeled app.kubernetes.io/part-of=symposium, which makes
# them subject to the helm-installed `symposium-allow-otel` NetworkPolicy
# (policyType: Egress, only ports 4317/4318 allowed). With kindnet enforcing,
# that implicitly blocks egress to Ollama on the docker bridge.
#
# This policy adds the missing egress paths (DNS + Ollama) for symposium pods.
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: symposium-allow-ollama
  namespace: symposium-system
spec:
  podSelector:
    matchLabels:
      app.kubernetes.io/part-of: symposium
  policyTypes:
    - Egress
  egress:
    # DNS
    - to:
      - namespaceSelector: {}
        podSelector:
          matchLabels:
            k8s-app: kube-dns
      ports:
        - port: 53
          protocol: UDP
        - port: 53
          protocol: TCP
    # Ollama via the host-ollama shim Service (Endpoints → 172.18.0.1:11434)
    - ports:
      - port: 11434
        protocol: TCP
    # Kubernetes API (skill-k8s-ops needs it)
    - ports:
      - port: 443
        protocol: TCP
      - port: 6443
        protocol: TCP

```

G.6 An agent as a Kubernetes object

Chapter 13's central claim is that an agent is an object like any other. This is one, whole — model, endpoint, tools, policy, and the briefing it is given.

```
# 06-symposium/cluster2-agent.yaml (excerpt – the Agent object)
apiVersion: symposium.ai/v1alpha1
kind: Agent
metadata:
  name: cluster2-agent
  namespace: symposium-system
  annotations:
    symposium.ai/description: "Cluster2 management-plane agent (KubeVirt VMs, CDI, CAPI, cluster
nodes)"
spec:
  agents:
    default:
      model: qwen2.5:7b
      baseURL: http://host-ollama.symposium-system.svc.cluster.local:11434/v1
      sandbox:
        enabled: false
  authRefs:
    - provider: ollama
      secret: llm-credentials
  skills:
    - skillPackRef: web-endpoint    # drives declarative serving (requiresServer=true)
    - skillPackRef: k8s-ops        # kubectl/virtctl sidecar for task (chat) runs
  policyRef: sandbox-restricted
  observability:
    enabled: false
  memory:
    enabled: false
    maxSizeKB: 256
    # Environment briefing – keep SHORT (7B model). Edit here; re-run the
    # installer (or kubectl apply -k 06-symposium/) to reproduce.
    systemPrompt: |
      You operate "cluster2", the management cluster of a KubeVirt multi-cluster platform.
Topology:
  - cluster2 (where your kubectl/virtctl run): a Kind cluster hosting Cluster API (CAPI),
KubeVirt + CDI (VMs run as pods), MetalLB, and Symposium. Your in-cluster tools talk to THIS
cluster's API only.
  - target-cluster: a k3s cluster whose nodes are KubeVirt VirtualMachines hosted on cluster2,
provisioned by CAPI. Inspect from cluster2 with: kubectl get clusters,machines -A (group cluster.x-
k8s.io); kubectl get virtualmachines,virtualmachineinstances -A (group kubevirt.io); kubectl get
datavolumes -A (CDI). target-cluster API VIP is 172.18.255.215.
  - cluster1: a SEPARATE Kind cluster (Istio ambient mesh, httpbin/sleep). You have NO kubectl
access to it from here.
```

Always use full resource names (virtualmachines, virtualmachineinstances, clusters, machines). Read-only unless the user explicitly approves a change.

Read that briefing against Chapter 19. It describes `cluster1` as running `httpbin/sleep` — which was true before the September rebuild and is not true now. The agent is told a fact about its own environment that has expired, and nothing in the platform notices.

G.7 The documented knob that controlled nothing

Chapter 19's most quotable find. The outer variable, with the warning added after the discovery:

```
# bake-common.sh (excerpt – header comment)
#     K3S_VERSION    – DOES NOT actually control the bake (found 2026-09-16,
#                     upgrading v1.31.4+k3s1 -> v1.37.0+k3s1): the version
#                     that reaches the VM is a SEPARATE hardcoded literal
#                     inside the embedded /usr/local/bin/bake.sh cloud-init
#                     content below (search this file for the second
#                     K3S_VERSION= assignment) – that content block is
#                     written to the VM verbatim, so this outer bash
#                     variable never reaches it. Keep both literals in sync
#                     by hand until someone wires the interpolation through.
#                     (default here: v1.37.0+k3s1)

K3S_VERSION="${K3S_VERSION:-v1.37.0+k3s1}"
```

And the literal that actually reaches the machine, a hundred lines below it. The `<< 'CLOUDINIT'` — note the quotes — is what makes the block above inert: a quoted heredoc expands nothing.

```
# bake-common.sh (excerpt – the embedded cloud-init content)
cat << 'CLOUDINIT'
# ...
- path: /usr/local/bin/bake.sh
  content: |
    #!/bin/bash
    set -e
    K3S_VERSION="v1.37.0+k3s1"
    # ...
    curl -sfl https://get.k3s.io | \
      INSTALL_K3S_VERSION="$K3S_VERSION" \
      INSTALL_K3S_SKIP_START=true \
```

Two assignments of one name, in one file, a hundred lines apart. Setting the documented one changes nothing at all.

G.8 The dead image, and the patch that survives a rebuild

Chapter 19's first real obstacle, and the fix that now runs automatically — quoted with its comment, because the comment records the uncertainty honestly rather than claiming a clean diagnosis.

```
# 02-capi-init/init-management-cluster.sh (excerpt)
# ship their kube-rbac-proxy sidecar as `gcr.io/kubebuilder/kube-rbac-proxy:v0.16.0`,
# a reference that no longer resolves (ImagePullBackOff, "not found") --
# apparently a dead/retired gcr.io path, though clusterctl fetches the release
# manifest fresh on every `init` so it's unclear whether this was ever pullable
# from this repo's install path. Whatever the history, the *deployment* that
# was actually running before this rebuild used
# `quay.io/brancz/kube-rbac-proxy:v0.16.0` (confirmed pullable), so that's the
# known-good pin here. Without this, both provider Deployments sit at 1/2
# ready forever and the `kubectl wait --for=condition=available` calls below
# silently time out into their `|| true` fallback -- so `make capi-init`
# reports success while KThreesControlPlane's mutating webhook is dead,
# and the first `target-cluster` deploy fails opaquely on
# "failed calling webhook ...kthreescontrolplane: connection refused".
for ns in capi-k3s-bootstrap-system capi-k3s-control-plane-system; do
    deploy="$(kubectl get deploy -n "$ns" -o name 2>/dev/null | head -1)"
    [[ -n "$deploy" ]] || continue
    if kubectl get "$deploy" -n "$ns" -o jsonpath='{.spec.template.spec.containers[?(@.name=="kube-
rbac-proxy")].image}' 2>/dev/null \
        | grep -q '^gcr.io/kubebuilder/'; then
        echo "==> Patching $deploy ($ns): kube-rbac-proxy gcr.io image -> quay.io/brancz (known
pullable)"
        kubectl patch "$deploy" -n "$ns" --type=strategic \
            -p='{ "spec": { "template": { "spec": { "containers": [ { "name": "kube-rbac-
proxy", "image": "quay.io/brancz/kube-rbac-proxy:v0.16.0" } ] } } } }'
    fi
done
```

The `|| true` that the comment mentions is the pattern Appendix C.2 catalogues: a wait that fails, a fallback that swallows the failure, and a script that reports success.

Appendix H — Sources

The book draws on three kinds of source, and they are not equally checkable. This appendix says which is which.

H.1 The companion repository

The **Build log** sections come from the platform repository, `sovereign_cloud`. It is the only source a reader can run. Every build-log claim names the file or command behind it in that chapter's **Open the repo** block, and the load-bearing files are quoted in Appendix G.

One caveat, recorded in Chapter 19's draft notes: most of the material this book cites lives on the working branch `upgrade/k8s-1.37-istio-1.31-symposium-0.10.75`, not on `main`. A reader who clones the default branch will not find much of it.

H.2 Briefing material (not public, not in the repository)

Four items are the author's own briefing material. They are described here rather than linked, because no path or URL would resolve for a reader. Where the book quotes them, it quotes them verbatim and says so.

Referred to in the book as	What it is
The strategy deck	<i>Reclaiming Margins with Sovereign Cloud</i> , a strategic briefing deck dated February 2026, in five chapters. Parts I and V adapt its argument; Chapter 20 takes its TCO slide apart. None of its figures carries a source, and the book says so each time it quotes one.
The sovereign-cloud flyer	A one-page campaign flyer, version 2, unbranded — six promises and a " <i>Save \$Millions — Creating a Cloud Owning Culture</i> " band. The epilogue reads it claim by claim. A separate version of the same flyer carries a real company's branding; the book does not quote that one, and does not name the company.
The recorded factory demonstration	A screen recording, 8 minutes 56 seconds, of the cluster-provisioning path in Chapters 4 and 5. Used only as a source of stills.
The recorded agent demonstration	A screen recording, 2 minutes 50 seconds, of the agent console in Chapter 13. Used only as a source of stills.

The cross-cluster mind map in Chapter 10 was previously cited from this material as well. It is in the repository, at `05-istio/cross-cluster-mindmap.html`, and the book now cites it there. The briefing copy is byte-identical, which is exactly Chapter 10's point: a diagram duplicated into two places goes stale in both.

H.3 The public record

Public record sections cite published, checkable sources: the Andreessen Horowitz cloud-cost paper (2021), Dropbox's 2018 IPO filing, 37signals' published cloud-exit accounts, Gartner's public-cloud spending forecast, Flexera's annual cloud reports, the EU Data Act, the AT&T/Broadcom litigation coverage, Grand View Research's sovereign-cloud market sizing, and the hyperscalers' own server-depreciation disclosures. Each is named where it is used, with its date.

Their verification status, as re-checked on **18 September 2026**:

Source	Status
Andreessen Horowitz, <i>The Cost of Cloud, a Trillion-Dollar Paradox</i>	Verified against the article itself — date (27 May 2021), authors, the 50%-of-cost-of-revenue benchmark, the one-third-to-one-half repatriation figure, and the \$100B / >\$500B market-value estimates all confirmed as quoted
37signals, <i>Our cloud exit savings will now top ten million over five years</i>	Verified against the post itself — date (17 October 2024), \$3.2M/year run rate, \$1.3M, about \$700,000 of servers, "well over ten million dollars over five years", and the "no hidden dragons" sentence all confirmed as quoted
EU Data Act, Regulation (EU) 2023/2854	Partly verified, and one error corrected. In force 11 January 2024; obligations applicable 12 September 2025; charges withdrawn 12 January 2027. An earlier draft misread the in-force date as the date cost-only charging began (Chapter 2). Still to verify against the EUR-Lex text rather than commentary, and to have counsel review
Gartner, worldwide public cloud end-user spending forecast for 2025	Not verified — the press release returns HTTP 403 to direct retrieval. The \$723.4B figure comes from the citation as previously recorded. Confirm before publication
Dropbox 2018 IPO filing (\$74.6M Infrastructure Optimization saving)	Not re-verified — currently cited via secondary coverage. Cite the S-1 directly
Flexera annual cloud reports (29% waste, 85%)	Not re-verified this pass — URLs recorded in the prologue's draft notes
AT&T / Broadcom litigation and settlement	Not re-verified — trade-press coverage only. Cite the court filing if possible
Grand View Research, sovereign-cloud market sizing	Not verified — the report page returned HTTP 403 when attempted; figures come from a search-result summary
Hyperscaler server-depreciation disclosures (Microsoft FY2022, Alphabet 2023, Amazon Feb 2025)	Not verified — secondary coverage summarising filings. Cite the 10-K/10-Q text
GDPR, Schrems II, CLOUD Act, India DPDP Act	Not verified — stated from general knowledge. Have counsel review (Chapter 3)

A search summary or a news report is not the filing it describes, and every row above marked *not verified* is a row where the book is currently trusting one.